Automating and accelerating scientific discovery in HEP with generative models

Tobias Golling

UNIVERSITÉ
DE GENÈVE

FACULTY OF SCIENCE

The Large Hadron Collider (LHC)

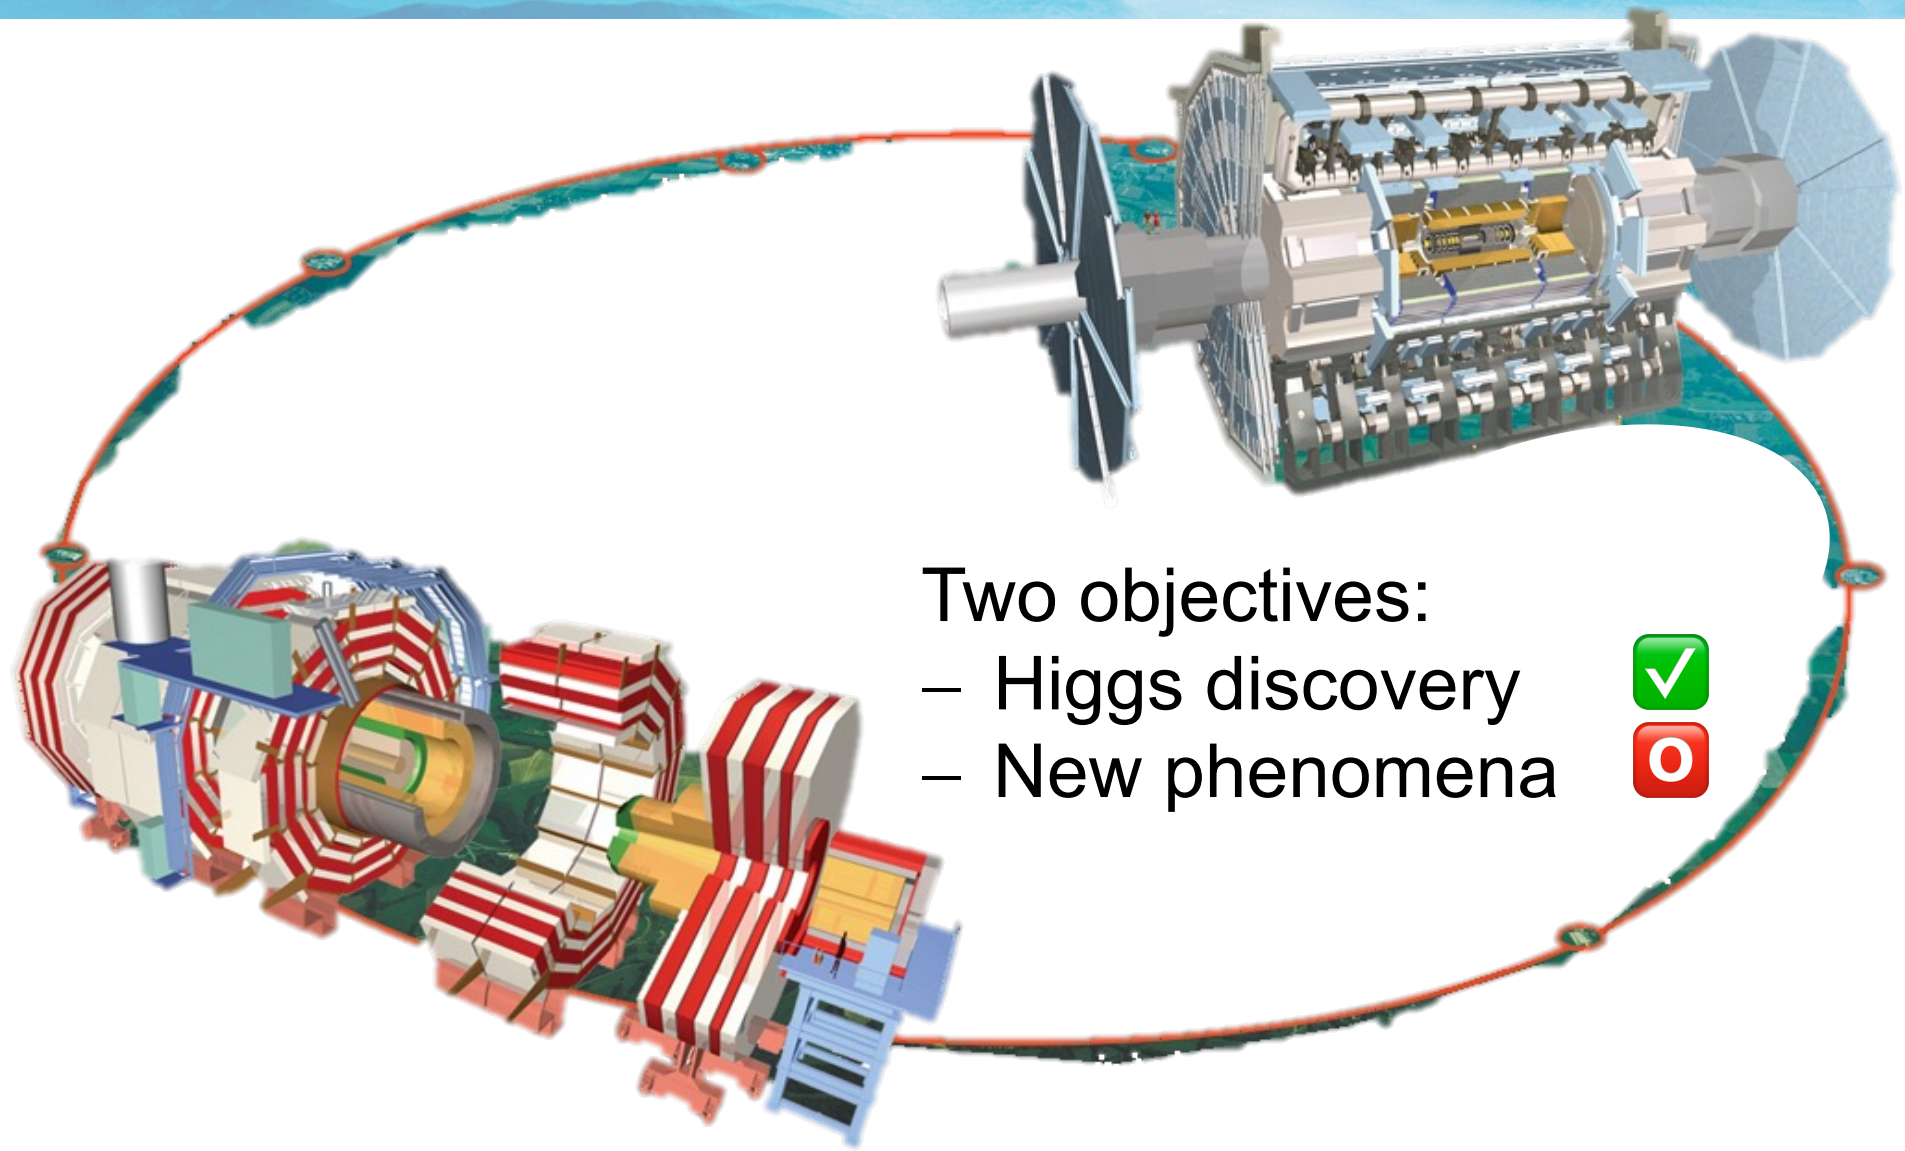

Two objectives:

- Higgs discovery ✅
- New phenomena 🅾️

The SM* is complete

Why keep going?

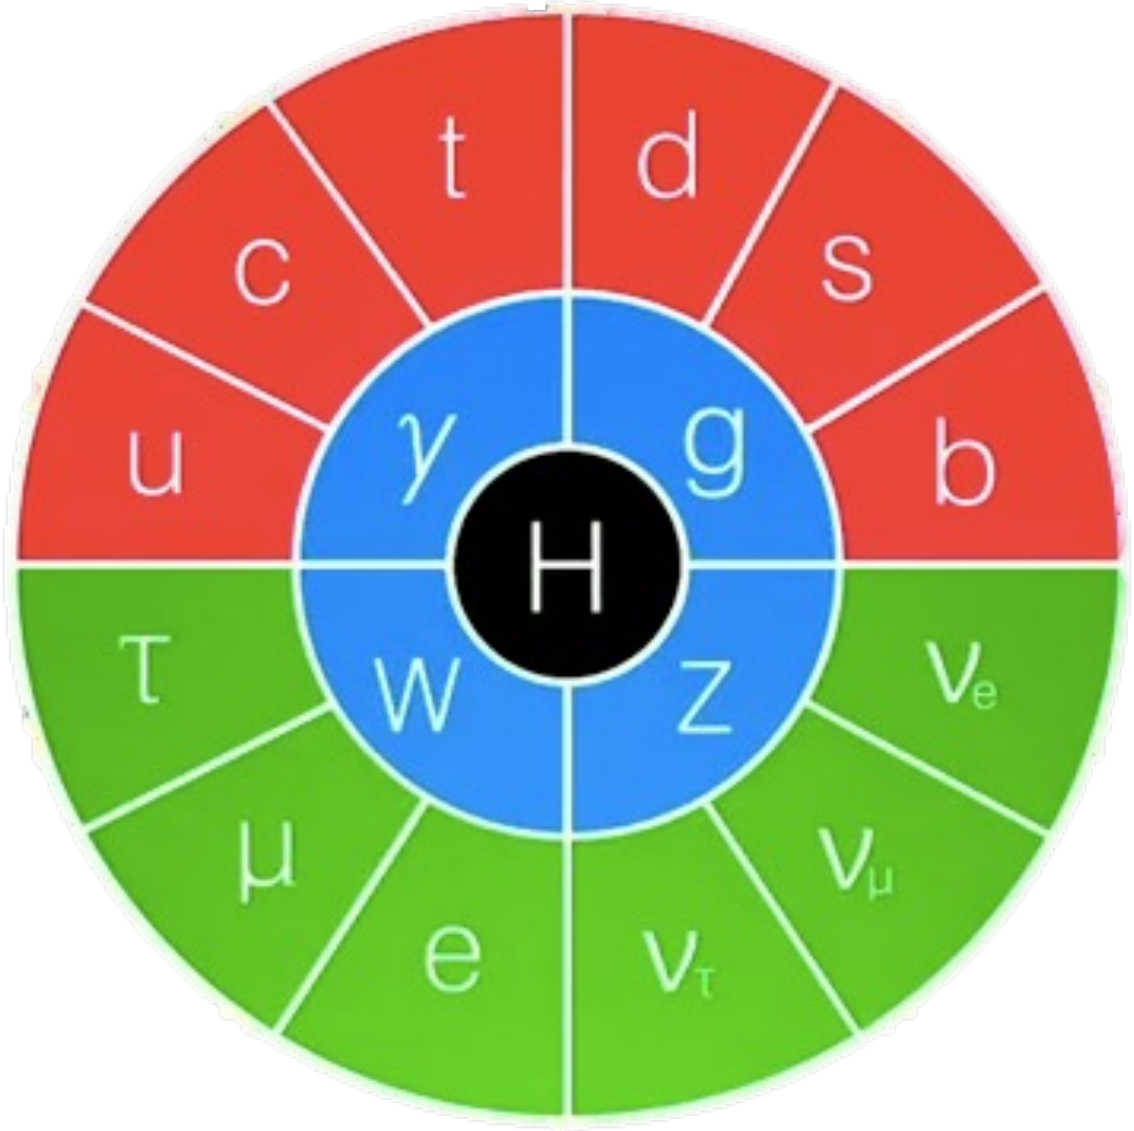

*the Standard Model of particle physics

Open mysteries remain

Unexplained observed phenomena

- Dark matter
- Dark energy
- Matter-antimatter asymmetry

Unsatisfactory SM

- Quantum gravity, naturalness,…

The theory guidance

Hypothesize SM extensions

Addressing SM shortcomings

→ *Testable* predictions

Plethora of BSM extensions

How to probe vast space

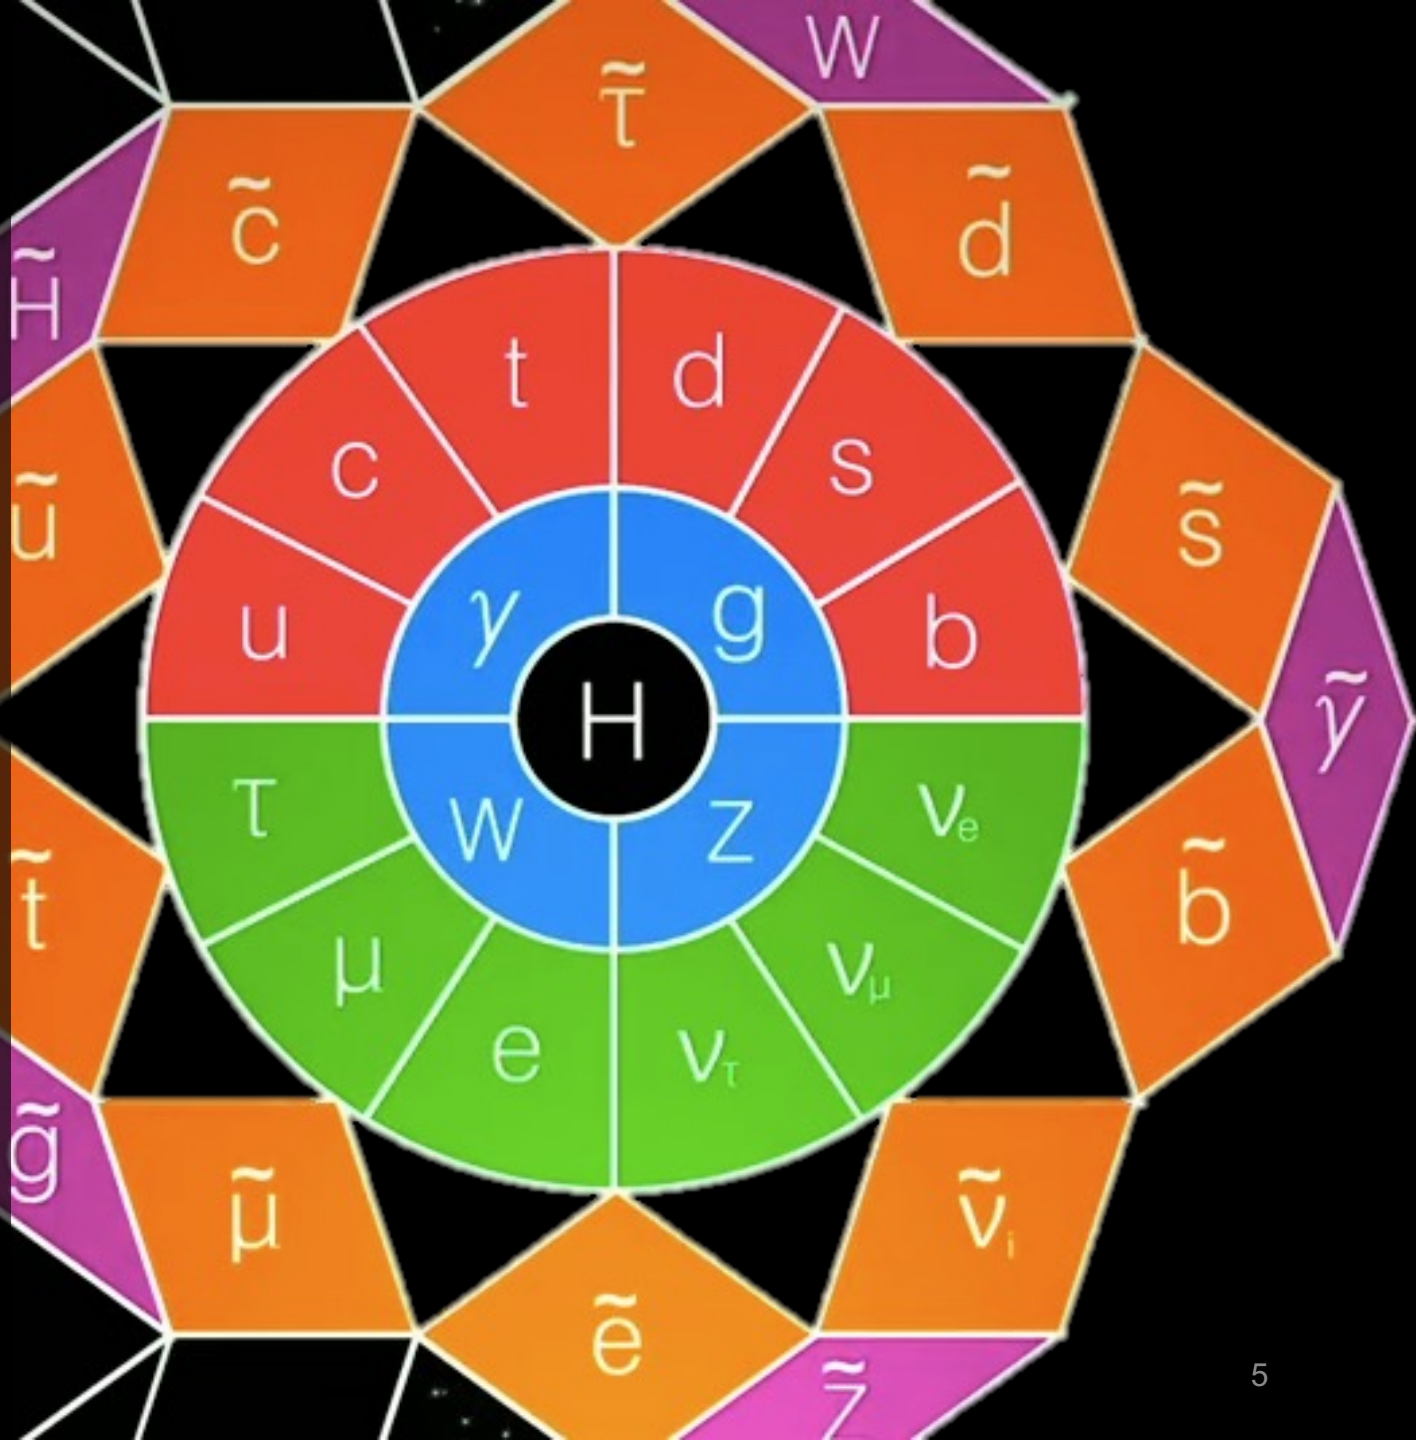

[Beyond-the-SM physics = BSM]

The need for synthetic data

Too complex to predict outcome from first principles:
Monte Carlo simulation

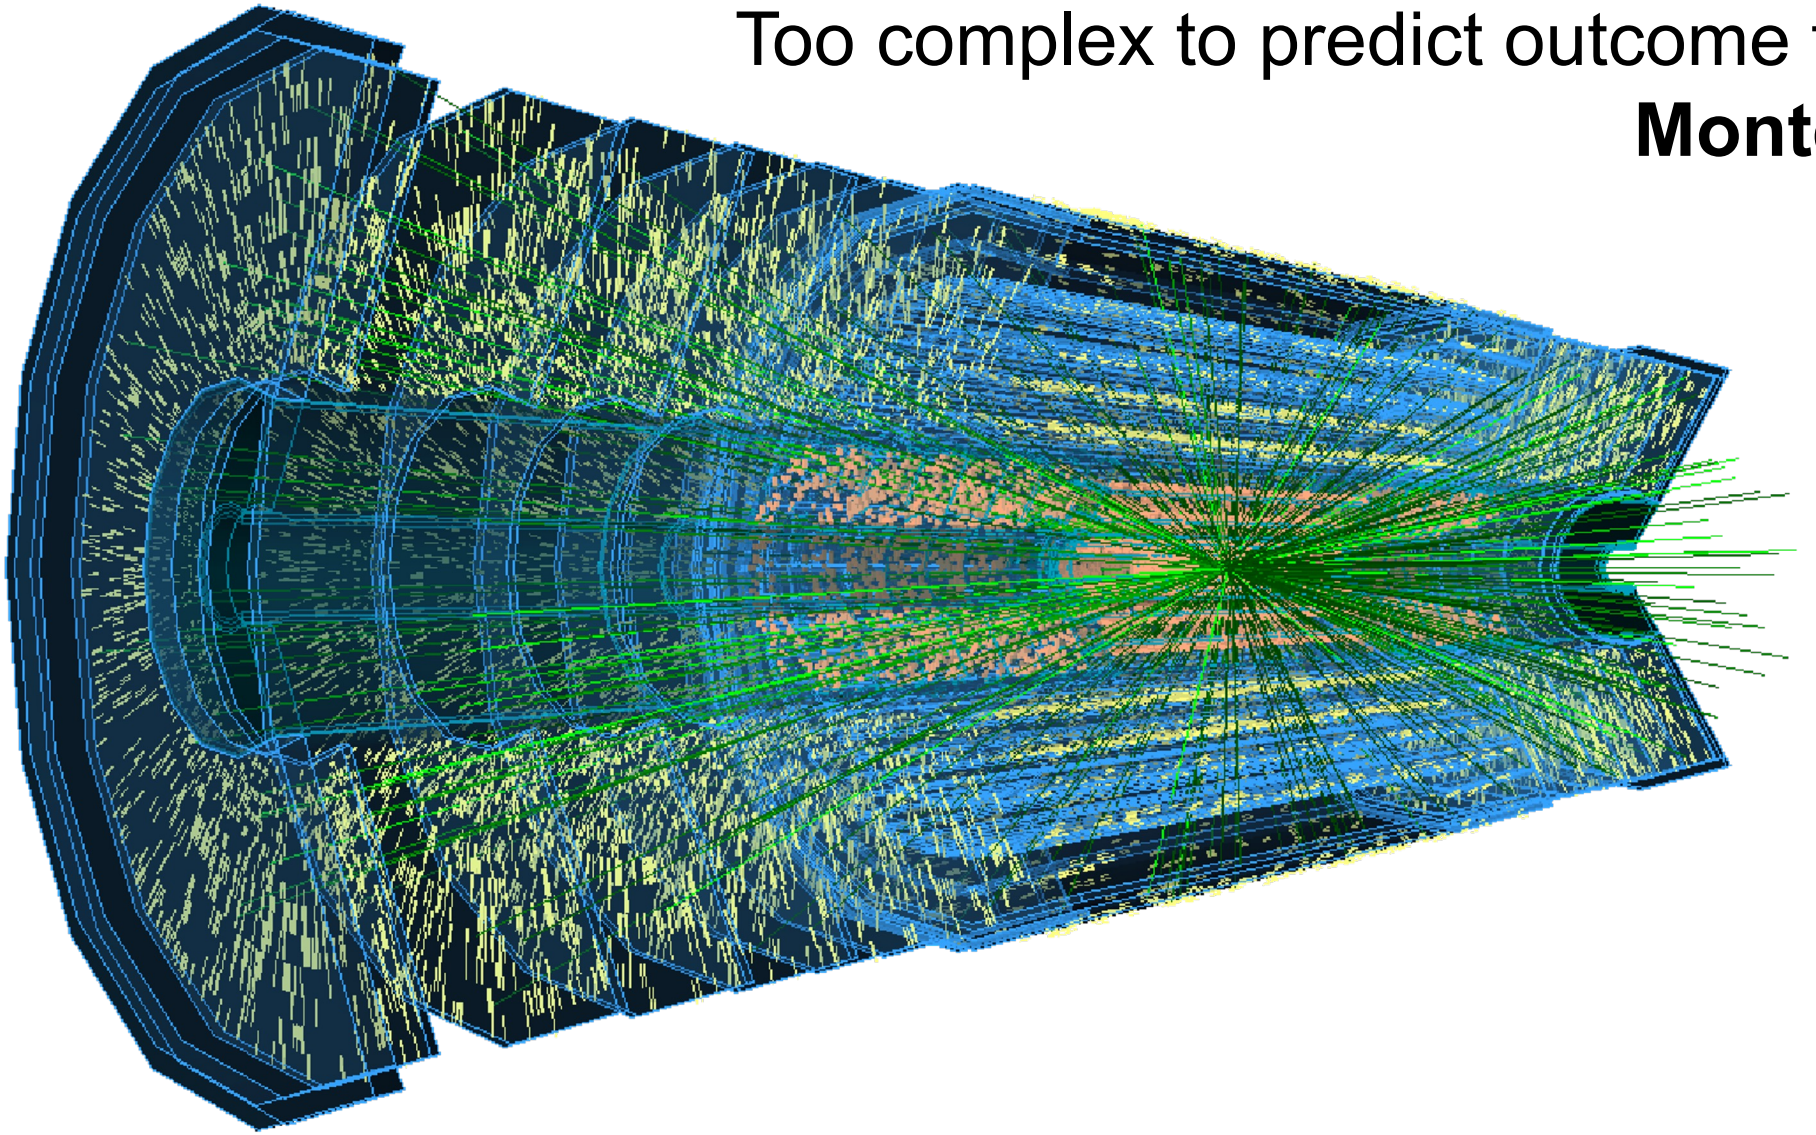

[Phys. Rev. D. 90 (2014) 052004]

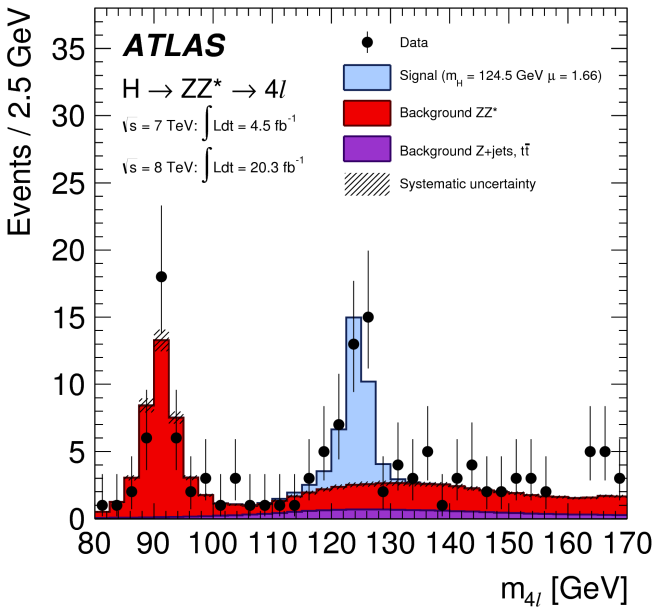

$$p(\text{data}\,|\,\text{theory})$$

LHC interim evaluation

Physics beyond the SM is **not around the corner**

Slow-growth era of LHC: energy & luminosity

Opportunity !
Turning crank → innovation

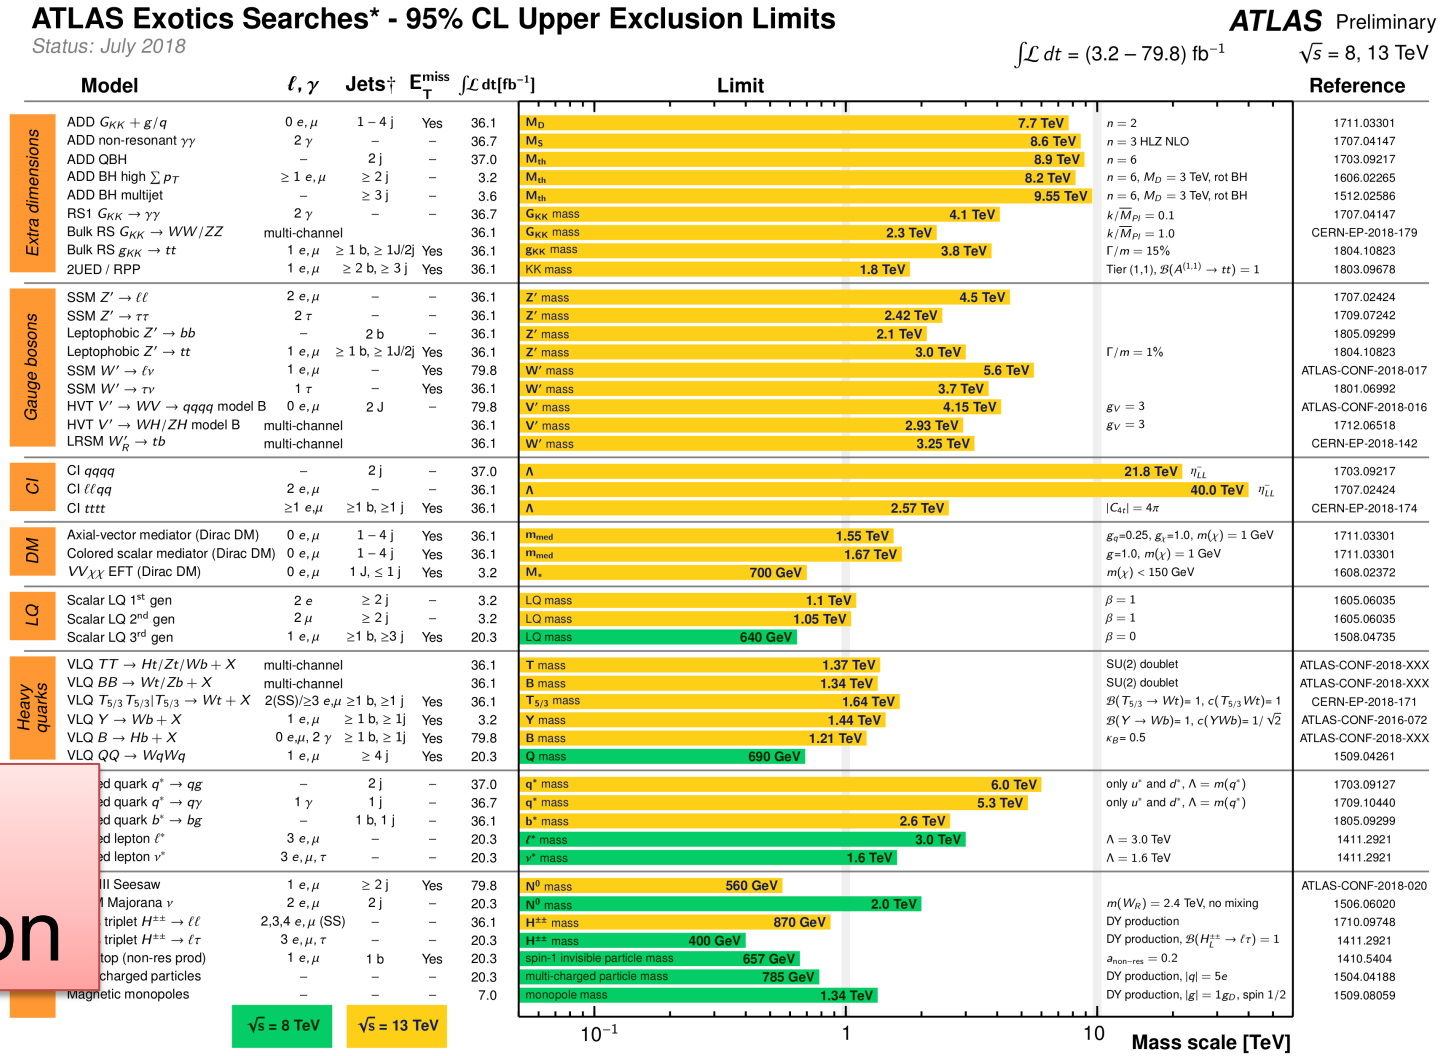

*Only a selection of the available mass limits on new states or phenomena is shown.
†Small-radius (large-radius) jets are denoted by the letter j (J).

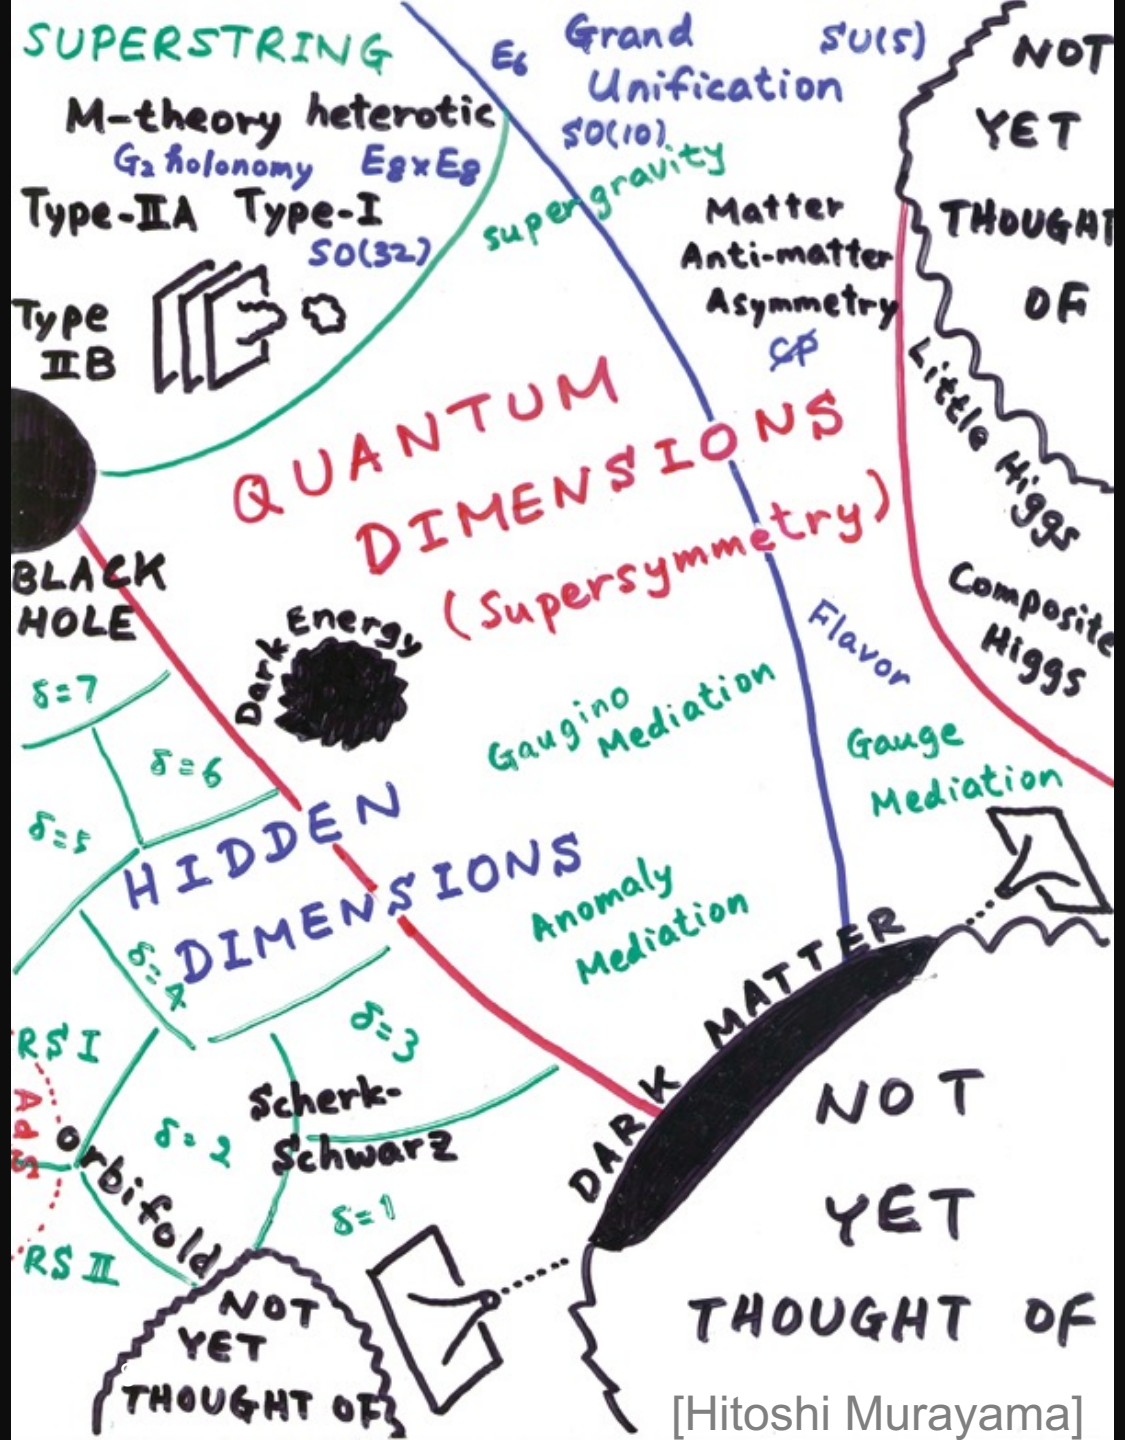

[Hitoshi Murayama]

How to maximize knowledge gain

Given person-power, compute, detector, time

How to invest?

Innovate vs. exploit?

Automating & Accelerating Scientific Discovery

Exploring vast spaces
with generative models

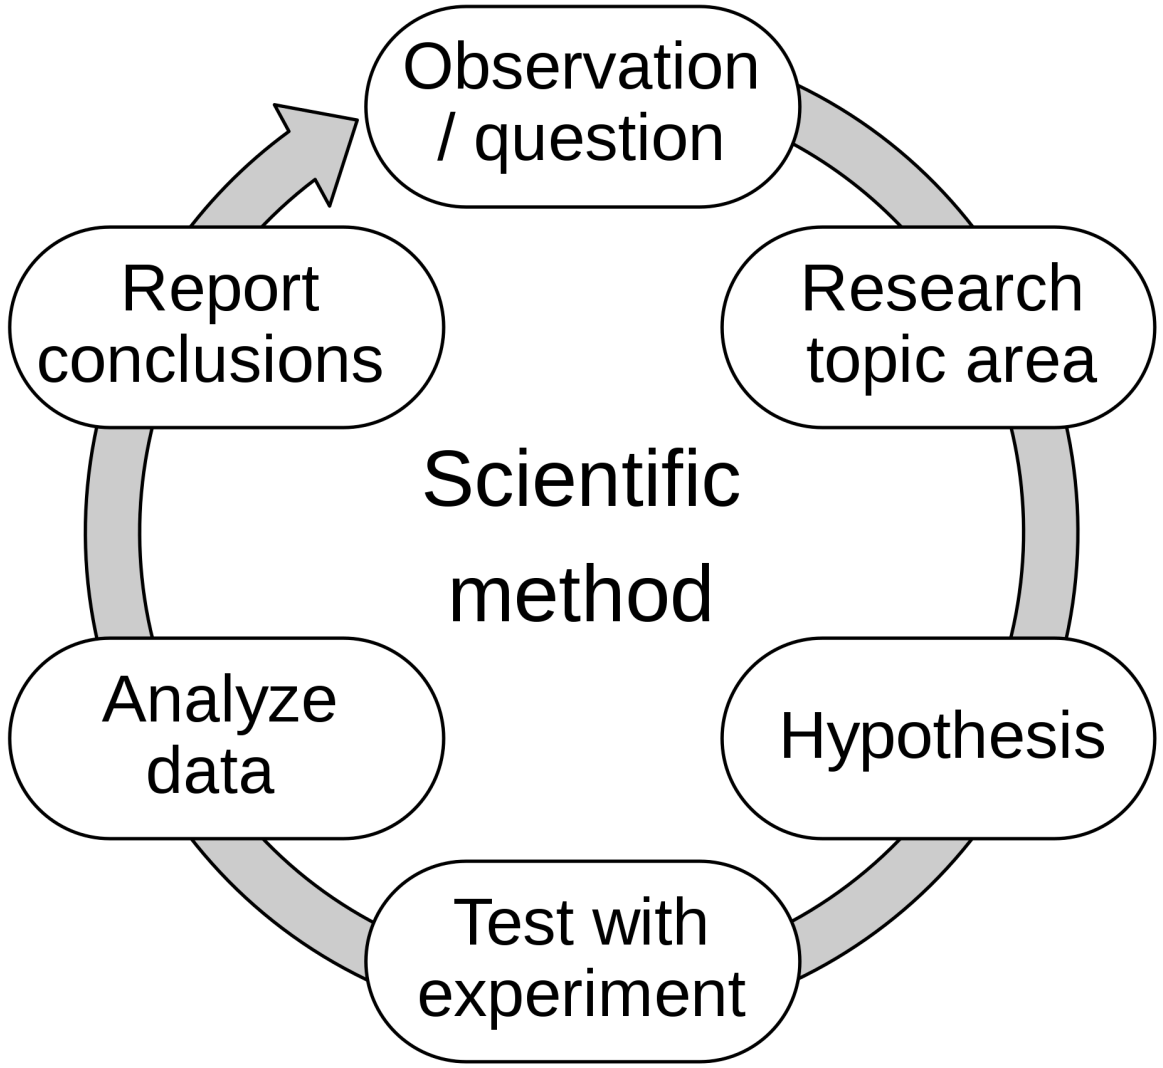

The ML@HEP success story

Constantly growing & cross-connecting

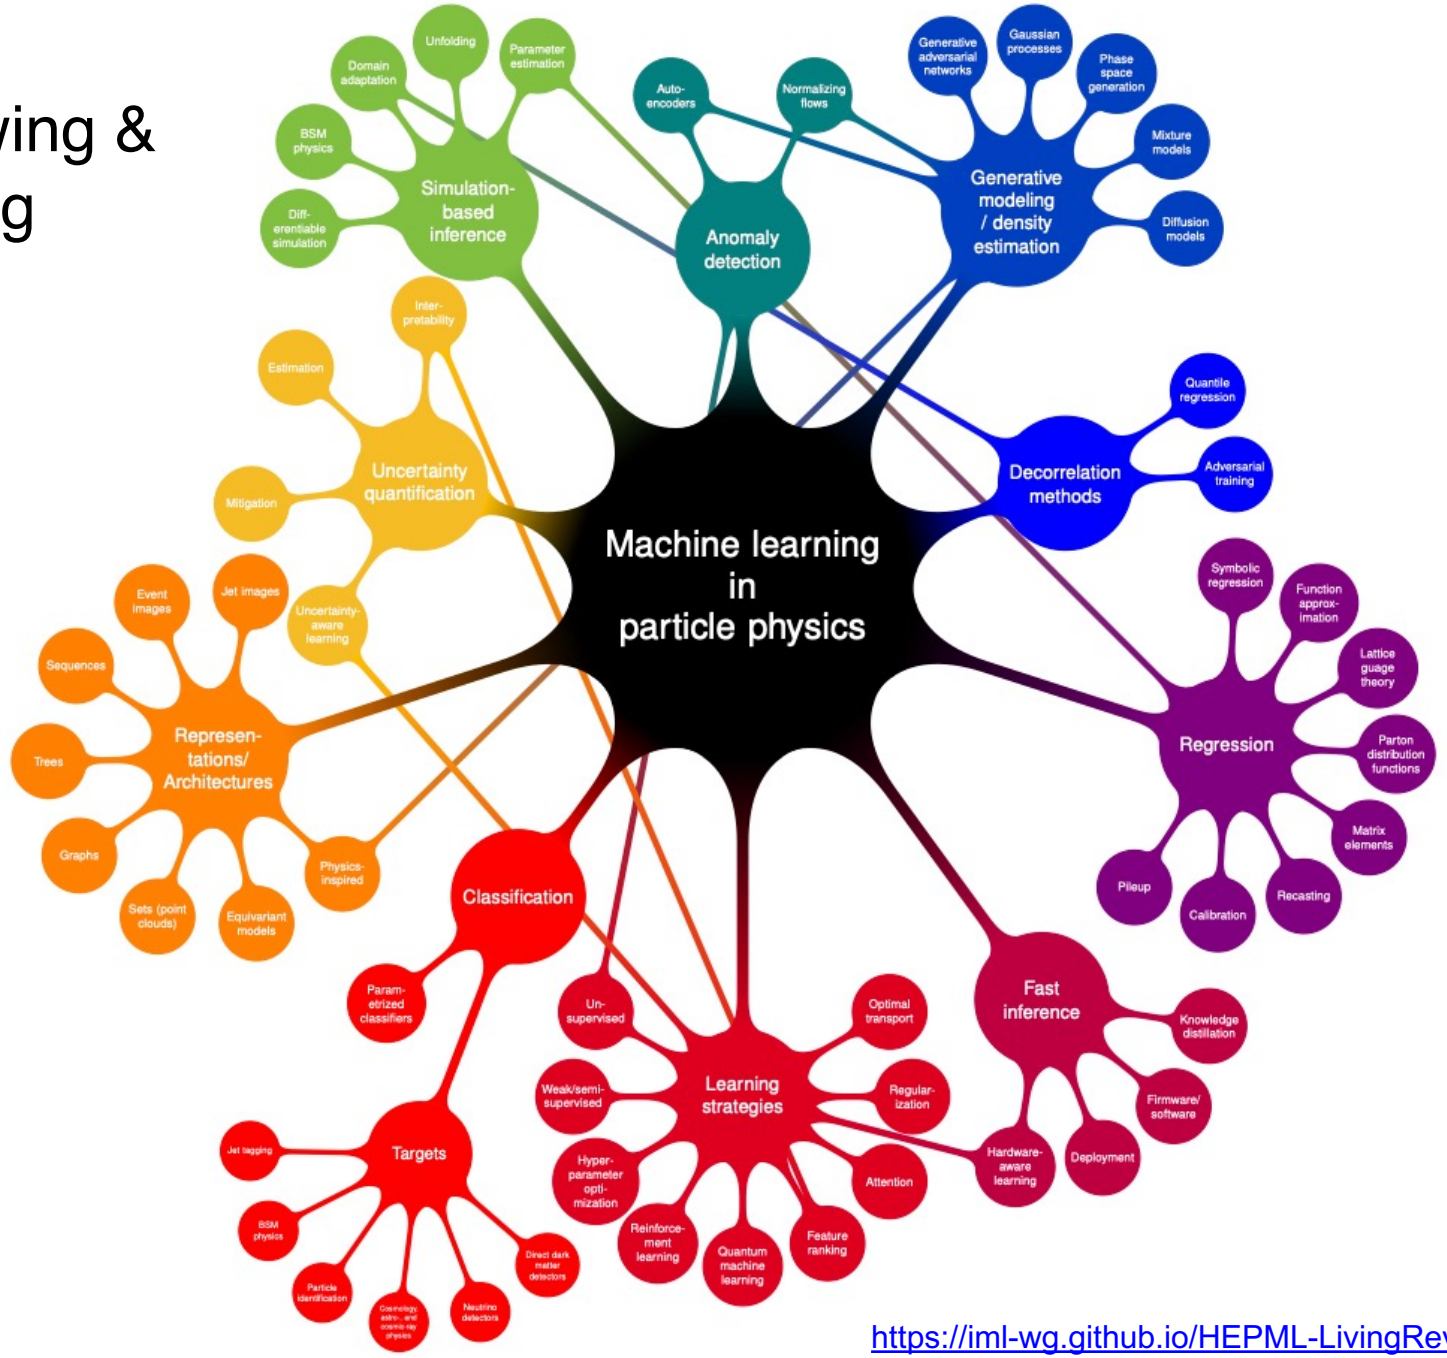

https://iml-wg.github.io/HEPML-LivingReview/
https://github.com/jmduarte/Nomological_Net_ML_Particle_Physics

[Credit: Kazuhiro Terao]

Today: AI/ML everywhere in our workflow

Differentiable Surrogate

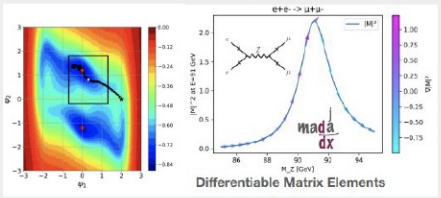

Design Optimization

Generative Models

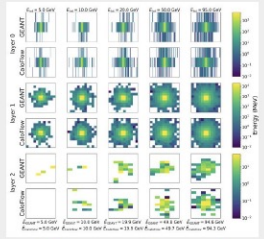

Fast Sim., Stochastic modeling

Ensemble, Bayesian NN, Temperature Scaling

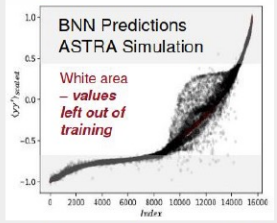

Uncertainty Quantification

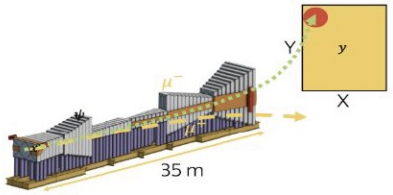

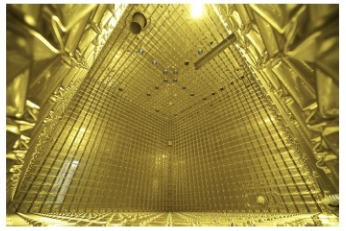

EXPERIMENT DESIGN → BUILDING / INSTALLATION → DATA TAKING / FACILITY OPERATIONS

SIMULATION

HYPOTHESIS BUILDING ← KNOWLEDGE UPDATE ← CALIBRATION PHYSICS INFERENCE

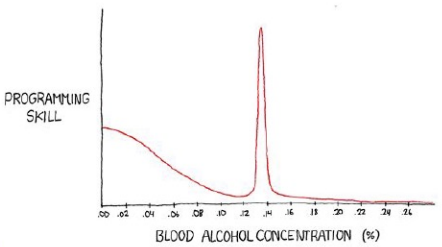

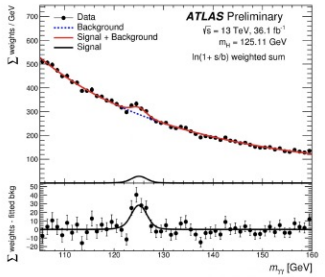

Fast/Edge-ML

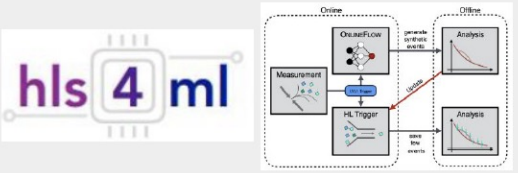

Trigger/Compression

Anomaly Detection

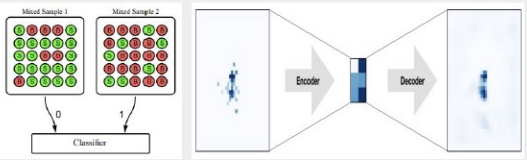

Rare Event, Diagnostics

BO/RL

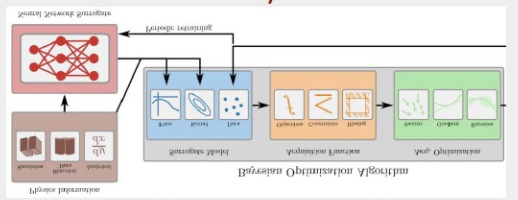

Control Optimization

CV, Geometric ML

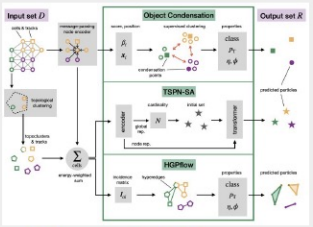

Rare Event, Diagnostics

Flagship ML@HEP

Flavor tagging

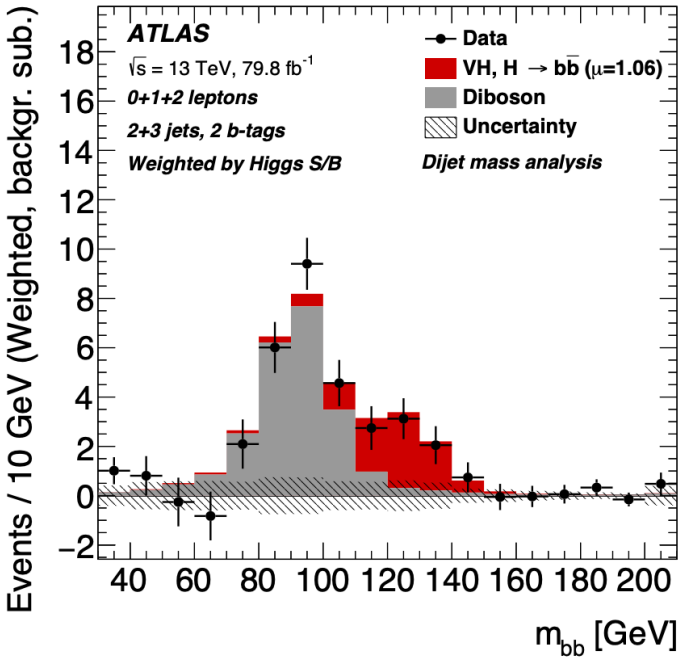

Enabler:

Higgs, top, new phenomena,…

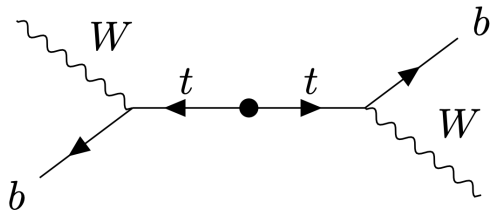

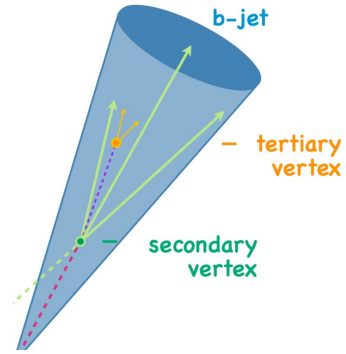

Long history of ML in flavor tagging

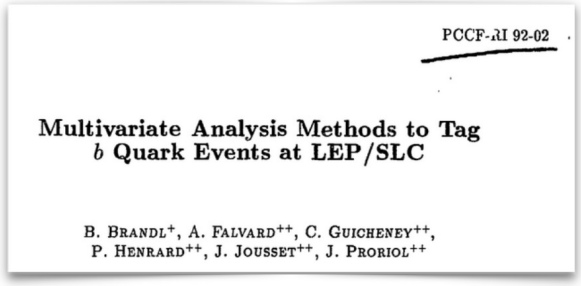
PCCF-RI 92-02

Multivariate Analysis Methods to Tag
b Quark Events at LEP/SLC

B. Brandl+, A. Falvard++, C. Guicheney++,
P. Henrard++, J. Jousset++, J. Proriol++

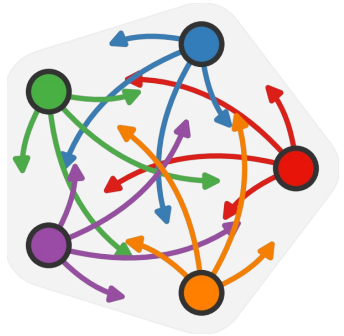

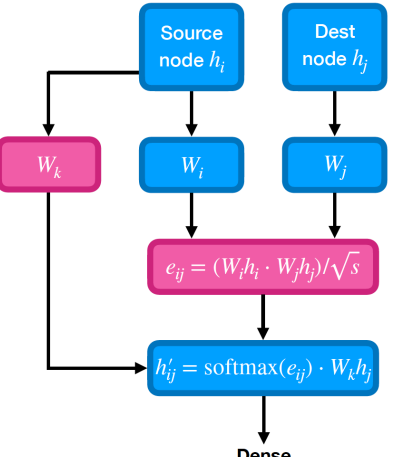

[1706.03762]

1992: Started with an **MLP** @LEP
2005: First ML b-tagging @hadron collider @D0
2007: CDF@Tevatron used **NN**
2012: ML @ATLAS: MV1
2015: **BDT** journey: MV2
2017: Back to **NN**: DL1
2017: CMS DL with DeepCSV
2019: CMS **ParticleNet**
2020: **Deep Sets**
2022: GN1 (**GNN**)
2023: GN2 (**Transformers**)
new training framework

A lot has been learned:
- Flexible multi-classification
- Hand-designed → end-to-end
- Benefit of auxiliary tasks
- Evolving data representations

Evolving data representations in HEP

Arbitrary inputs
FF NN

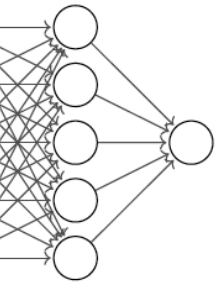

Images
CNN
[1511.05190]

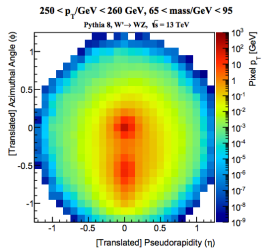

Sequences
RNN
[ATL-PHYS-PUB-2017-003]

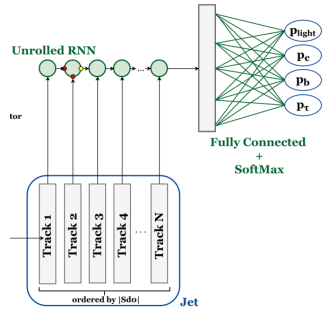

Point clouds

Deep Sets
[1810.05165]

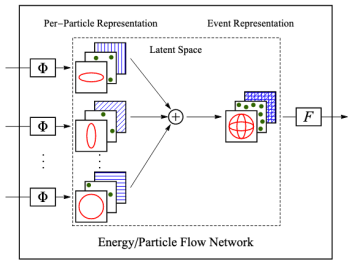

GNN / Transformer
[2203.01189]

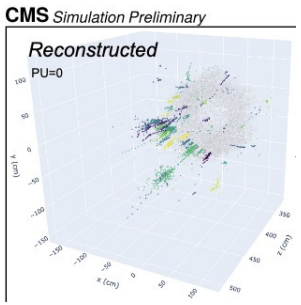

…

Impact on physics

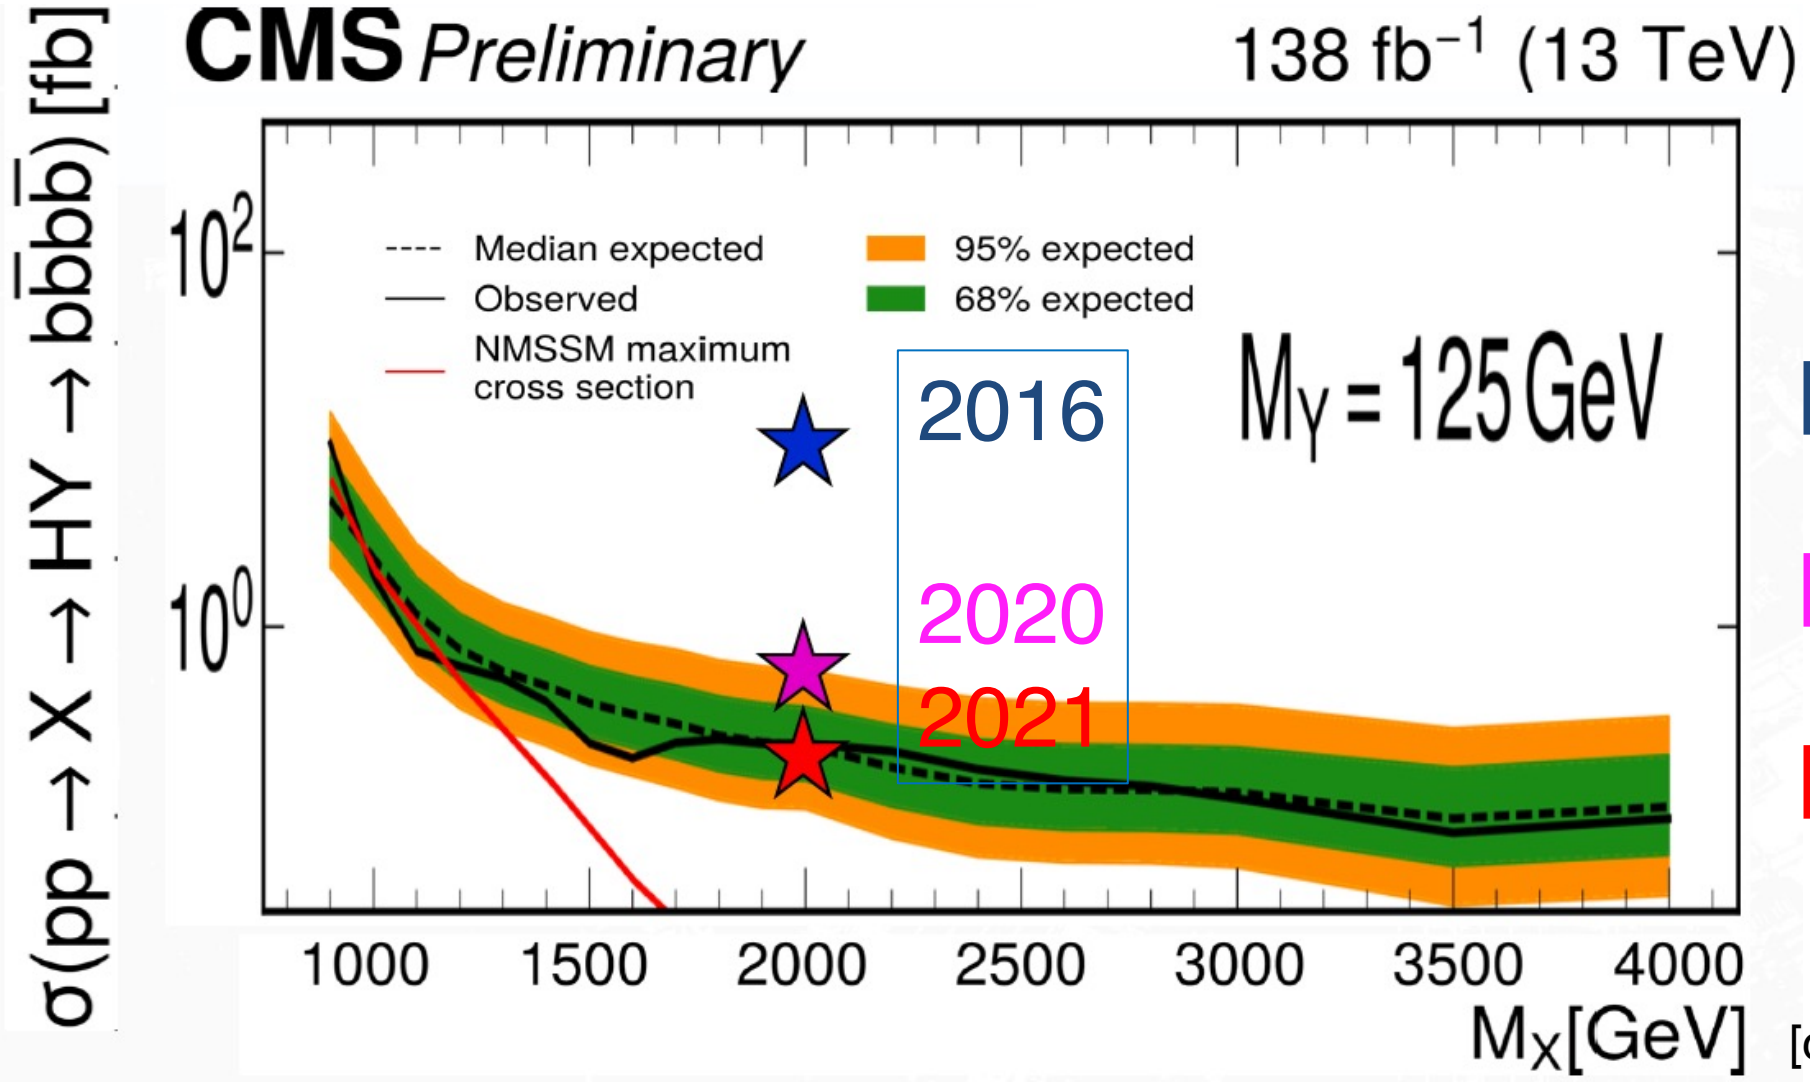

[1710.04960]

[CMS-B2G-20-004]

[2204.12413]

[credit: Loukas Gouskos]

When are we reaching a plateau?

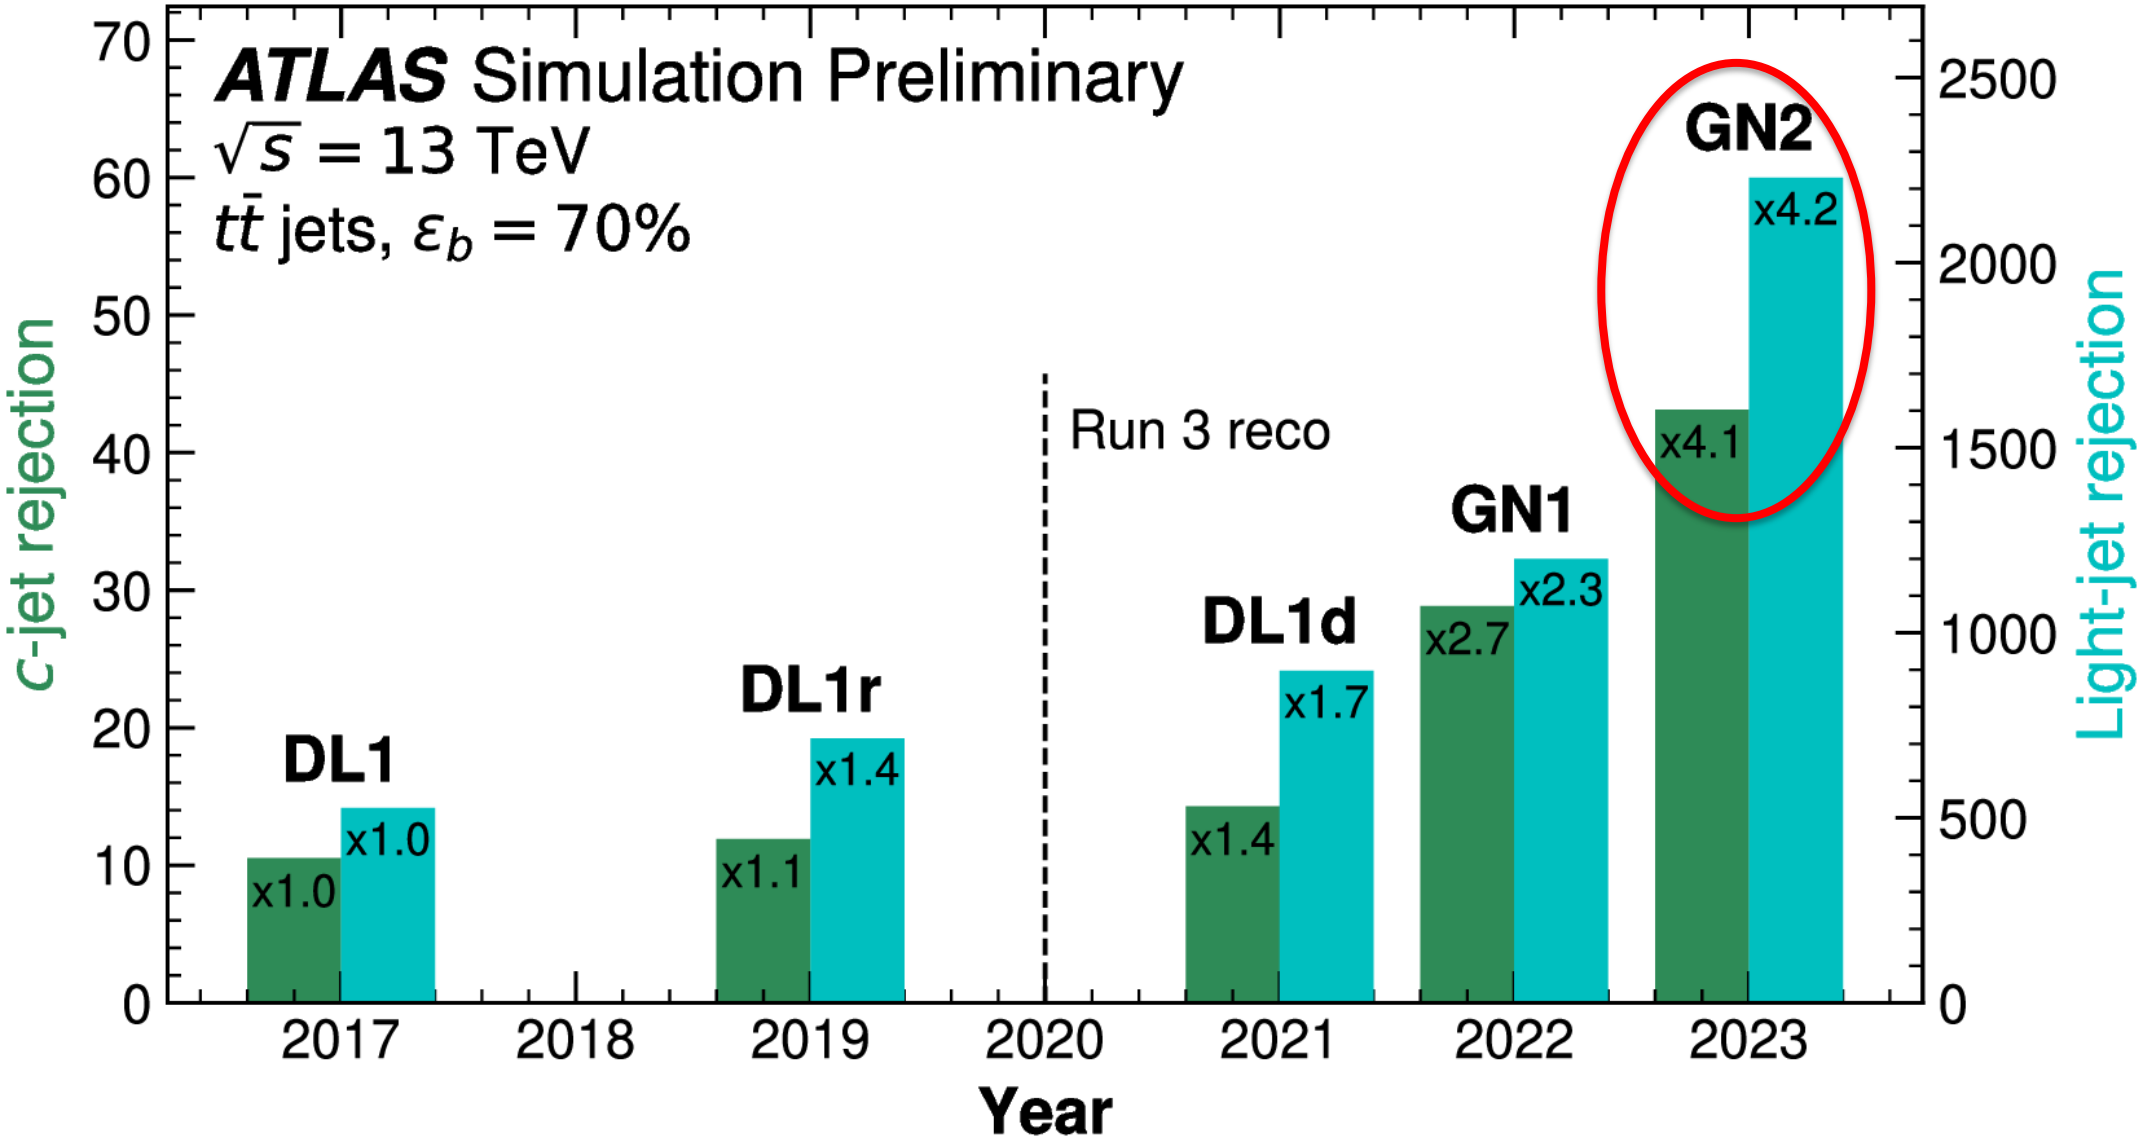

[Transformer-based GN2, see FTAG-2023-01]

Physics-aware AI
[the edge of science]

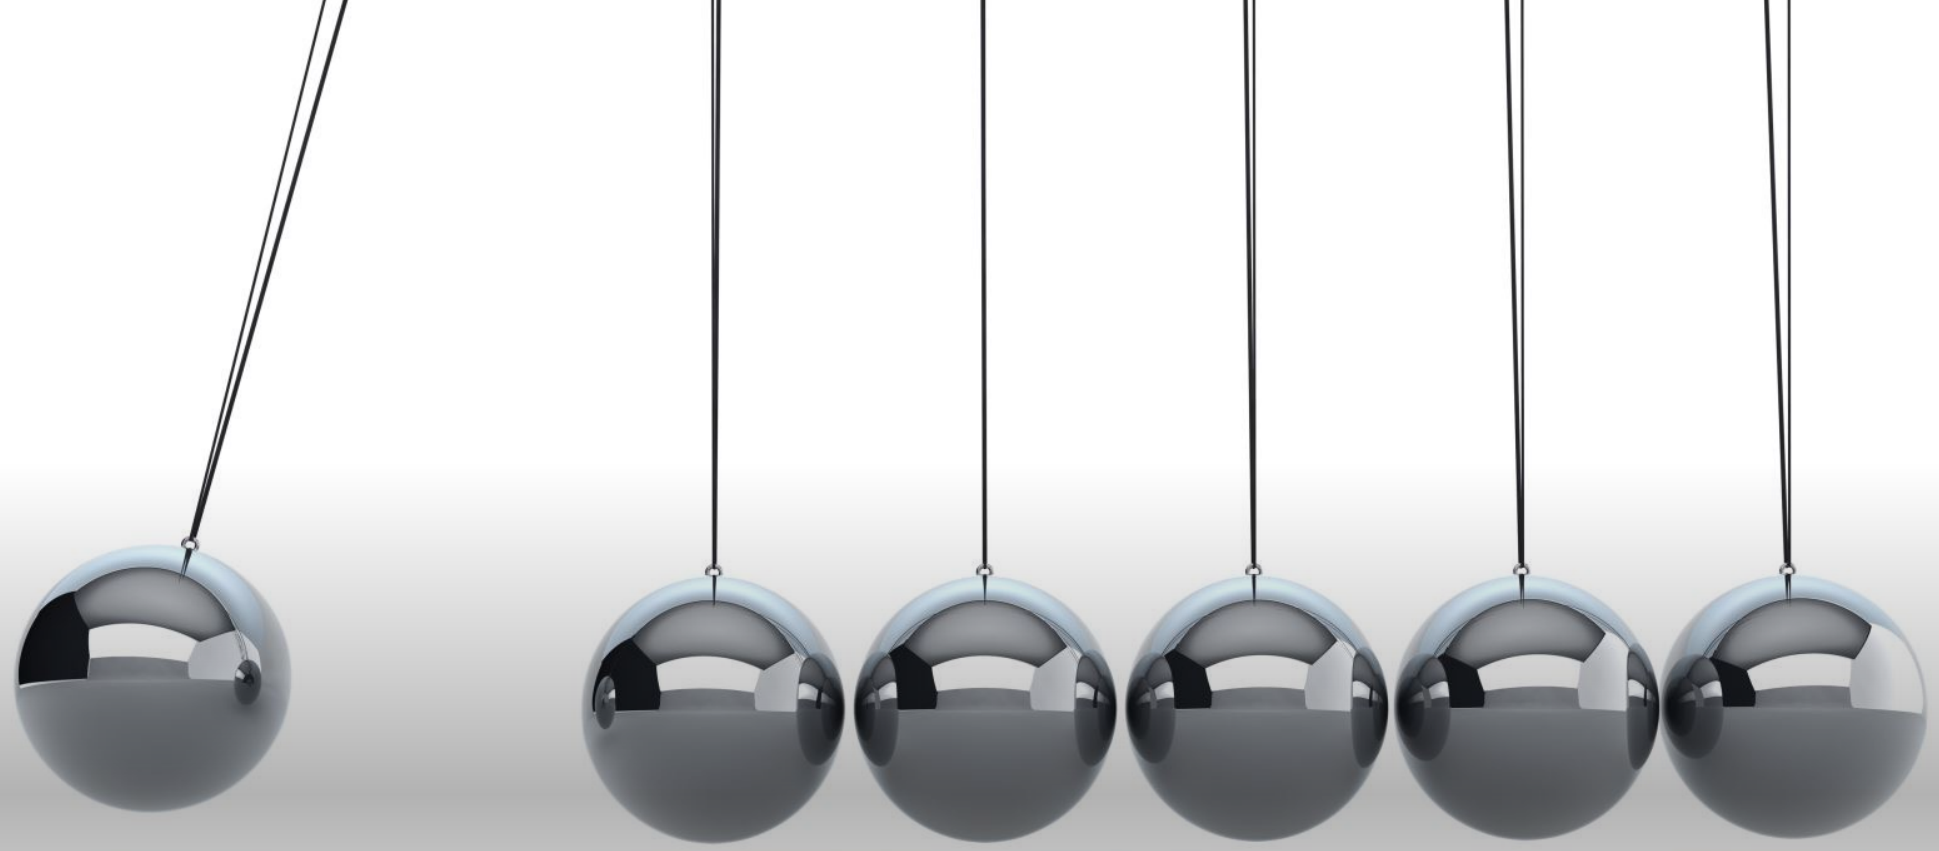

The difference between language models & PP?

We have a model

Invariance to transformation: contrastive learning

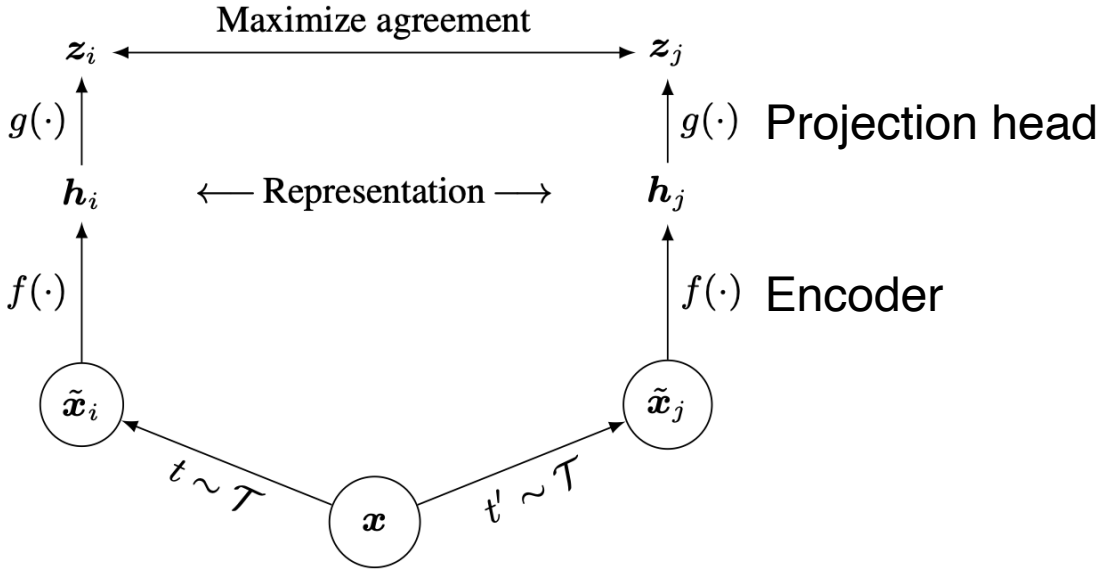

[JetCLR [2108.04253] (based on SimCLR Hinton et al.)]

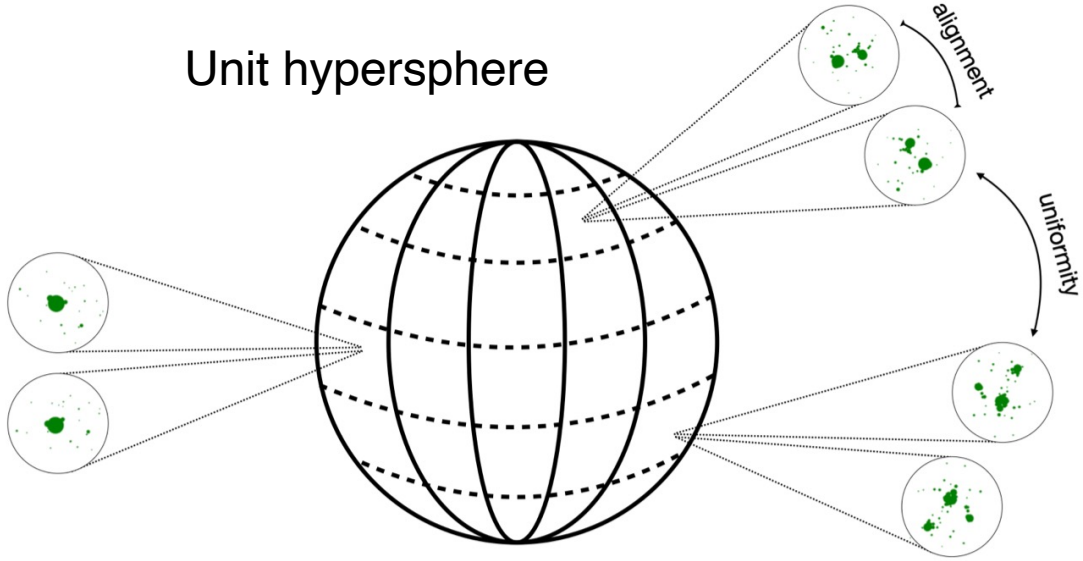

$$s(z_i, z_j) = \frac{z_i \cdot z_j}{|z_i||z_j|} = \cos\theta_{ij}$$

Augmentation	$\epsilon^{-1}(\epsilon_s = 0.5)$	AUC
none	15	0.905
translations	19	0.916
rotations	21	0.930
soft+collinear	89	0.970
all combined (default)	181	0.979

Encode physics into a GNN

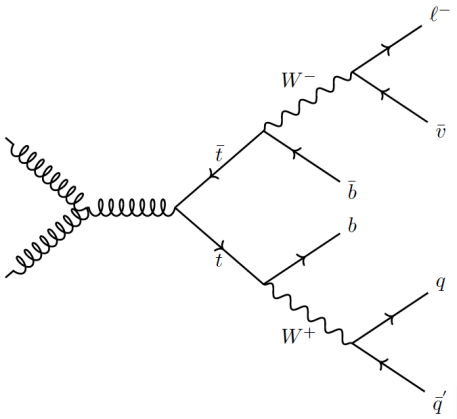

(a) Feynman diagram

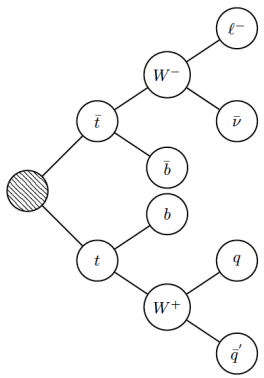

(b) Node and edge graph

Encode information by leaving out edges

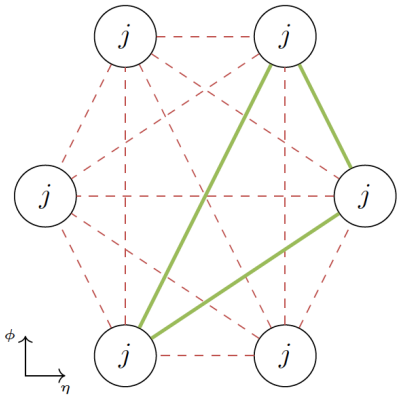

(a) Fully connected graph

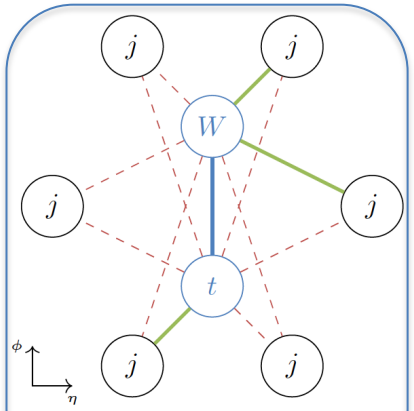

(b) Topograph

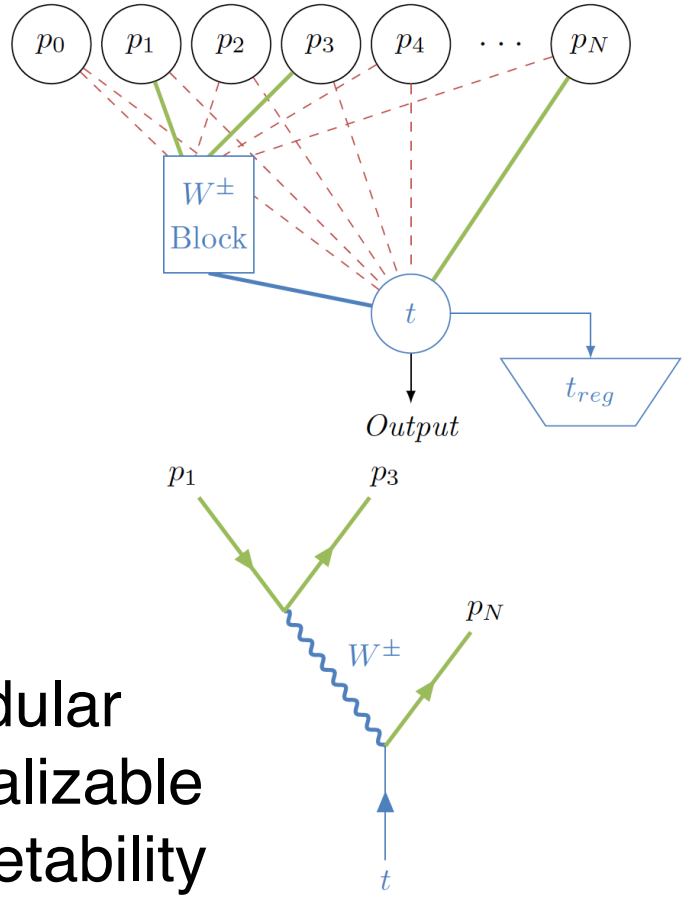

Modular
Generalizable
Interpretability
Combinatorics solving
Downstream tasks

[2303.13937]

Inject physics knowledge into AI

[1702.00748, 1711.02633]

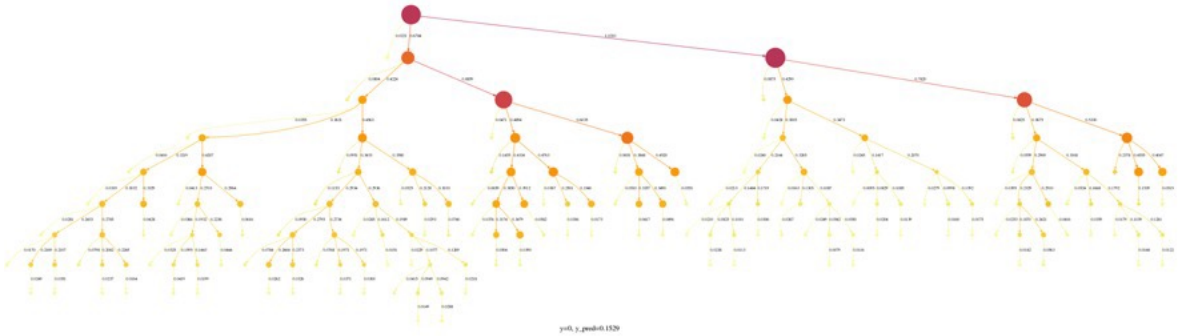

Tree structure of sequential recombination jet algorithms as Recursive NN

- Symmetries [rotation, translation, permutation,…]
  - Lorentz layers [2006.04780, 2201.08187]
  - GNNs: permutation symmetry [Energy flow network, ParticleNet]
  - PELICAN [2211.00454]
- Auxiliary tasks: energy conservation,…
- Observable construction with ML [1902.07180]

Grain of salt: inductive bias at which level?

θ → $p(z \mid \theta)$ → $p(x \mid \theta)$

Exact symmetries in latent space… but much harder to learn symmetries here

$p(x \mid z)$

Complex noise model (physics & detector folding)

Only approximate symmetries in data space

[credit: Kyle Cranmer @ Hammers & Nails 2022]

All models are wrong, but some are are useful.

– GEORGE BOX

Useful in what sense?

What is scientific understanding?

[We want more than an AI oracle]

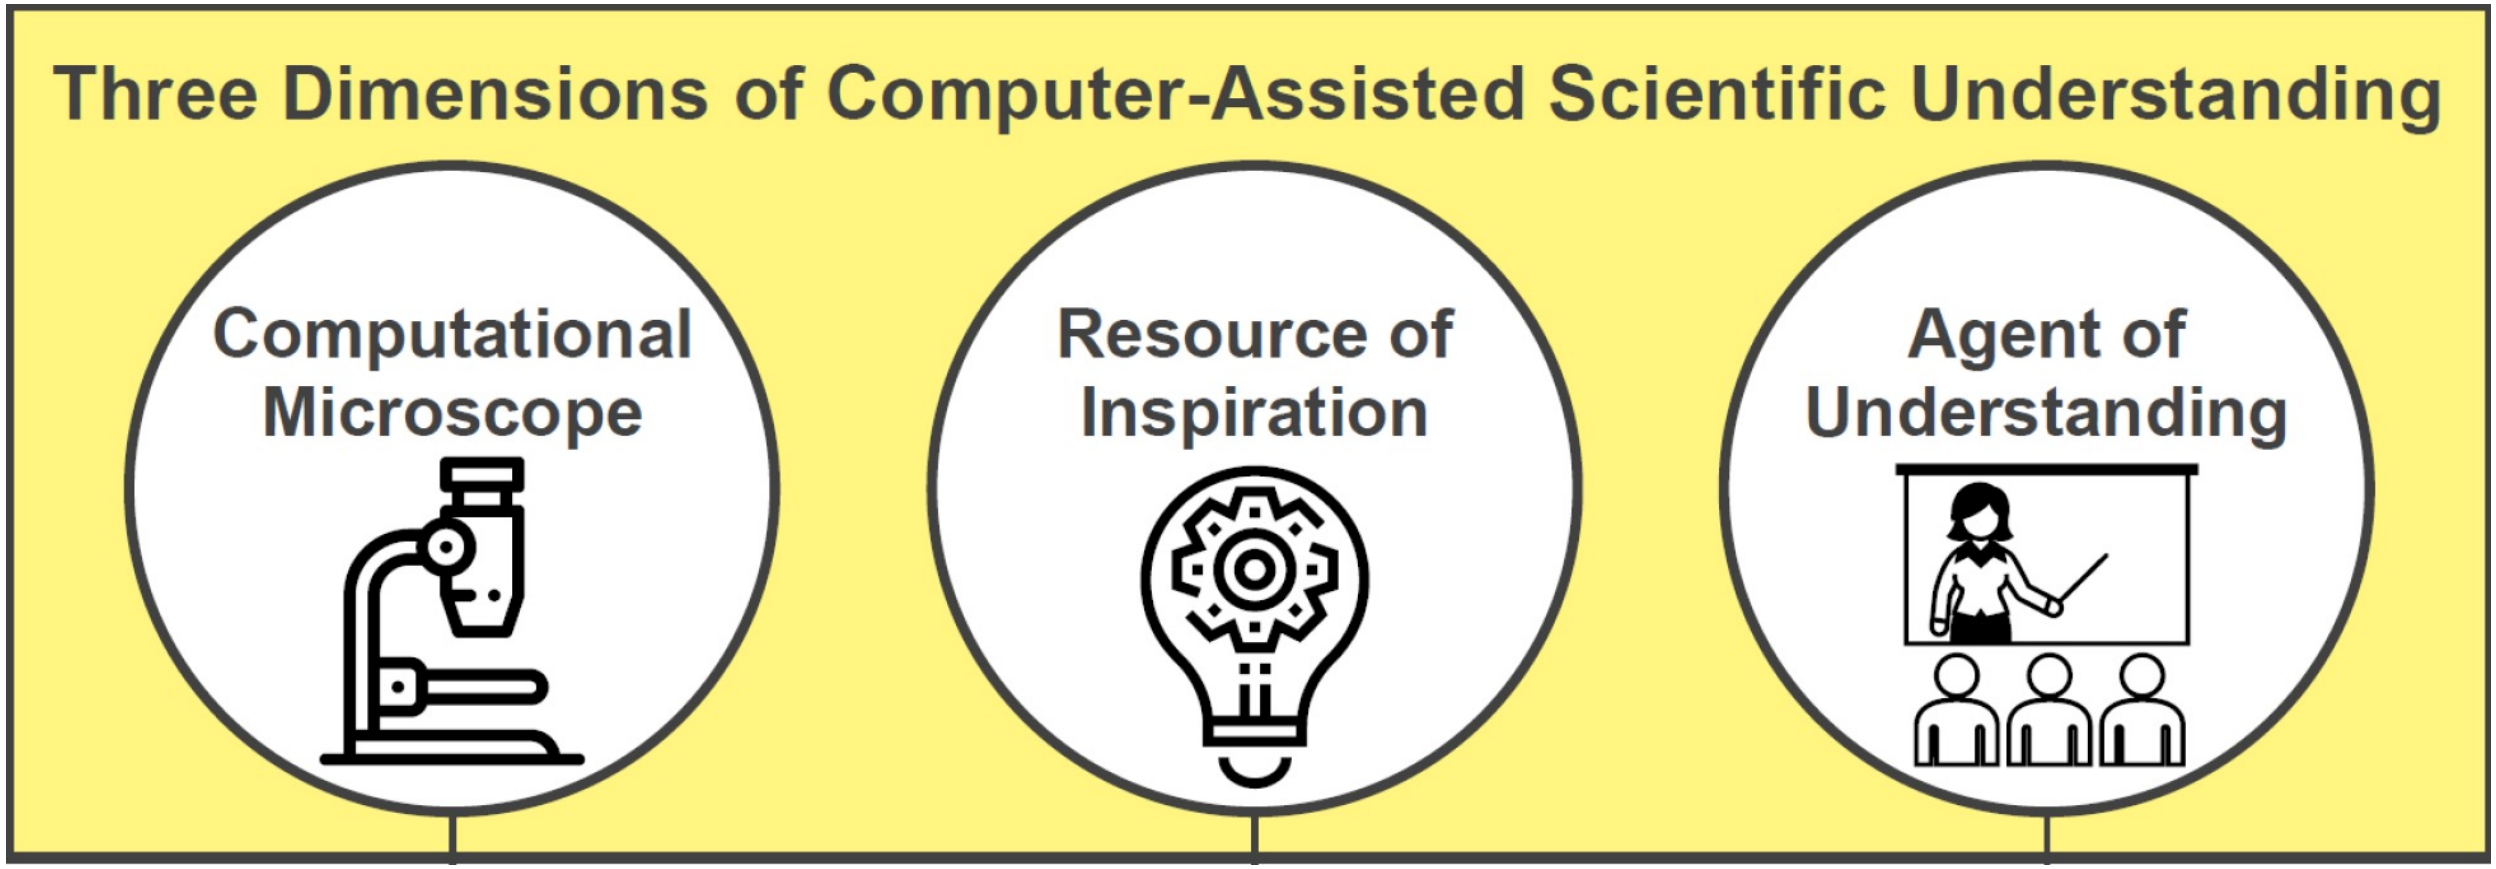

[2204.01467]

ML interpretability for science

Science

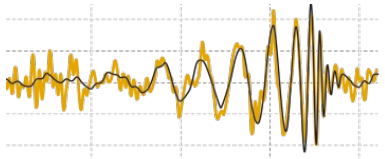

$$h = \frac{2G}{c^4}\frac{1}{r}\frac{\partial^2 Q}{\partial t^2}$$

Computer vision

???

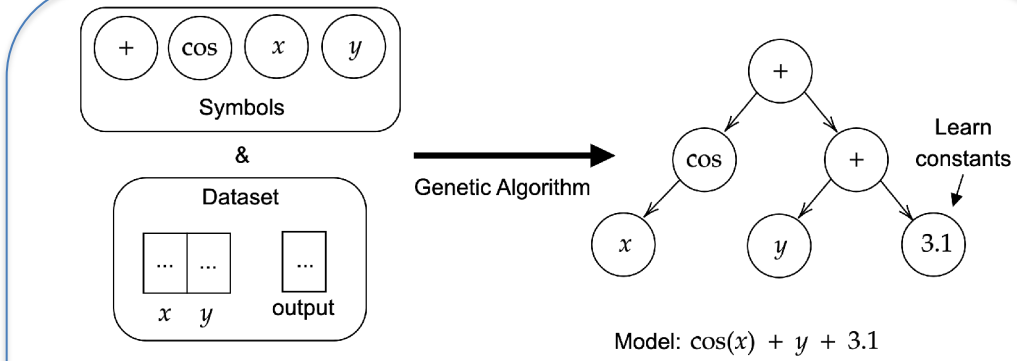

Analytic models often generalize better than NN
Symbolic regression as inductive bias

NN weights [black box] → PySR → **Analytic expression** [insights]

[Miles Cranmer – Hammers & Nails 2022]

Learn Newton's law from solar system

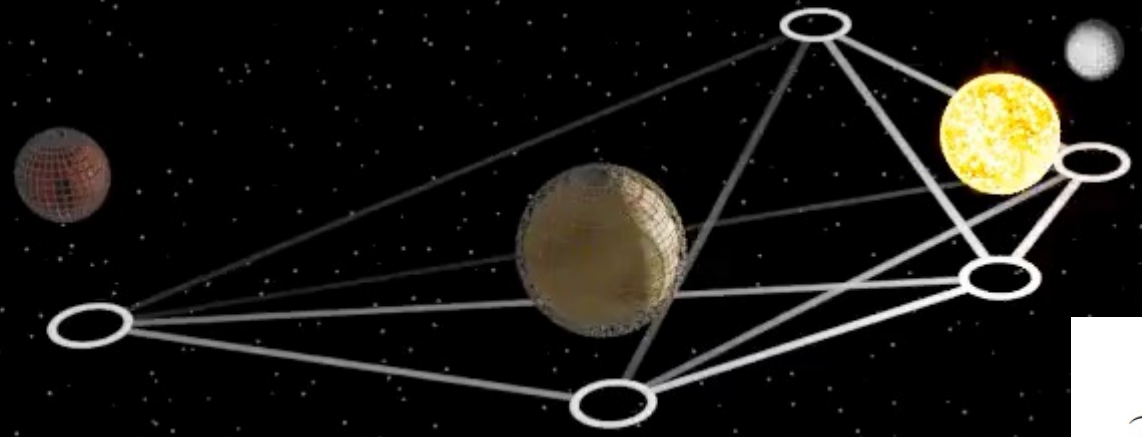

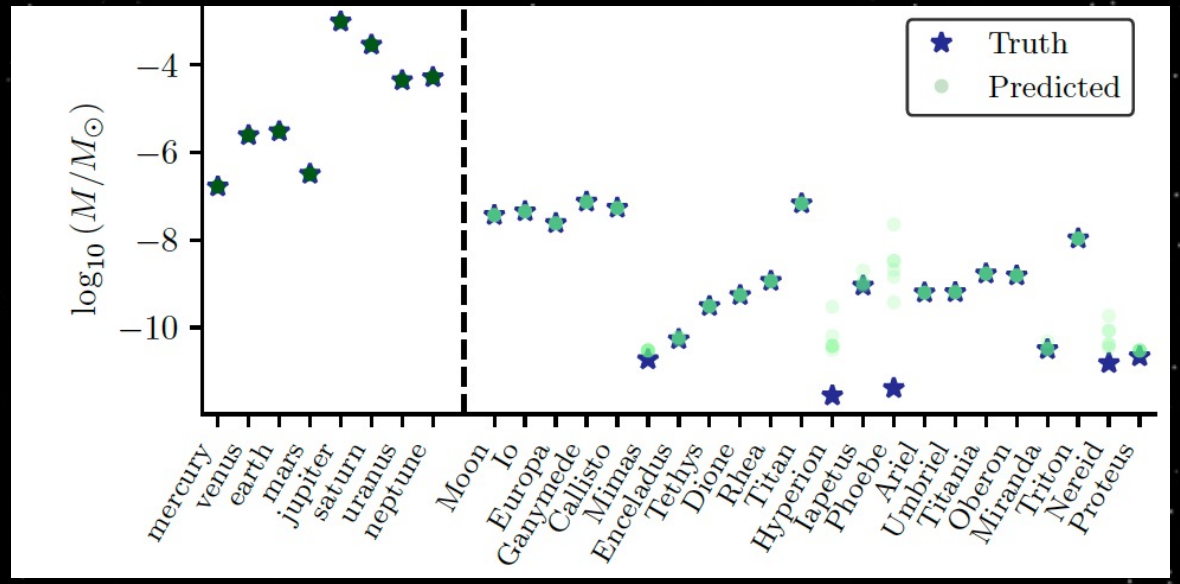

GNN → PySR → Learn masses + dynamics

2202.02306

True

Predicted

Surrogate modeling

ν-Flows: Conditional Neutrino Regression

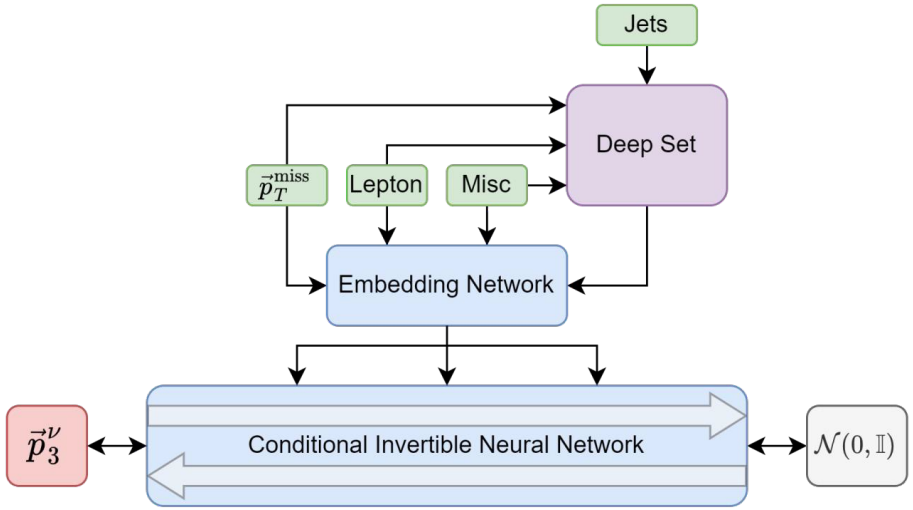

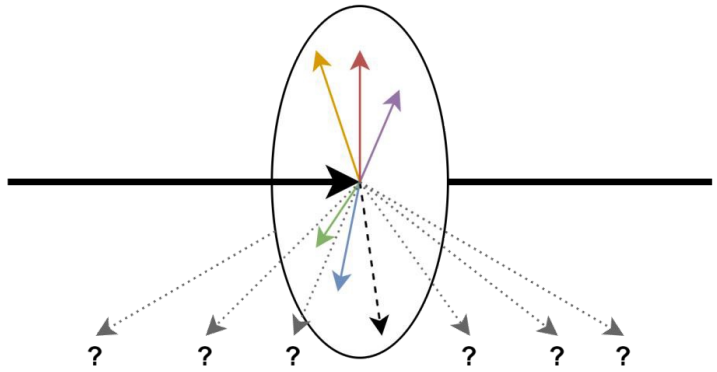

Learn ν **3**-momentum

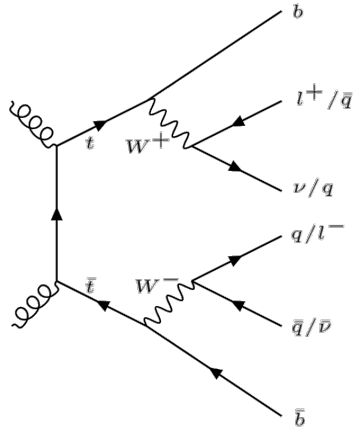

Cherry picked representative examples:

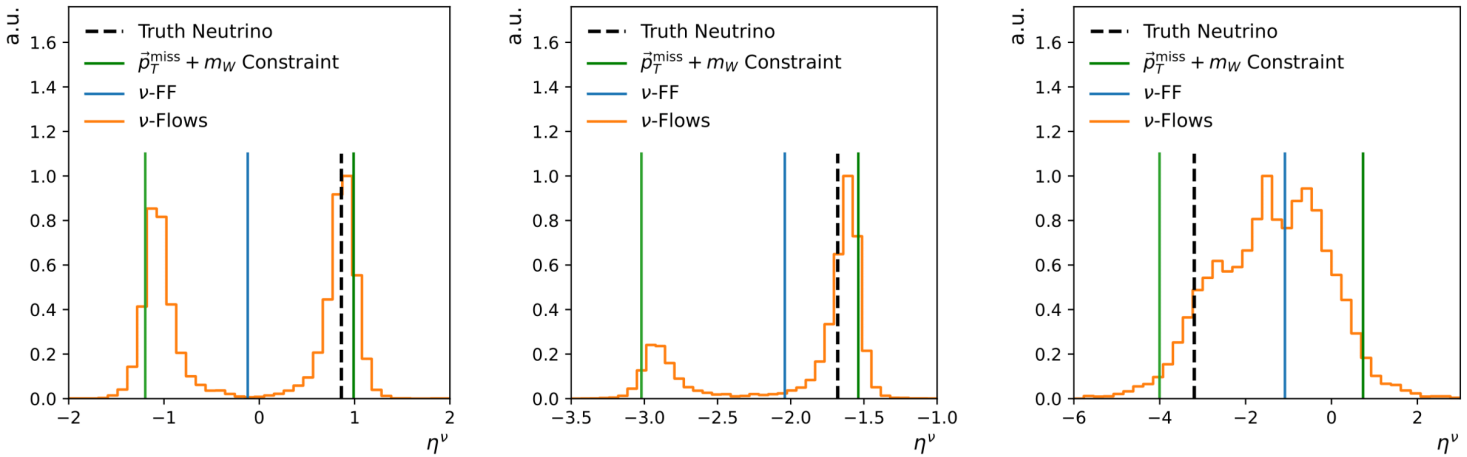

Conditional normalizing flow: learn **conditional likelihood** over neutrino momenta assuming an underlying process (inductive bias)

Improve over traditional method

[2207.00664]

Domain shift: calibrate synthetic to real data

1. Reweighting
- Tricks to battle curse of dimensionality
- Non-overlapping support

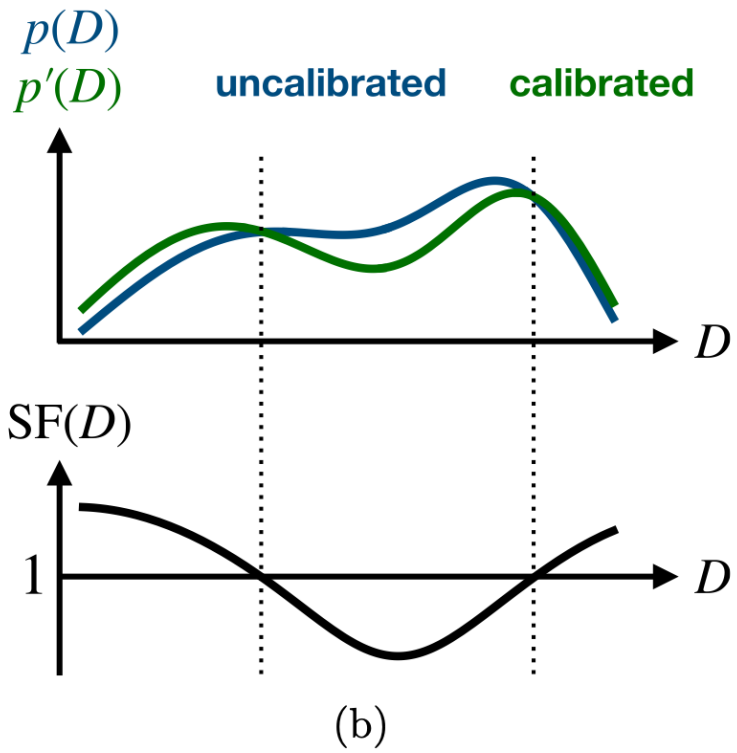

2. “Transport your problems away”

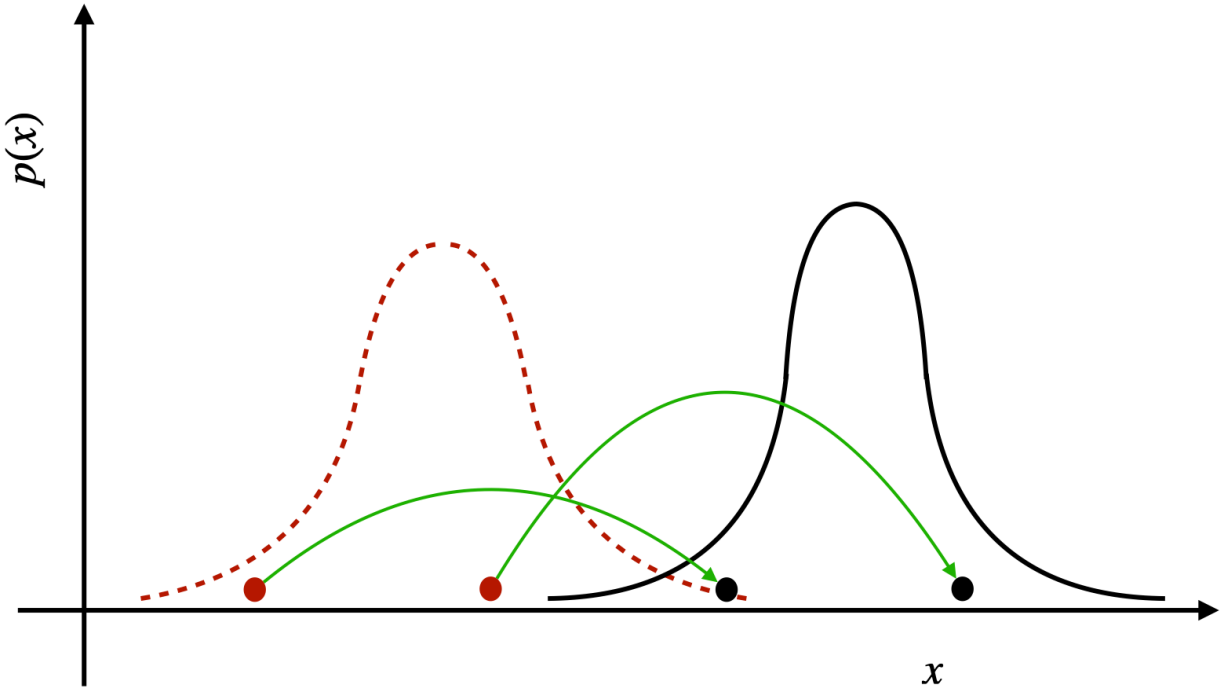

[2107.08648]

Full Sim & Reco

Bottleneck: computing budget

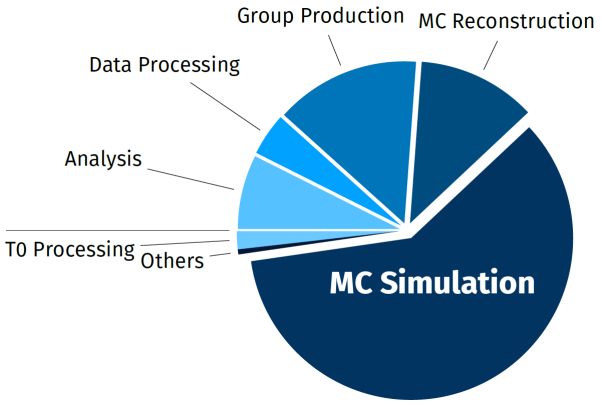

Up to 10 min / event

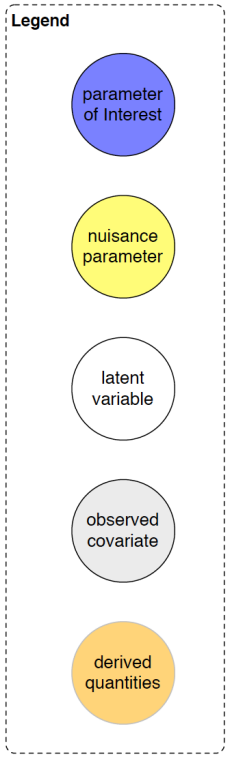

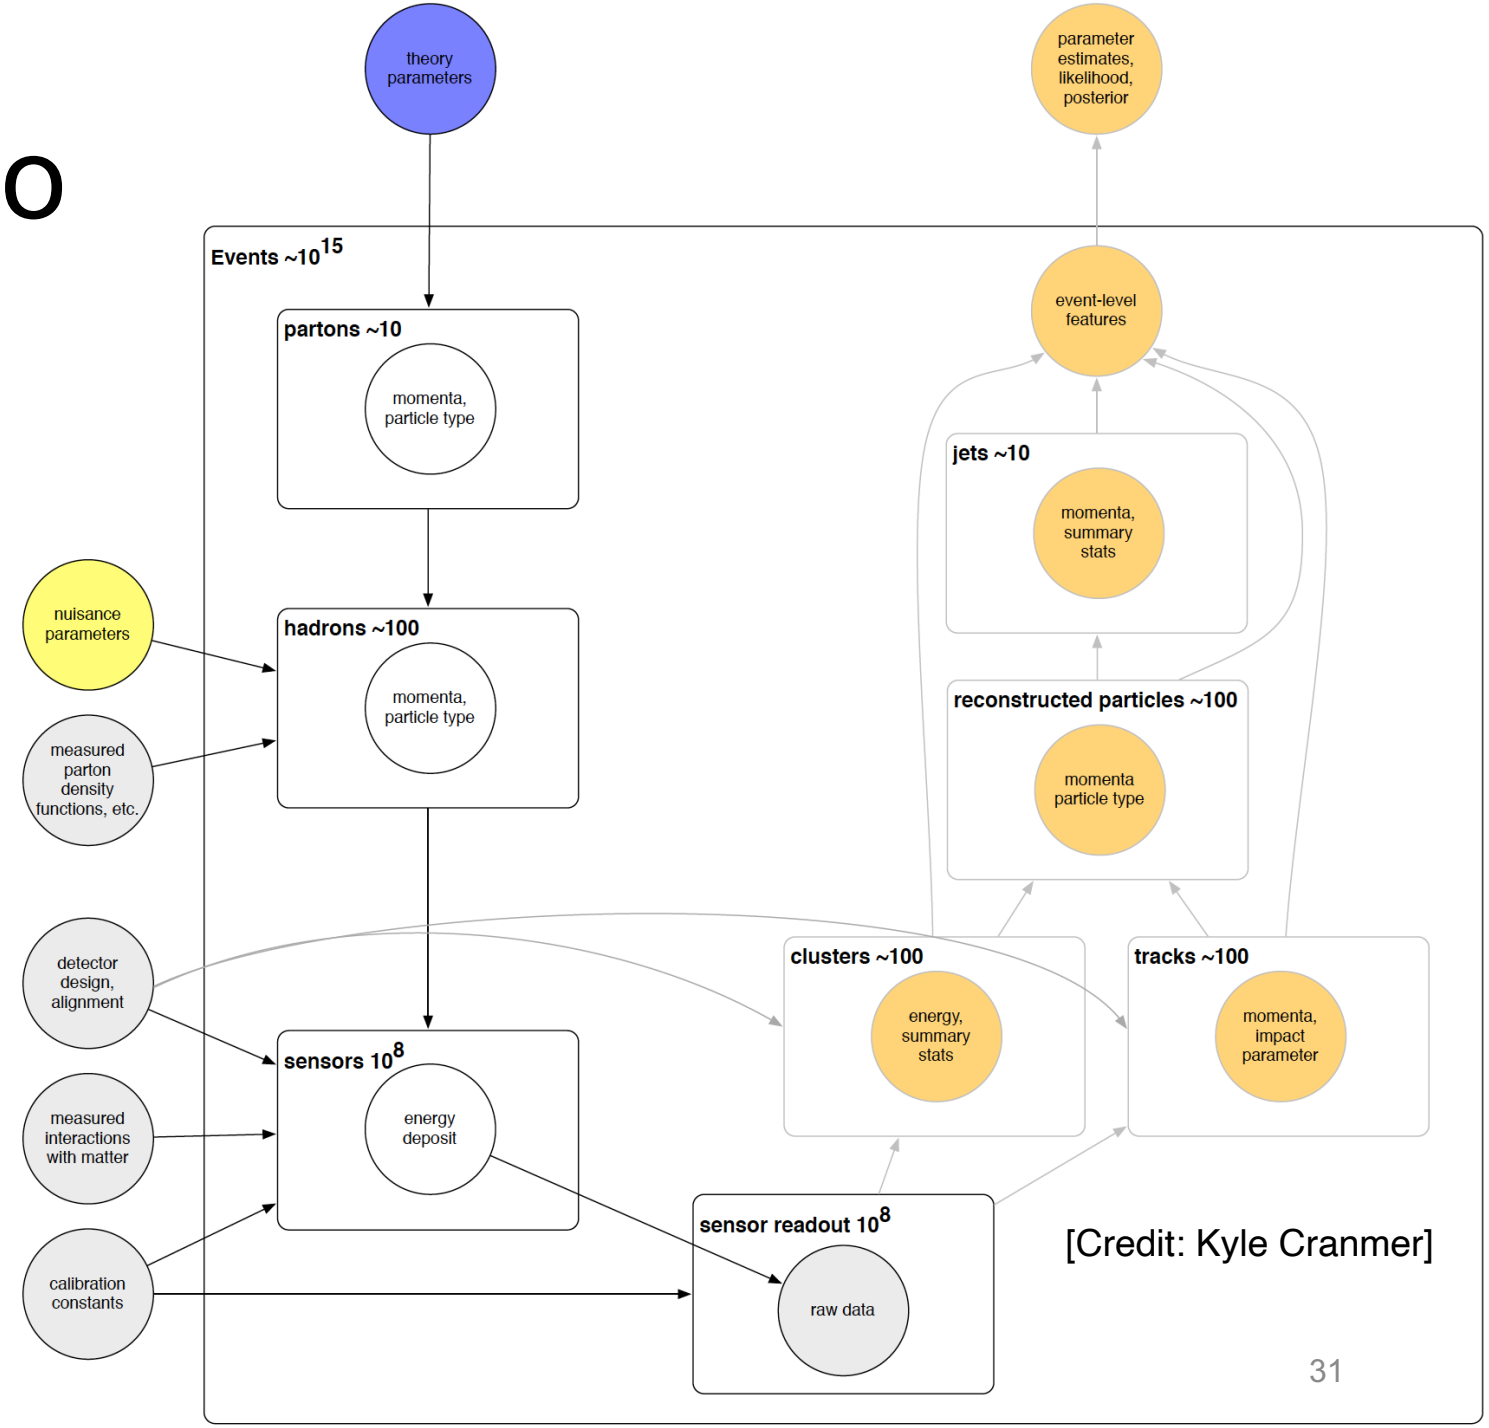

[Credit: Kyle Cranmer]

[LHCC-2022-005]

Fast Sim & Reco

Challenges:

- Fidelity, flexibility, portability
- Non-uniform geometry [FastCaloGAN, Geometry-aware]
- Sparse data
- Large dynamic range: tails
- Validation [2211.10295]
- Uncertainty
- Understanding inductive bias [GANplification]

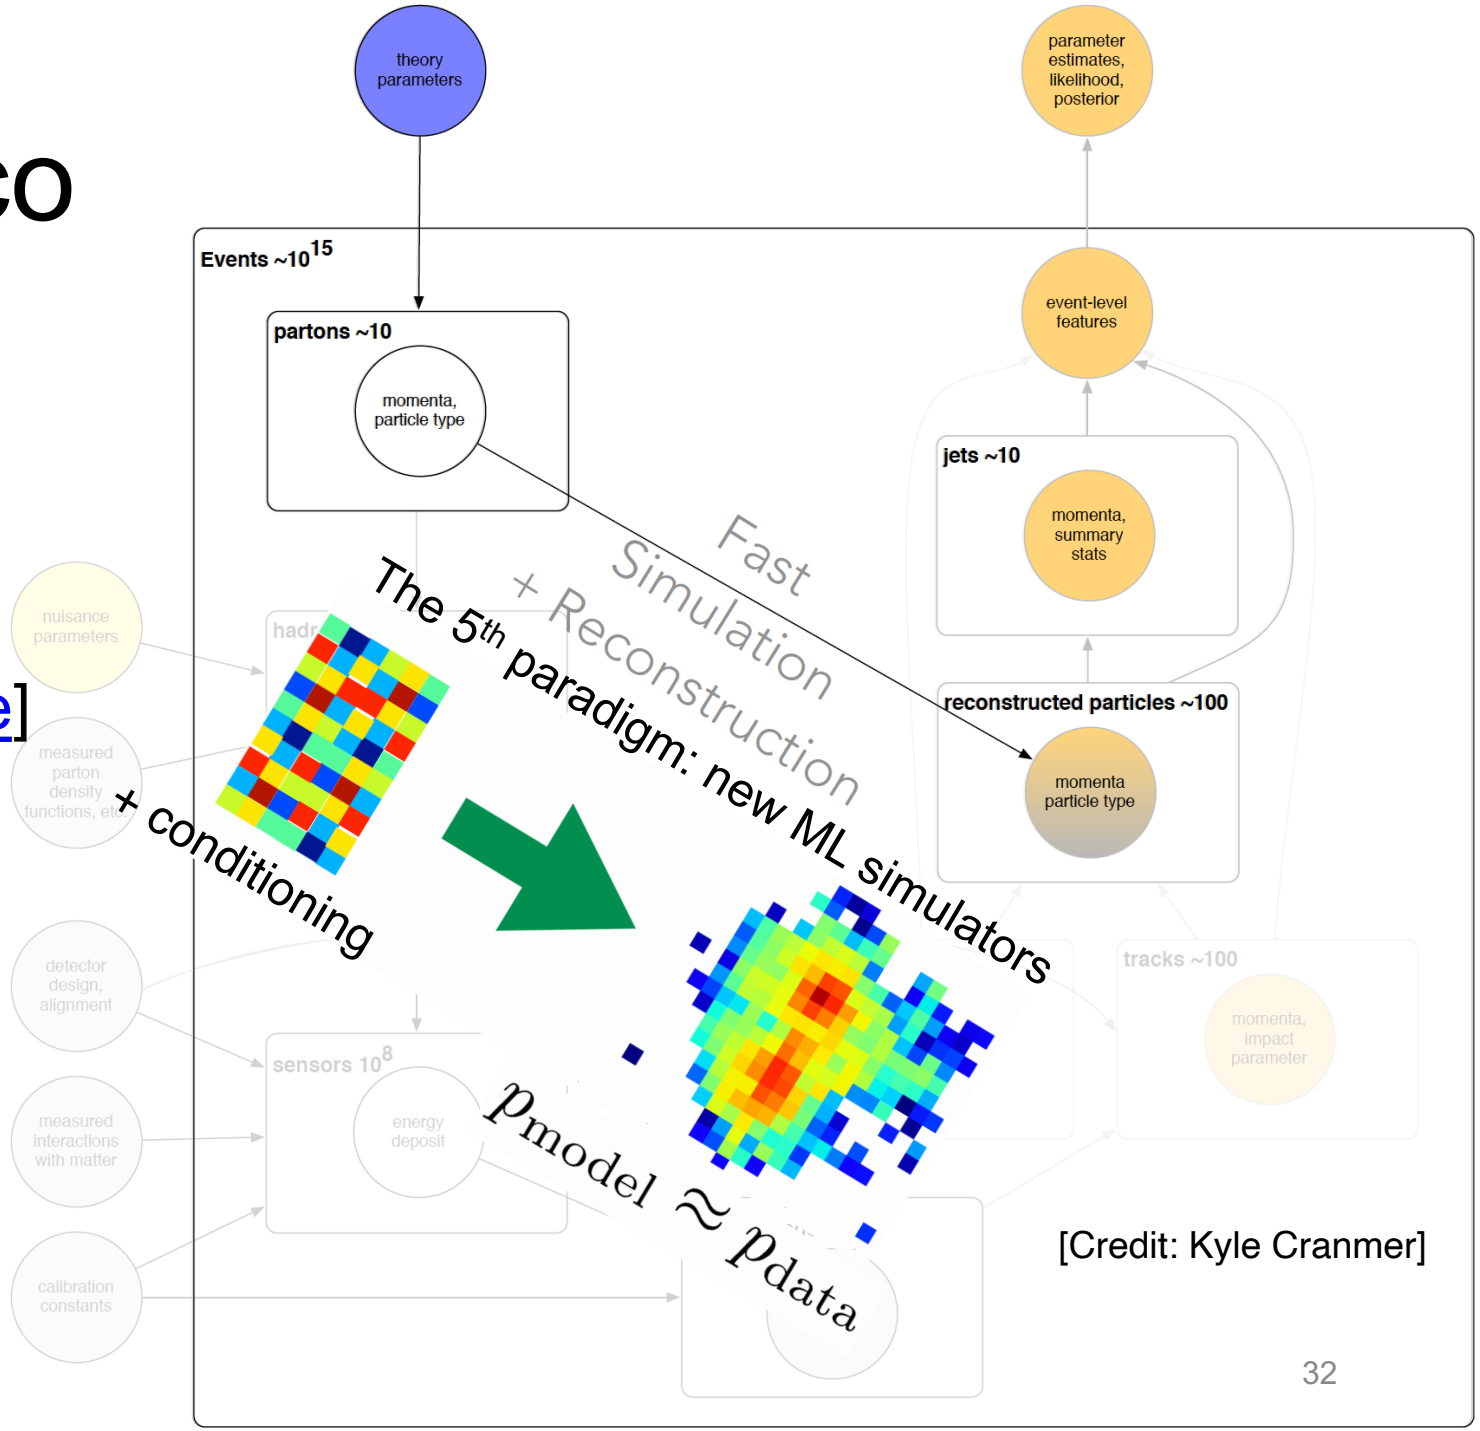

Toolbox: generative models

[Differentiable & fast]

Faces

[Karras et al., 2018]

Images of calo showers

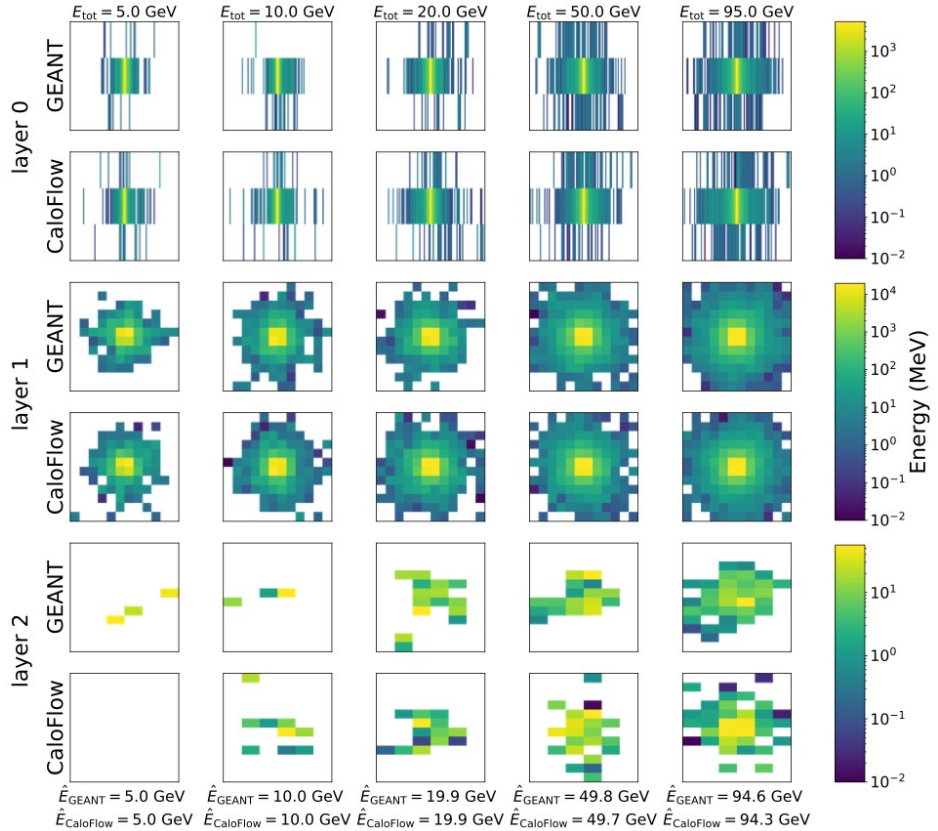

[CaloFlow]

VAEs, GANs, Flows, Diffusion,…

Generation from noise

Parton Interactions
$\mathcal{O}(10)$

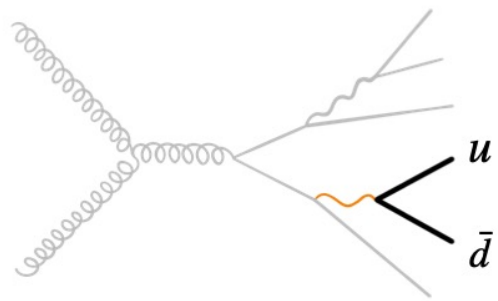

Showering
$\mathcal{O}(100)$

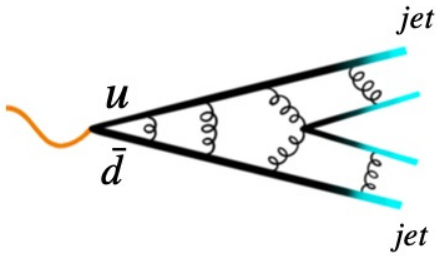

Detection
$\mathcal{O}(10^6)$

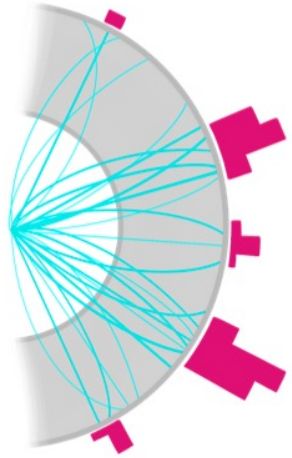

Reconstruction
$\mathcal{O}(10)$

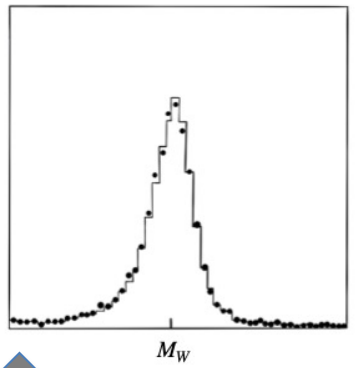

NOISE

NOISE

NOISE

→ Conditioning

[1907.03764,…]

[1701.05927,1712.10321,2005.05334, EPIC-GAN,…]

Create 4-vectors at analysis level
[1901.00875,1901.05282,…]

Images → Point cloud

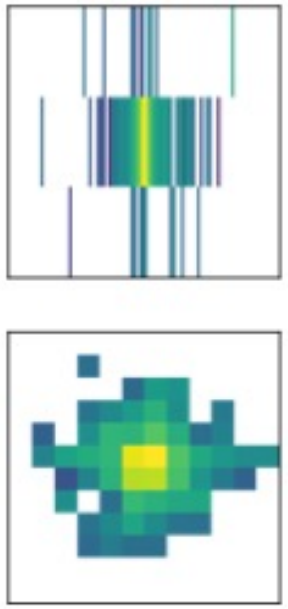

Decouple modeling
from detector geometry

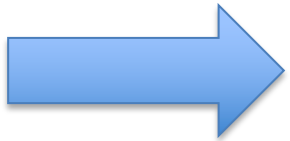

- Addresses sparsity issue
- Promotes portable solutions
- Encode symmetries (inductive bias)

New on the market: point cloud diffusion

[blog]

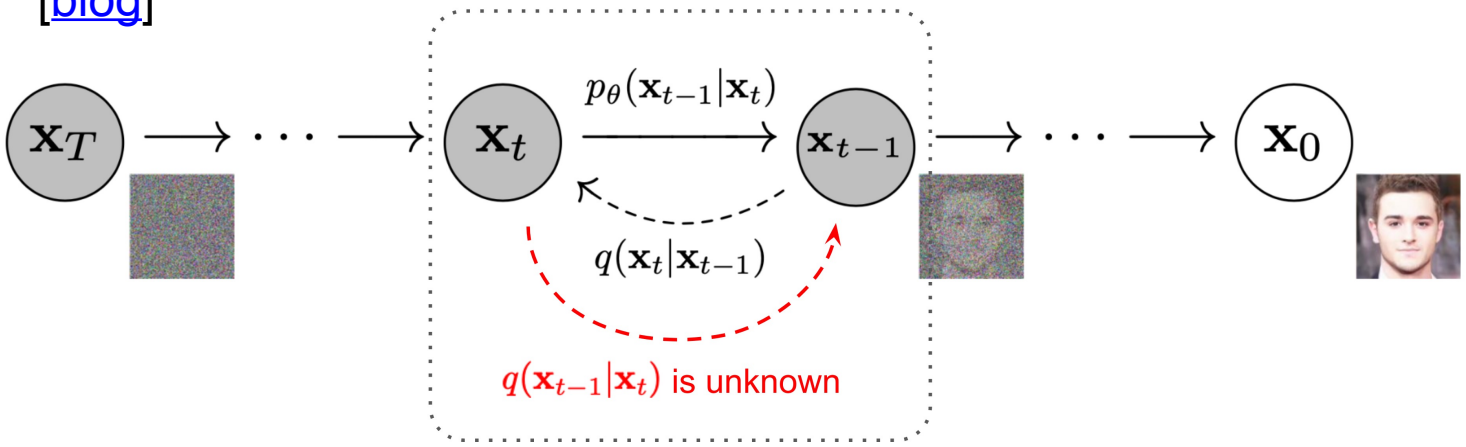

Gradually add Gaussian noise (right-to-left=forward)

Reverse "learn the noise"

1000 → 100 → ~20 steps (over last ~year)

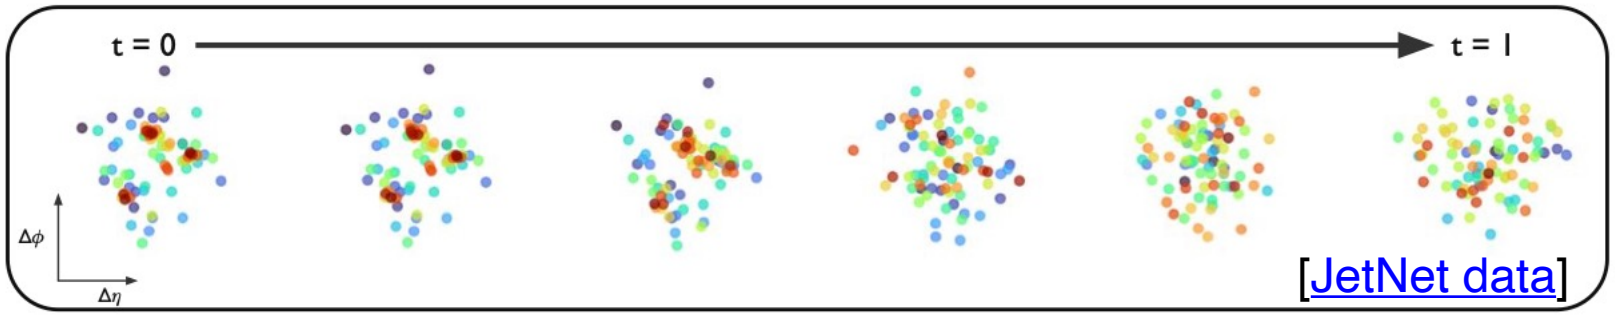

[JetNet data]

[PC-JeDi]

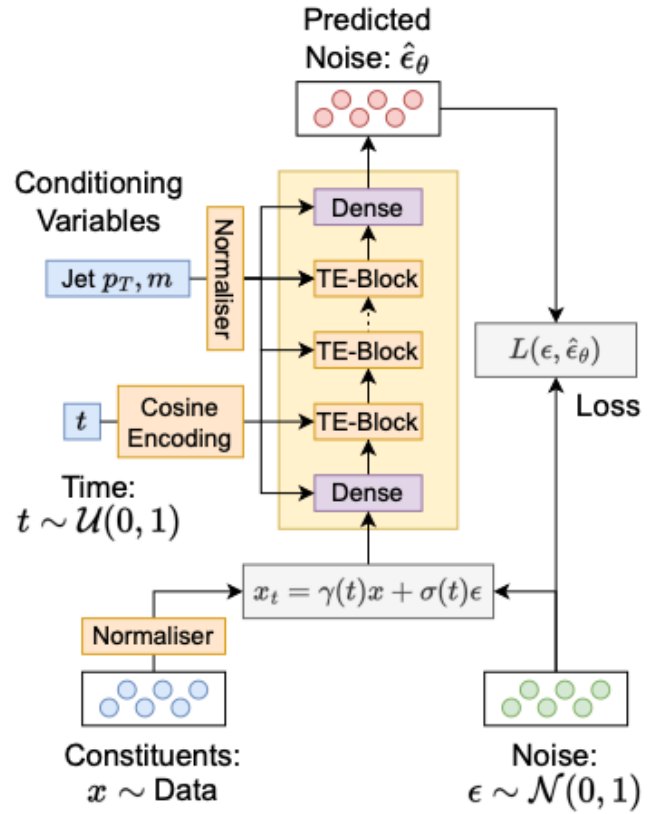

Transformer Encoder (TE) Block

[See also 2206.11898]

Invertible surrogates to solve inverse problem

Folding

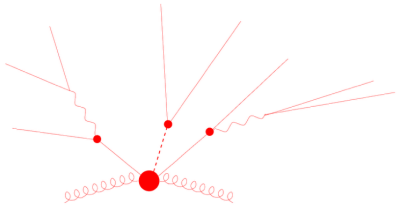

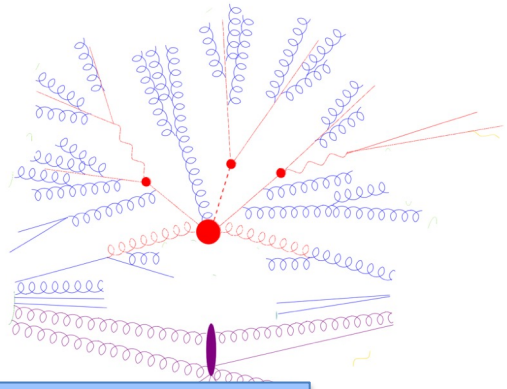

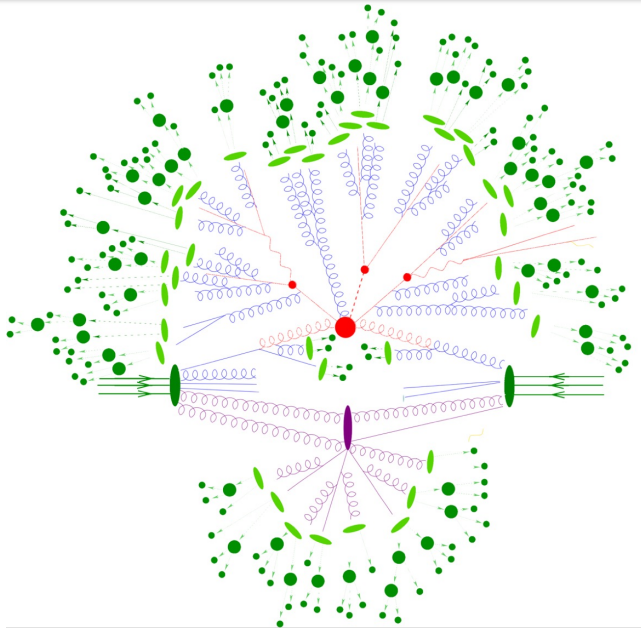

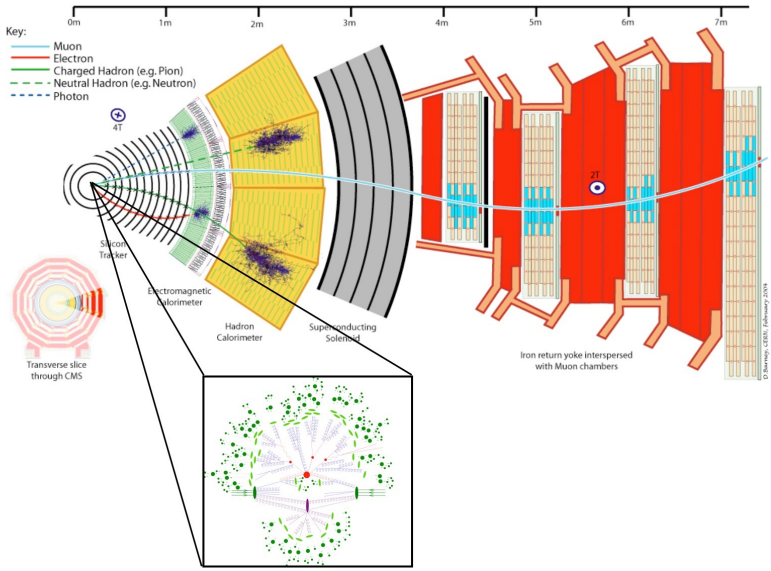

Unfolding allows to
- Compare at theory level
- Compare between experiments
- More *useful* data

Hard scatter

Radiation

Hadronization

Detector

Unfolding

Simulation-based inference: learn p(θ|x)

accounting for latent variables [parton shower, detector effects,…]

Replace **computationally expensive numerical integrals** (MEM, NNLO event weights etc.) with a **regression** phase (ML)

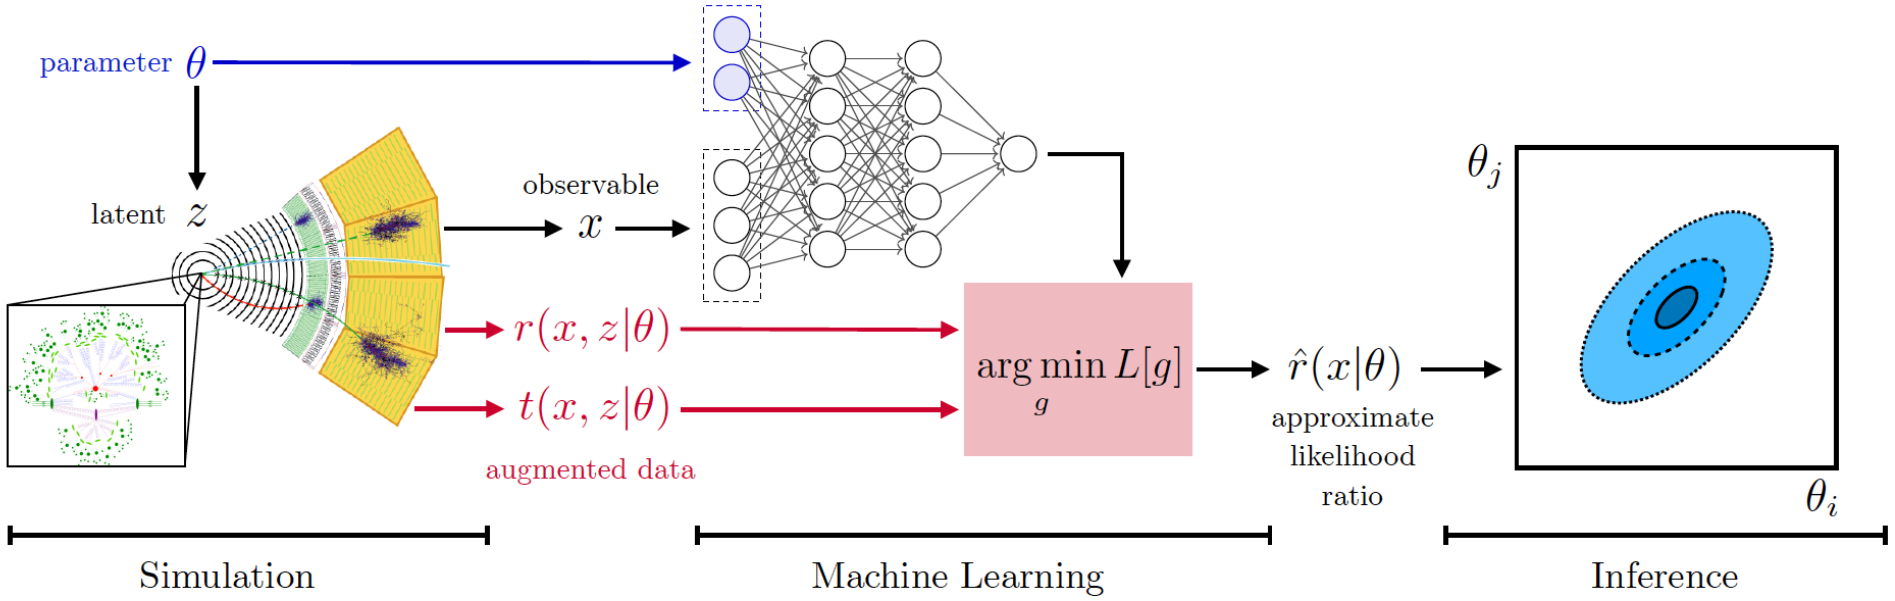

[1805.00013]

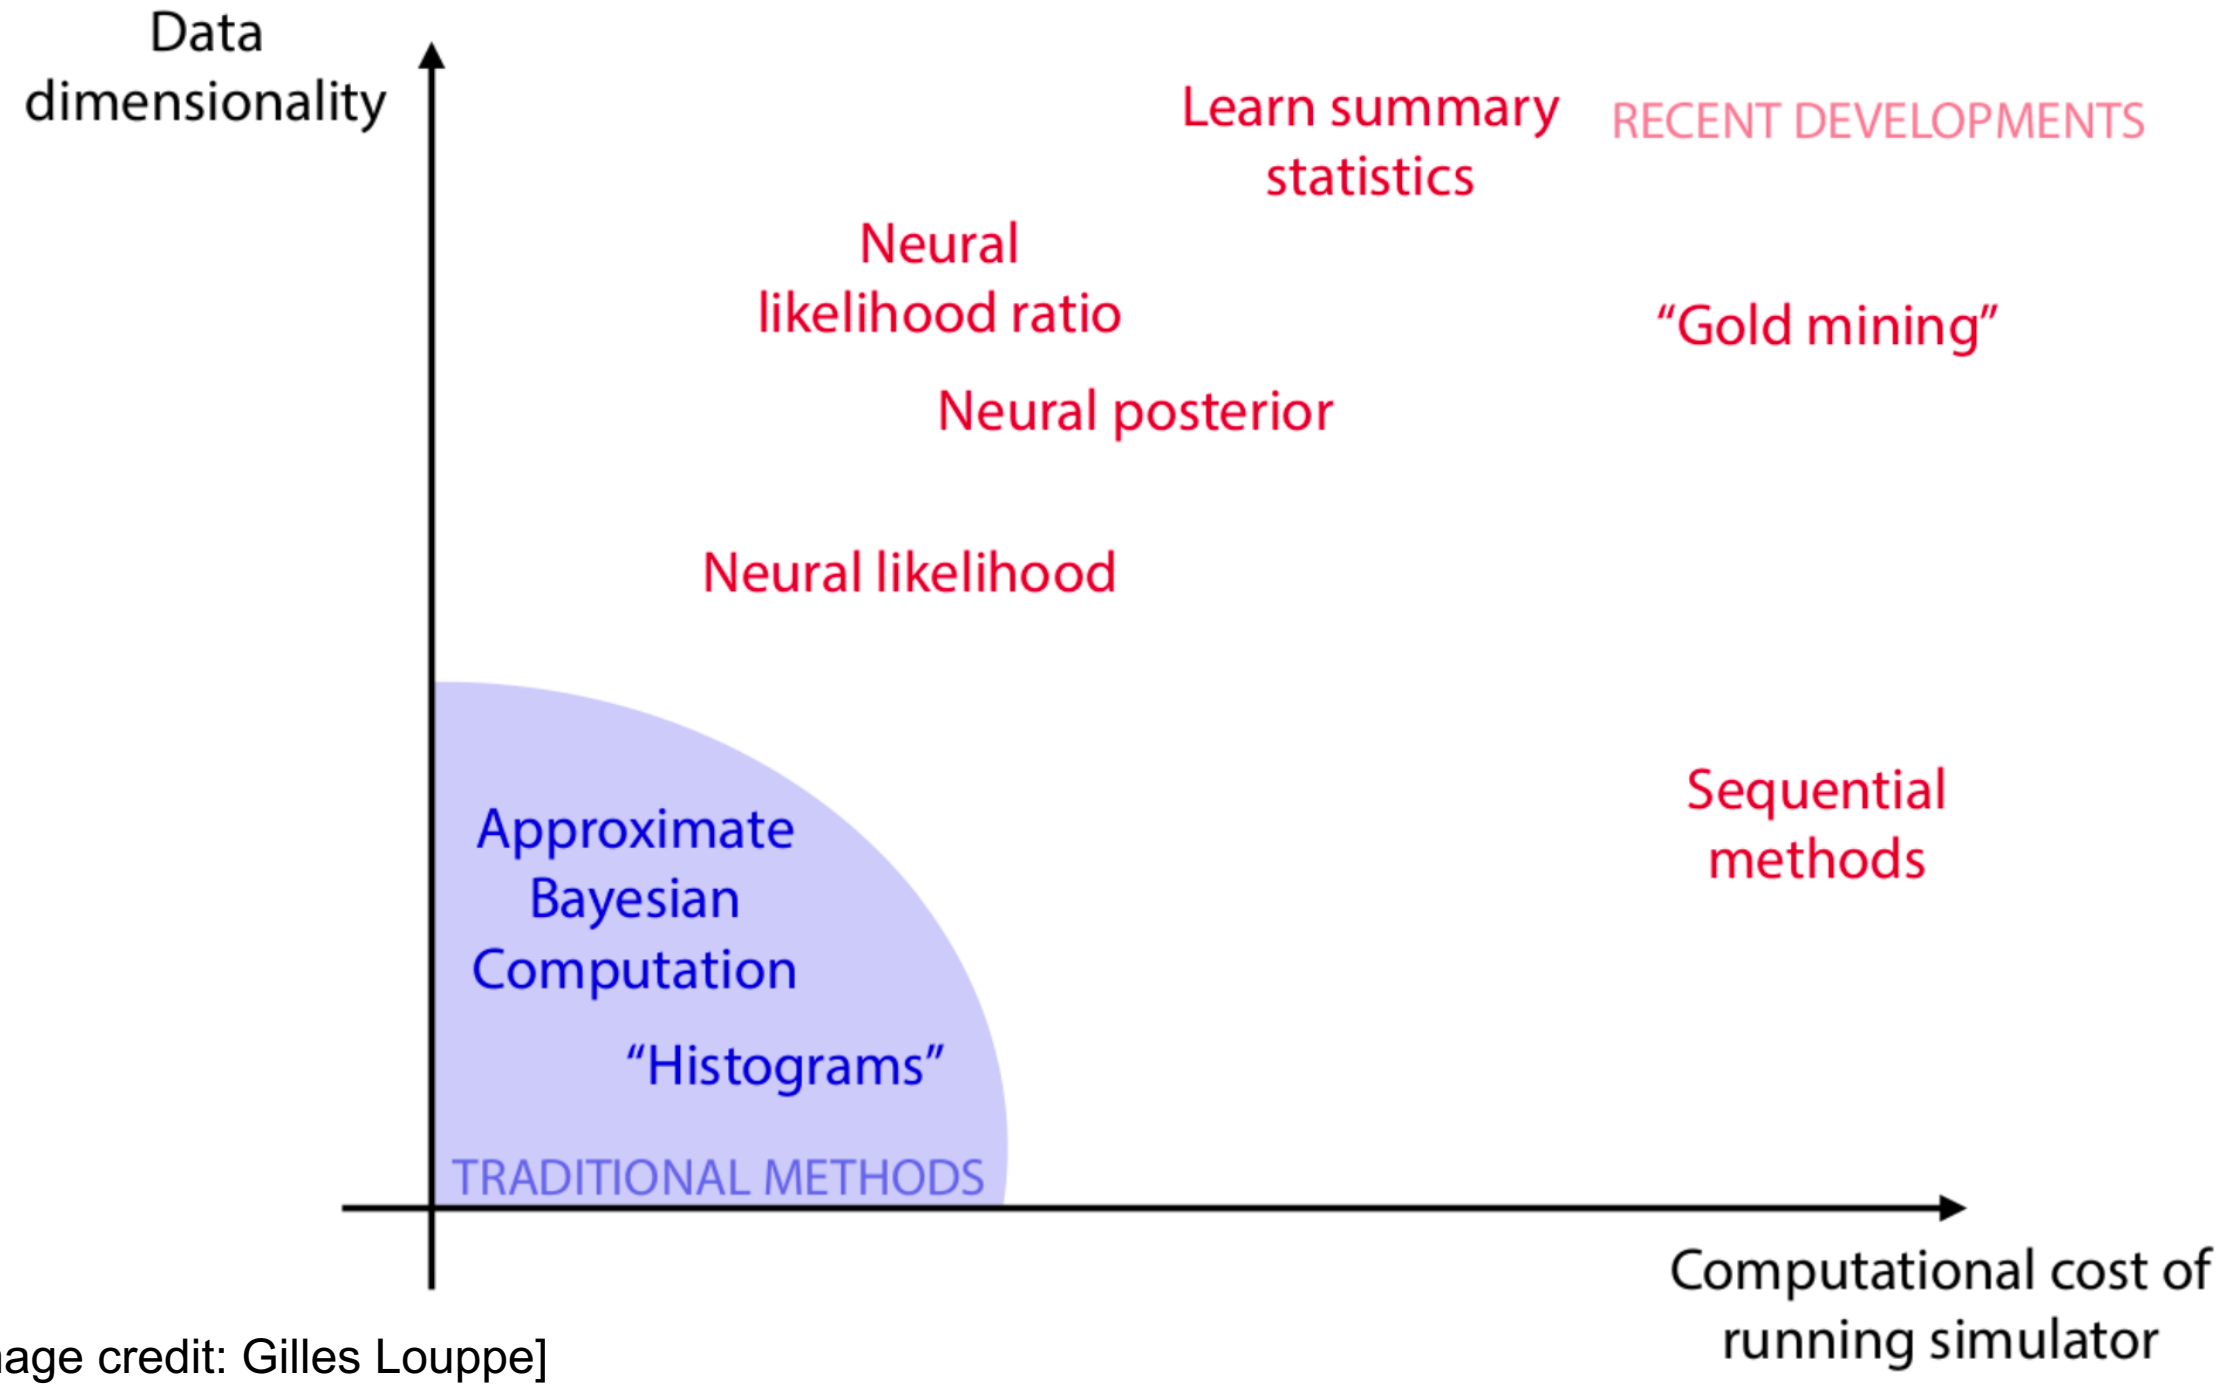

[Image credit: Gilles Louppe]

$$p(\text{theory}\,|\,\text{data}) = \frac{p(\text{data}\,|\,\text{theory})p(\text{theory})}{p(\text{data})}$$

[Lukas Heinrich - Detector design using differential programing]

Ultimate goal:
Learning about Nature

Optimizing the science output

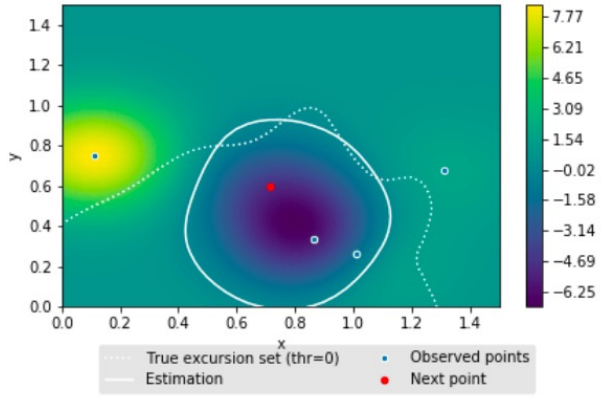

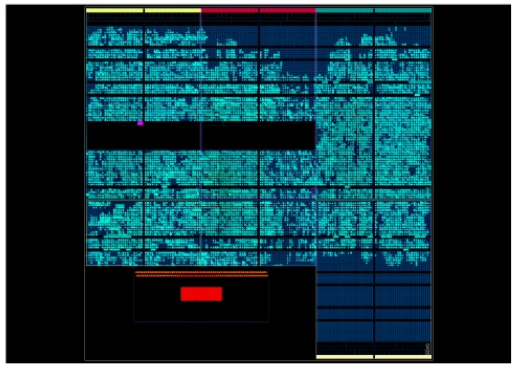

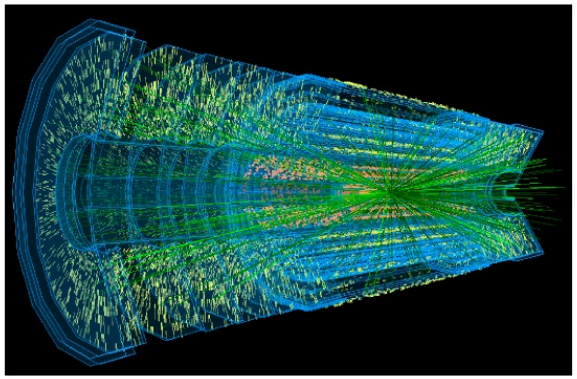

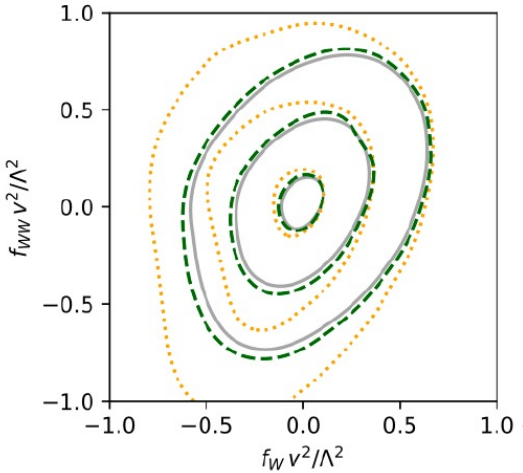

Optimal Theory Exploration

Optimal Data Taking / Experiment Operations

Optimal Reconstruction

Optimal Analysis

[Lukas Heinrich - Detector design using differential programing]

Natural limit: true posterior *p*(theory | data)

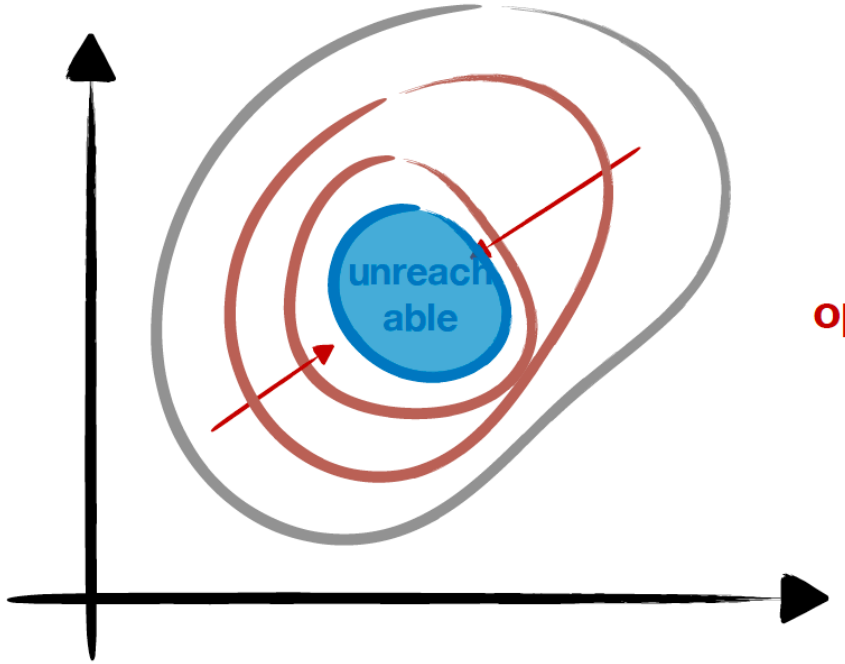

unoptimized

optimized (e.g. w/ ML)

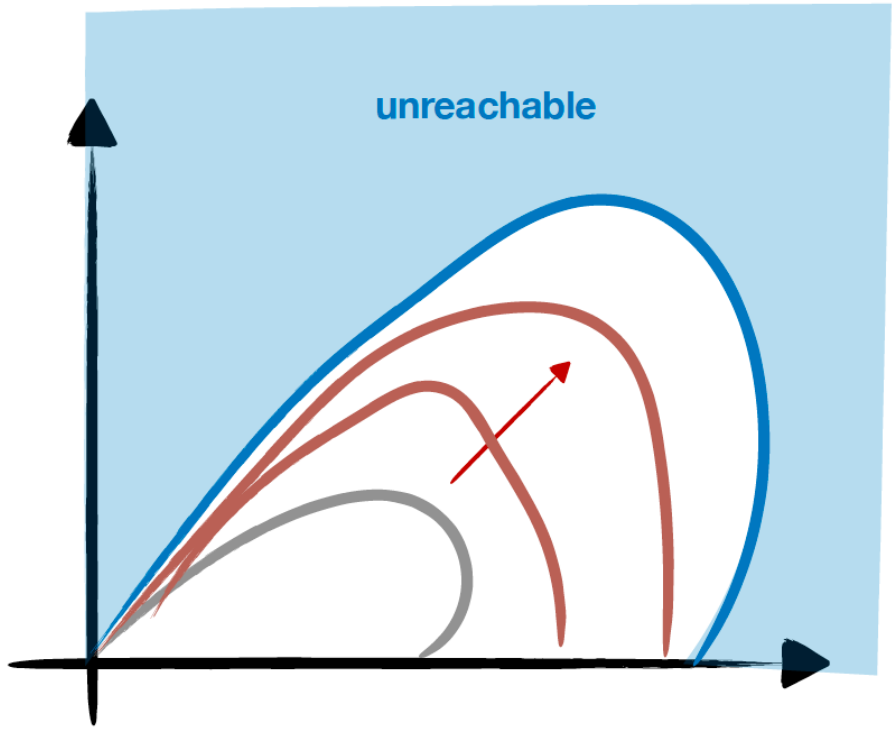

Measurements
(e.g. Higgs Couplings)

Searches
(e.g. Supersymmetry)

[Lukas Heinrich - Detector design using differential programing]

Opportunity: new ***optimal*** detector

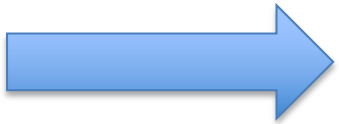

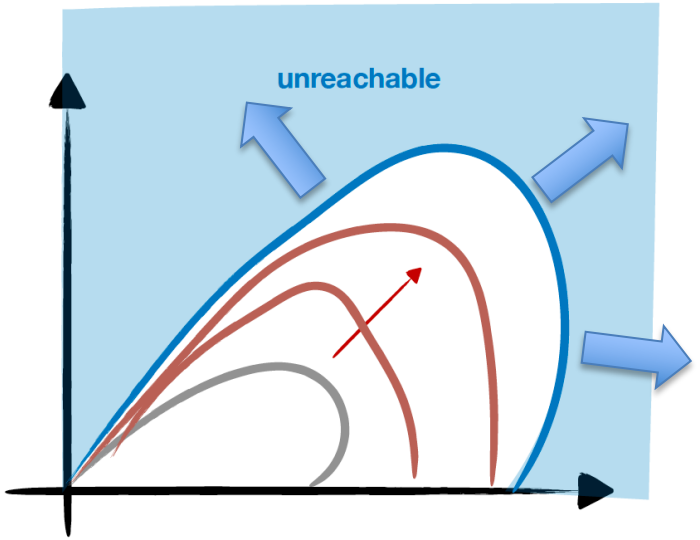

Goal: optimize
p(theory | data)

[Lukas Heinrich - Detector design using differential programing]

Need **design-conditional** model $p(x \mid \theta, \mathbf{D})$

Approximate $p(x \mid \theta, D)$ using **generative model**
→ **Fast**
→ **Differentiable**

Challenge:
$p(x \mid \mathbf{D})$ without already exploring all design space **D**

Solution:
train local models as you optimize [2002.04632]

Detector design is a challenging frontier in ML@HEP
Fine-tune human design → discovery of novel designs

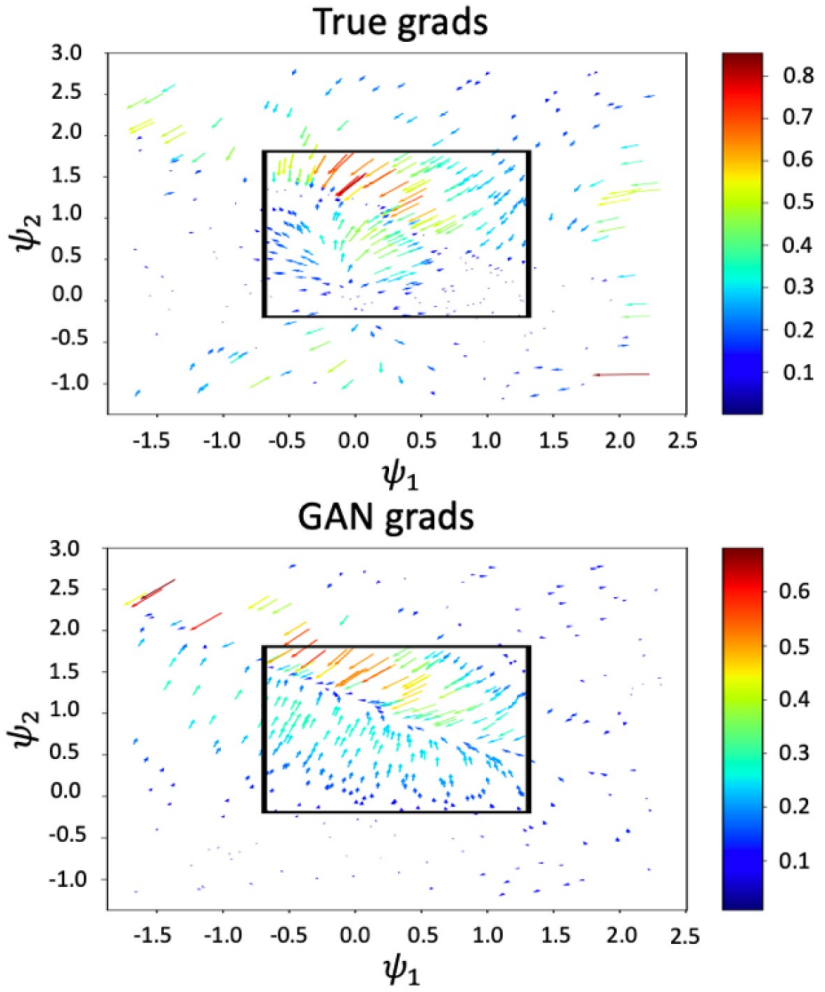

Search for the Unknown

Signal-driven approach

Works great if you know what you're looking for!

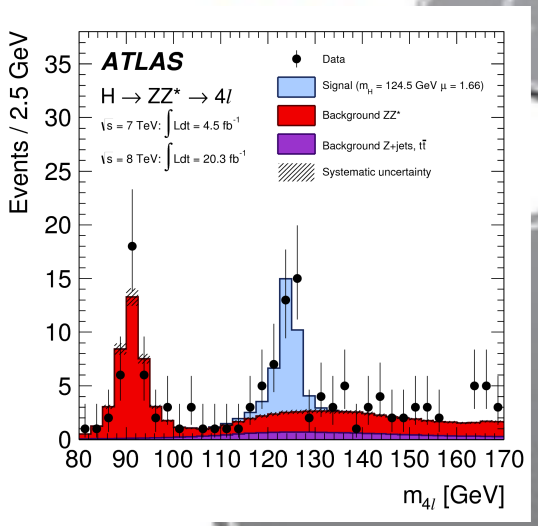

SUSY, etc.

Higgs

Top

W boson

Strategy breaks down as confidence in model decreases

Playing the lottery

How to maximize the discovery potential

Current approach is inefficient & incomplete

	e	μ	τ	q/g	b	t	γ	Z/W	H	BSM → $SM_1 \times SM_1$				BSM → $SM_1 \times SM_2$				BSM → complex			
										q/g	γ/π^0's	b	⋯	tZ/H	bH	⋯		$\tau qq'$	eqq'	$\mu qq'$	⋯
e	[37,38]	[39,40]	[39]	∅	∅	∅	[41]	[42]	∅	∅	∅	∅		∅	∅	∅		∅	[43,44]	∅	
μ		[37,38]	[39]	∅	∅	∅	[41]	[42]	∅	∅	∅	∅		∅	∅	∅		∅	∅	[43,44]	
τ			[45,46]	∅	[47]	∅	∅	∅	∅	∅	∅	∅		∅	∅	∅		[48,49]	∅	∅	
q/g				[29,30,50,51]	[52]	∅	[53,54]	[55]	∅	∅	∅	∅		∅	∅	∅		∅	∅	∅	
b					[29,52,56]	[57]	[54]	[58]	[59]	∅	∅	∅		[60]	∅	∅		∅	∅	∅	
t						[61]	∅	[62]	[63]	∅	∅	∅		[64]	[60]	∅		∅	∅	∅	
γ							[65,66]	[67–69]	[68,70]	∅	∅	∅		∅	∅	∅		∅	∅	∅	
Z/W								[71]	[71]	∅	∅	∅		∅	∅	∅		∅	∅	∅	
H									[72,73]	[74]	∅	∅		∅	∅	∅		∅	∅	∅	
BSM → $SM_1 \times SM_1$ q/g										∅	∅	∅		∅	∅	∅		∅	∅	∅	
γ/π^0's											[75]	∅		∅	∅	∅		∅	∅	∅	
b												[76,77]		∅	∅	∅		∅	∅	∅	
⋮																					
⋮																					

[1907.06659]

Vast signature space **un**explored

Rephrasing the problem:

Look for deviations from SM in model agnostic way

Cast a wide web

Inform future searches

Model-agnostic search portfolio

1. Unsupervised autoencoder-style outlier detection

2. Semi-supervised in-situ background modeling

Fabulous idea: outlier detection with autoencoders

Train on *normal* (=SM)

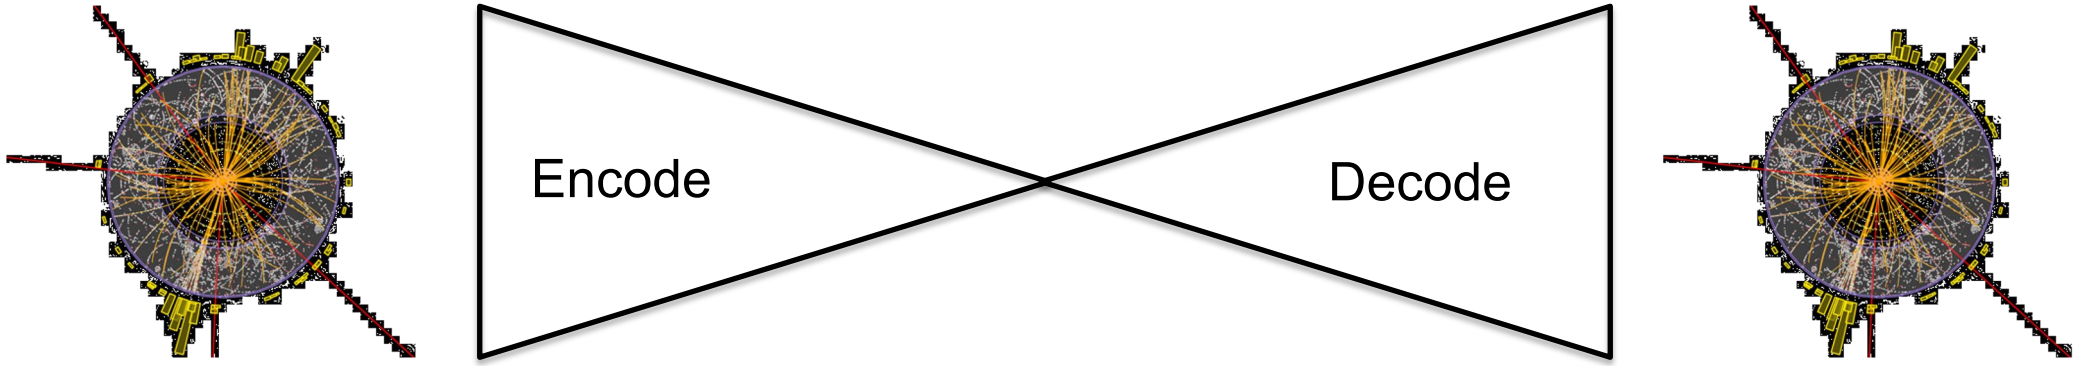

Poor reconstruction = *anomaly*

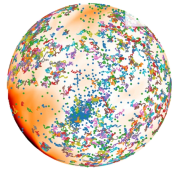

[NAE]

Challenges:

- Outlier in high-dimensional space
- Performance (e.g. anomaly metric dominated by mass)
- Add physics priors without becoming supervised

Jet level [1808.08979, 1808.08992, 2007.01850, 2301.04660…]
Event level [1806.02350, 2105.14027…]

Learning high-D background templates*

	Learn from simulation	Learn from data (SB)
Modeling the likelihood ratio	SALAD 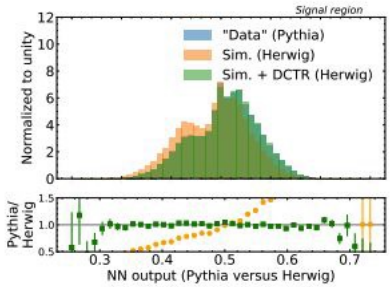	CATHODE* 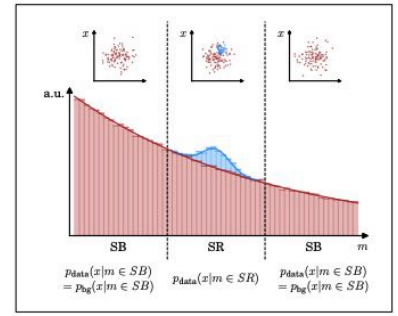 [*see also LaCATHODE & ANODE]
Morphing the features	FETA 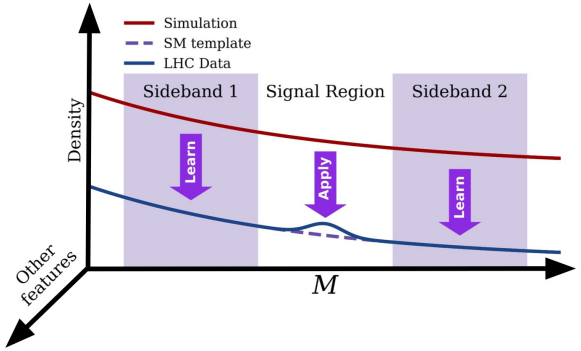	CURTAINs & Flow4Flows 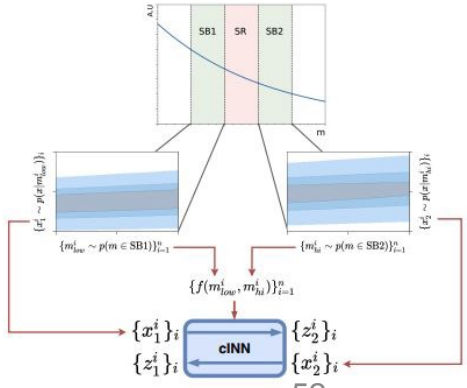

[*Fidelity of simulation alone insufficient]

In-situ background modeling for bump hunt

[Predicted by stable-diffusion-animation]

What would a **SB background datapoint [apple tree]** look like if it had a **SR mass [age]** value?

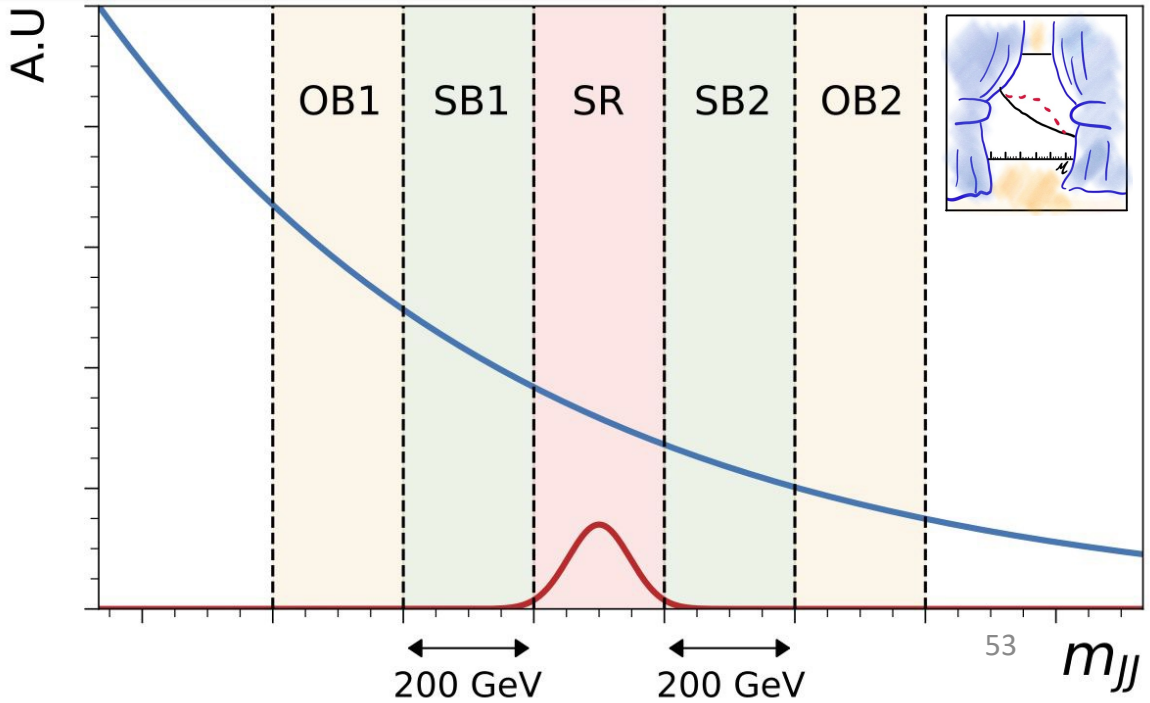

Classification without labeling (CWoLa)

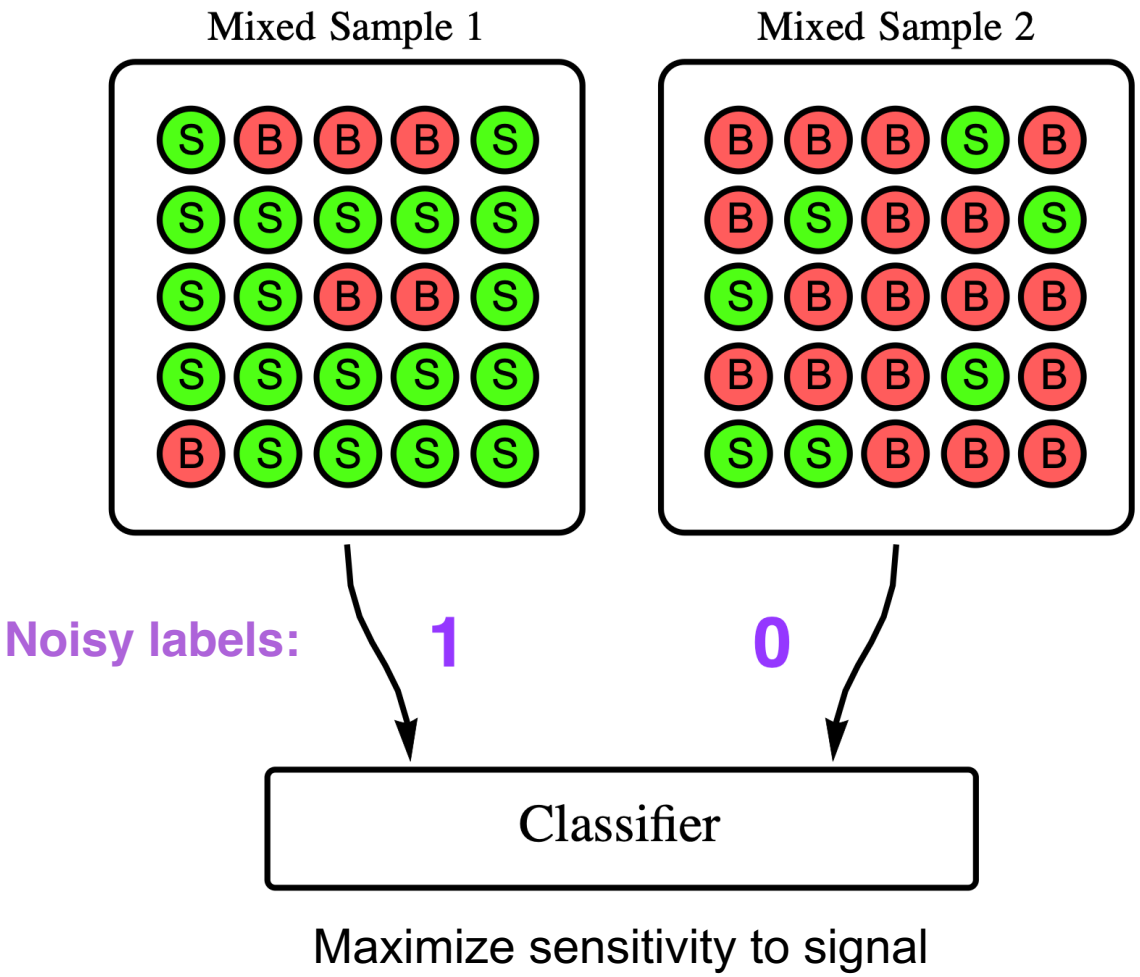

Abandon notion of *event label*

Noisy labels to be S or B

Bump hunt [1902.02634]
ATLAS analysis [2005.02983]

Beyond resonances
e.g. symmetries [2203.07529]

Comparison of methods

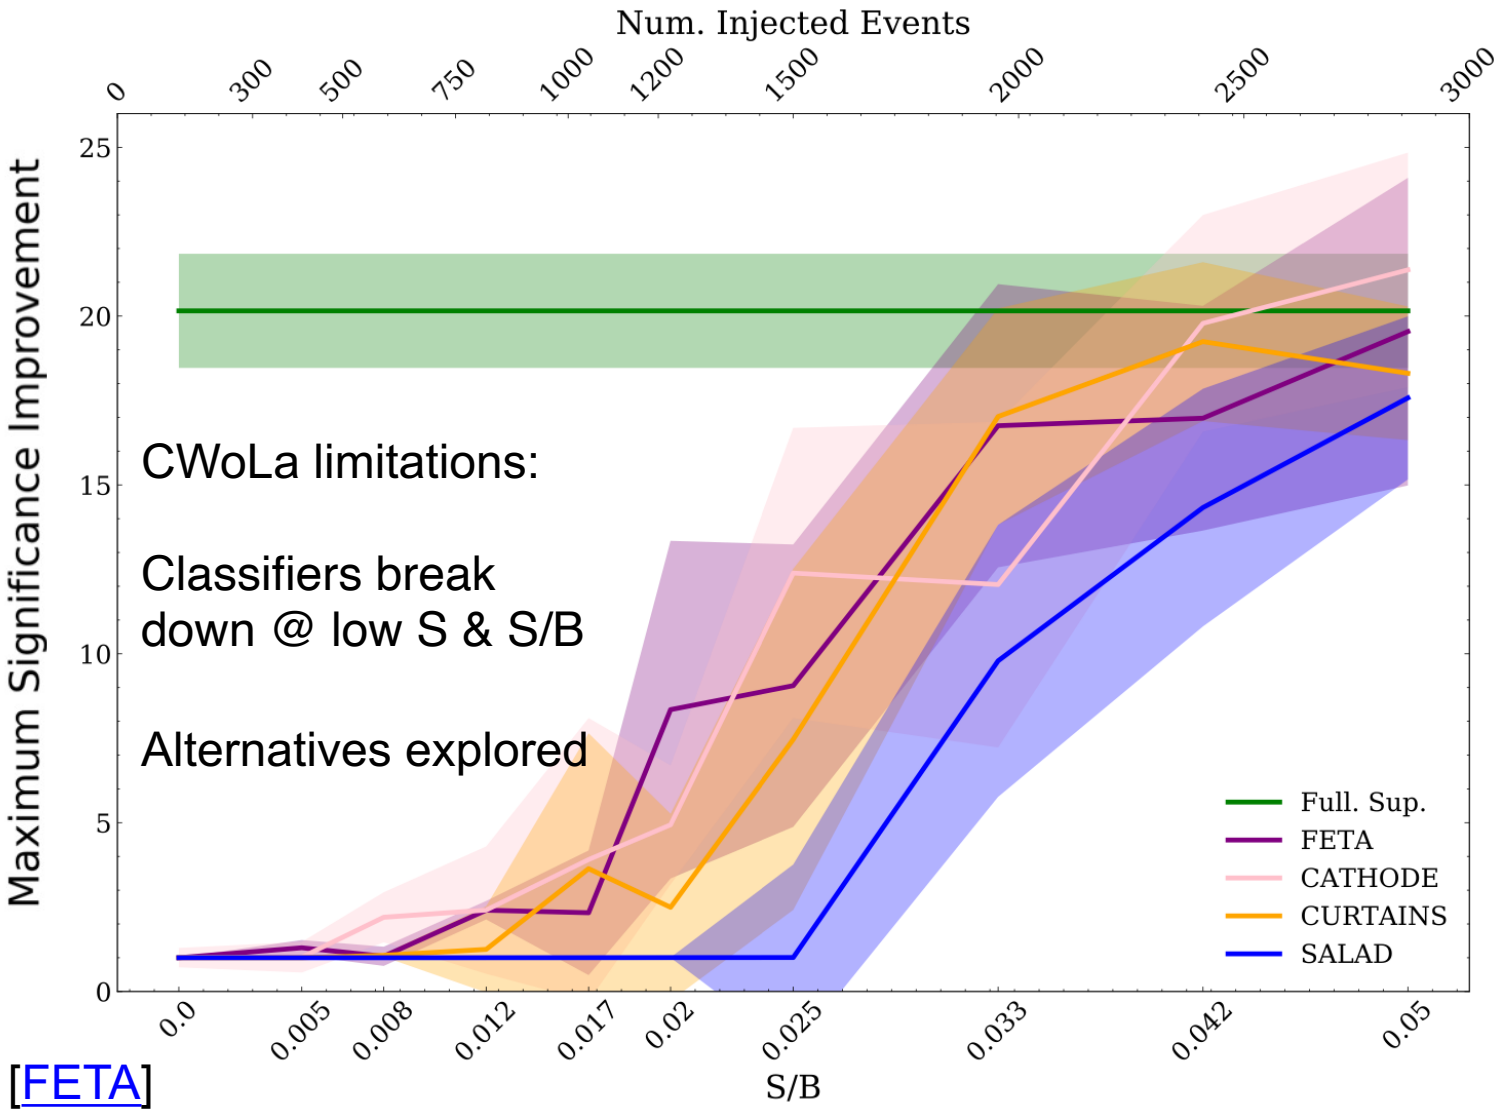

[FETA]

Similar performance of methods

Study complementarity & sensitivity to # & *noisiness* of features

Questions beyond in-situ modeling + CWoLa

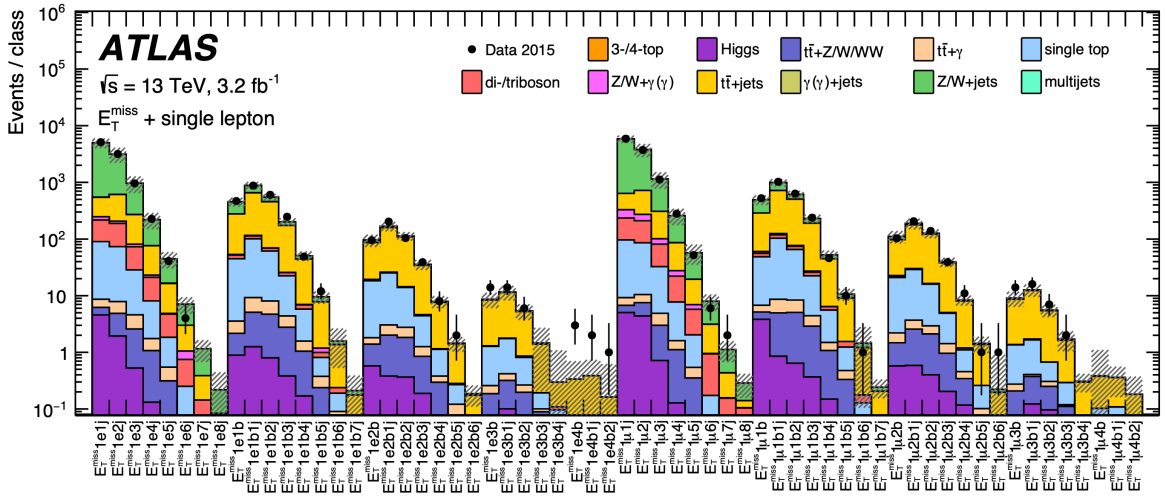

10^5 signal region [1807.07447]

Automation

The choice of feature space: there is no single good summary statistics

Data slicing & #tests [look elsewhere effect]

Dial up/down the physics prior

Interpretation w/o benchmarks

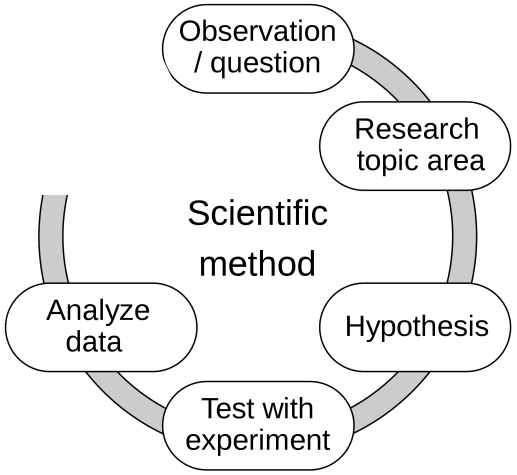

How to interpret null results?

We do not know what it is that we have not found

Poor man's assessment: *benchmarking*

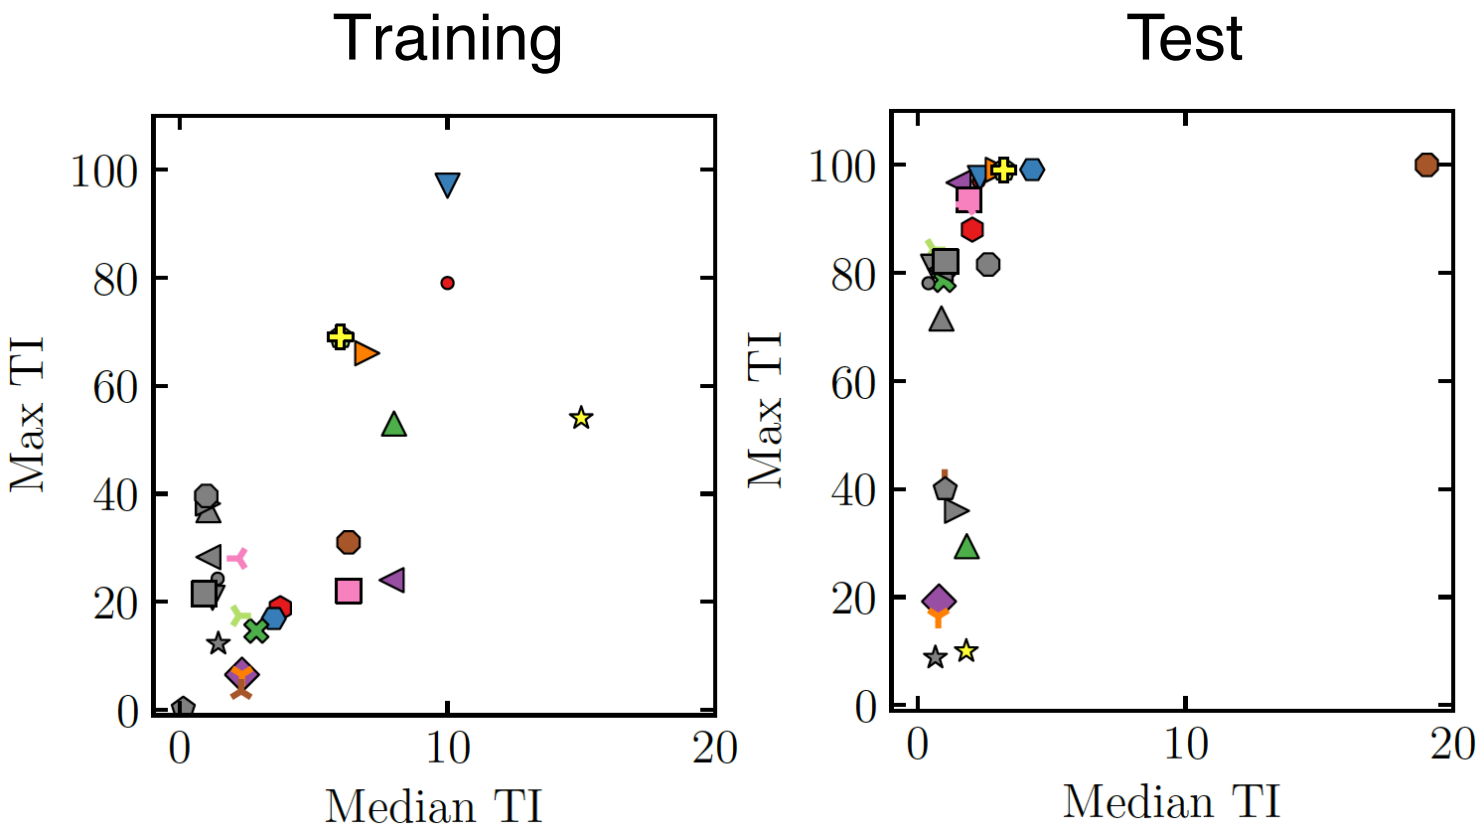

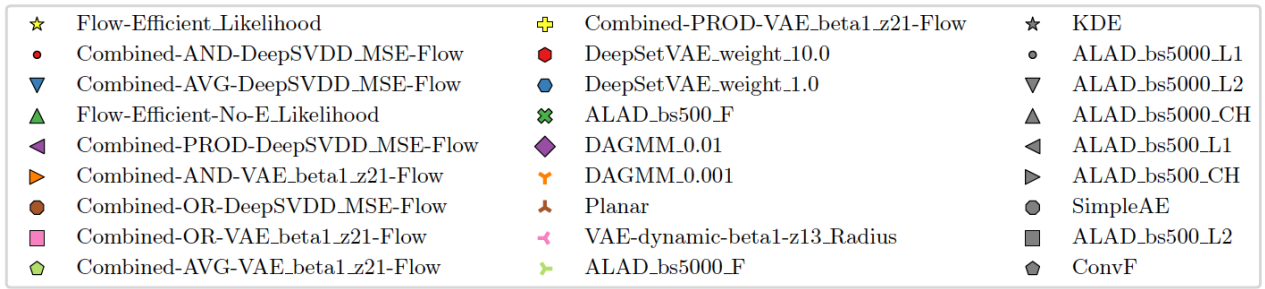

TI = Total Improvement: significance improvement over many signal benchmarks

- Max TI
- Median TI
- (Min TI)

<u>Beyond benchmarking</u>: Make *reinterpretation* possible

Test many anomaly models

[2105.14027]

Recastability to close the loop

Smart sampling with active learning: **simulate on demand**

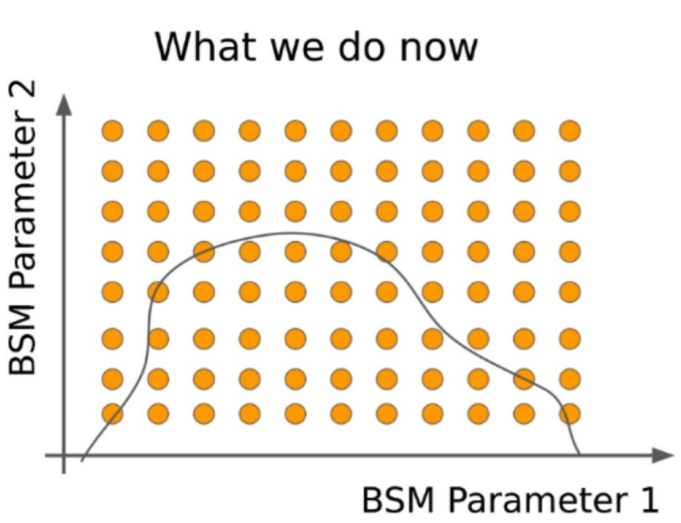

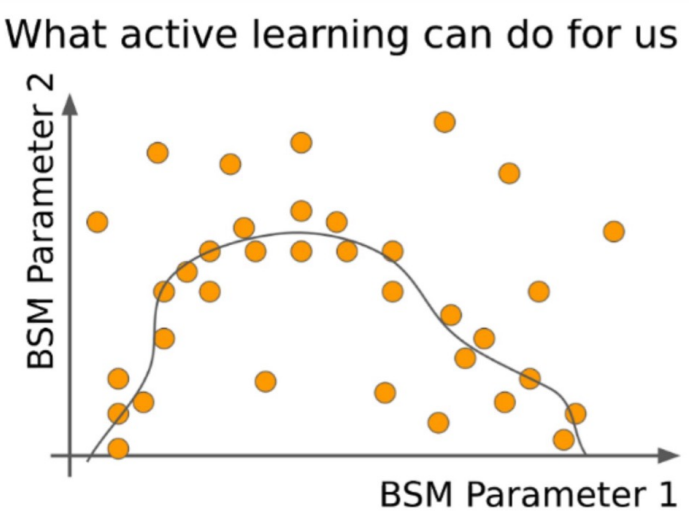

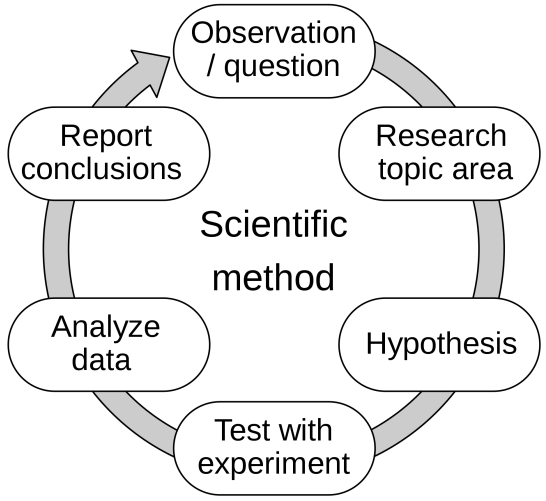

Thrives on high-dimensional space

[Regress on upper limit with GP for ATLAS mono-H(bb) search, ACAT22, ATL-PHYS-PUB-2022-045]

Quantifying search capability

*Volume in embedded space,
adjusted ROC: 2208.05484
Human-interpretable?

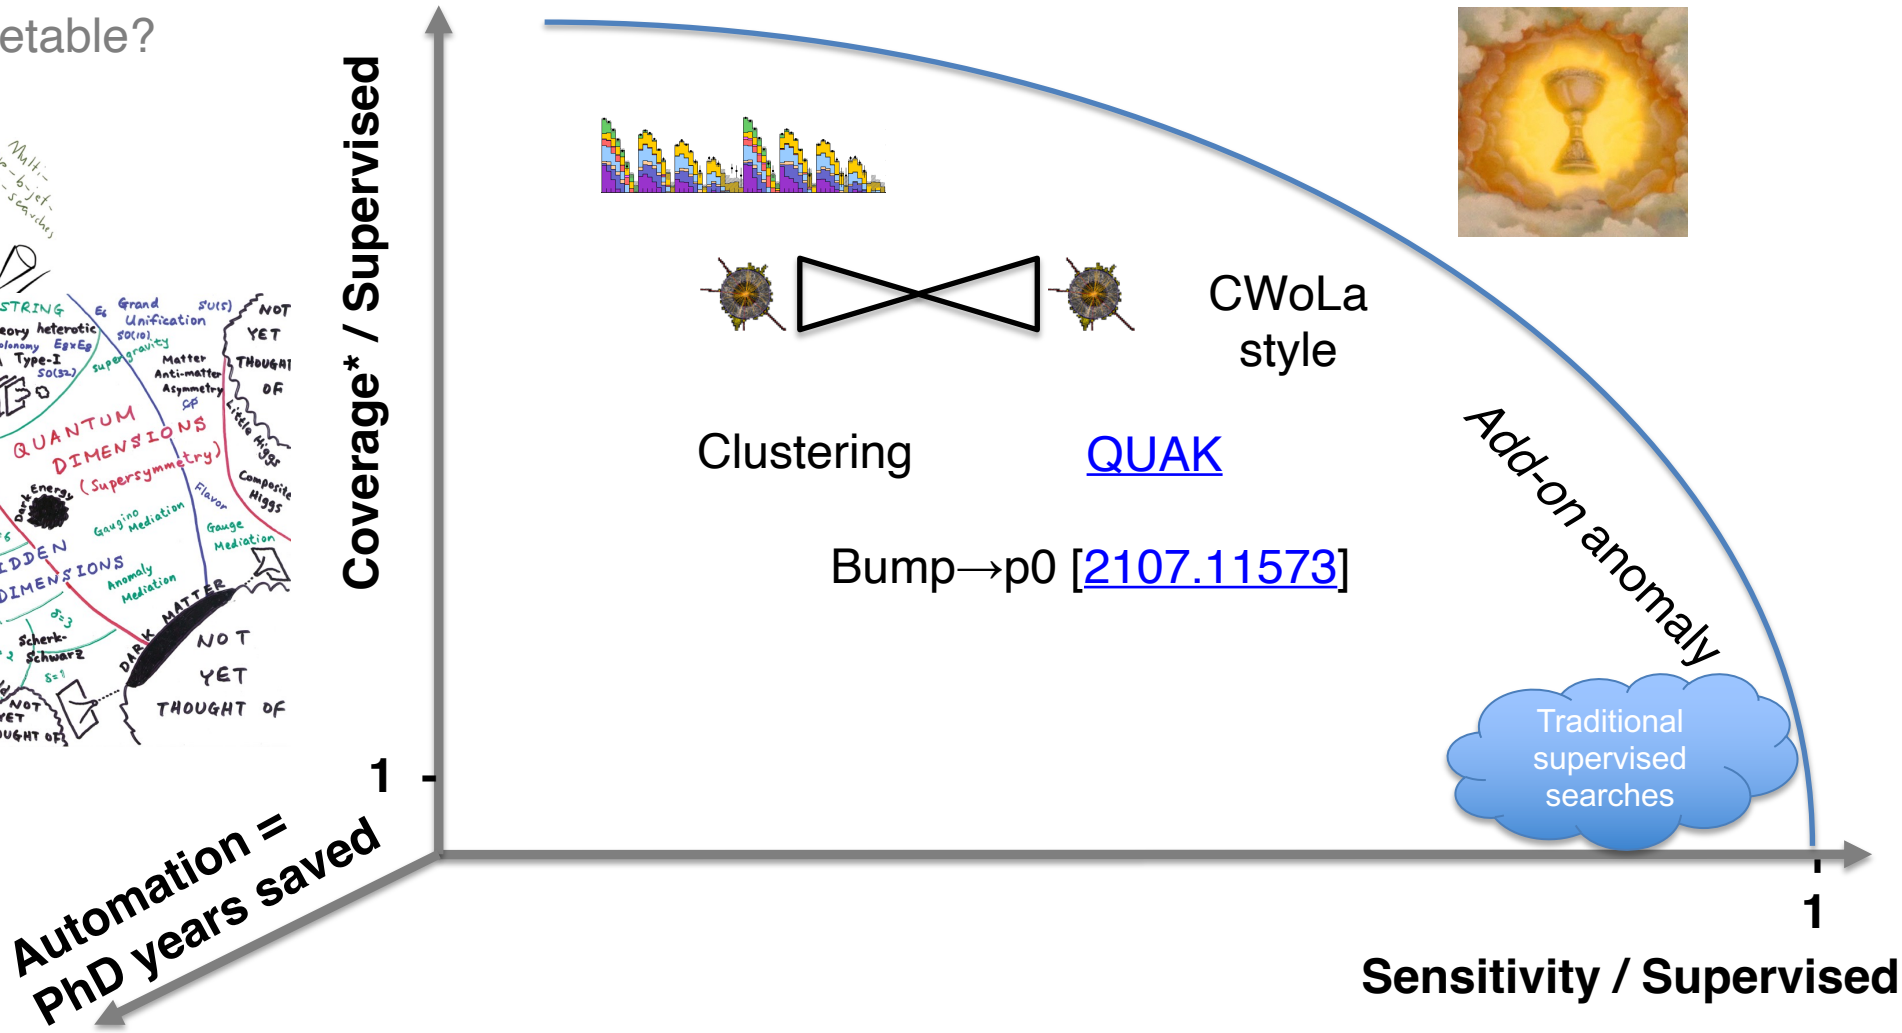

Looking ahead [speculative]

The power of foundation models [LLM]

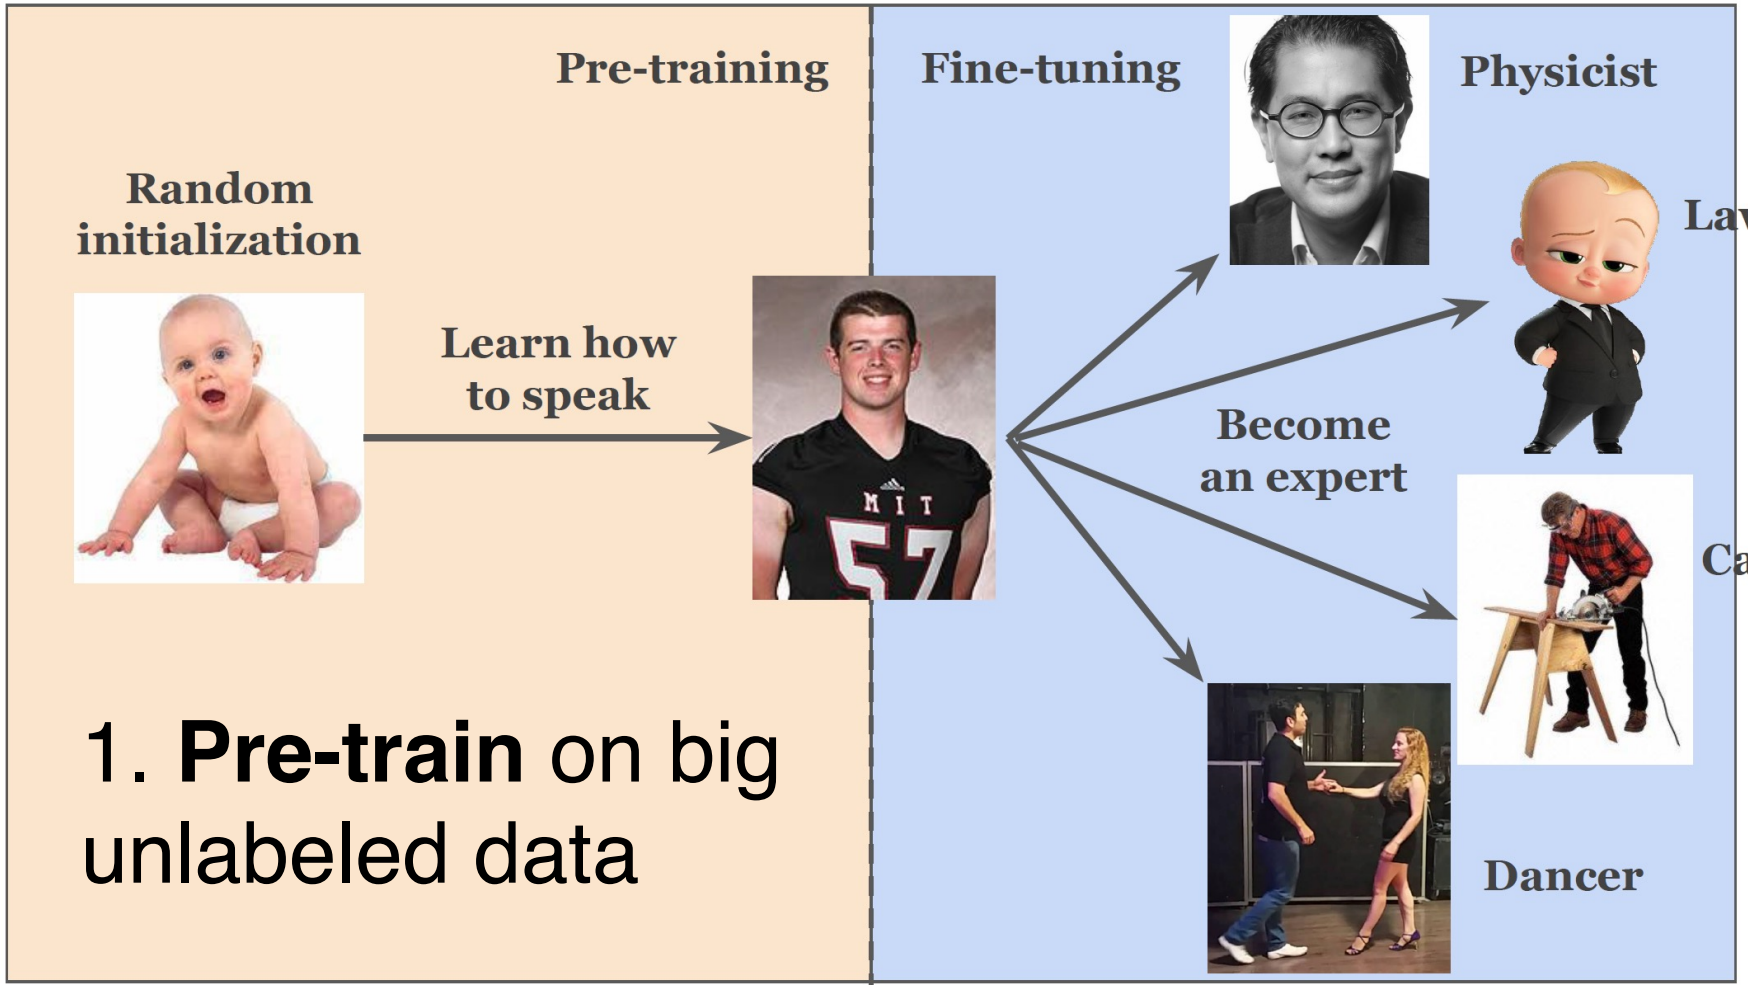

Data, memory & compute hungry

1. **Pre-train** on big unlabeled data

2. **Fine-tune** on labeled data + transfer learning

[Image credit: Kazuhiro Terao]

ChatGPT for HEP? – *Maximalist* ML

Big common pre-trained feature extractor:
Low-level features → Truth (e.g. *Higgs score)*

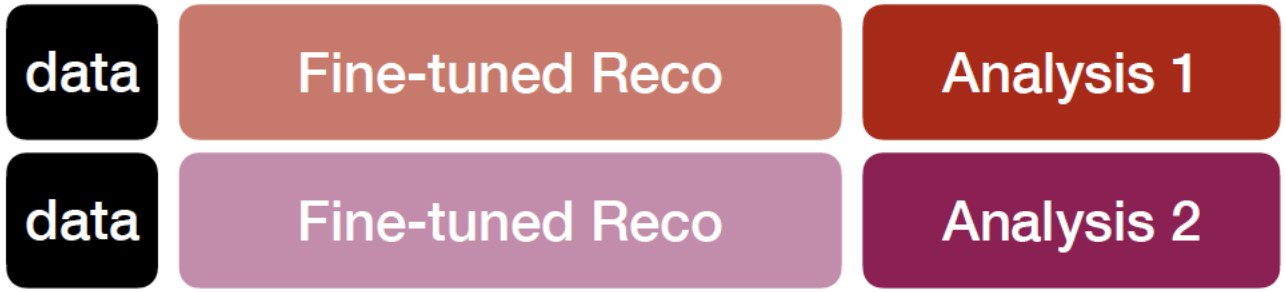

<u>Open question</u>:
One backbone > ∑ backbones per object ?

[Lukas Heinrich, H&Ns 2022]

Energy Frontier Data

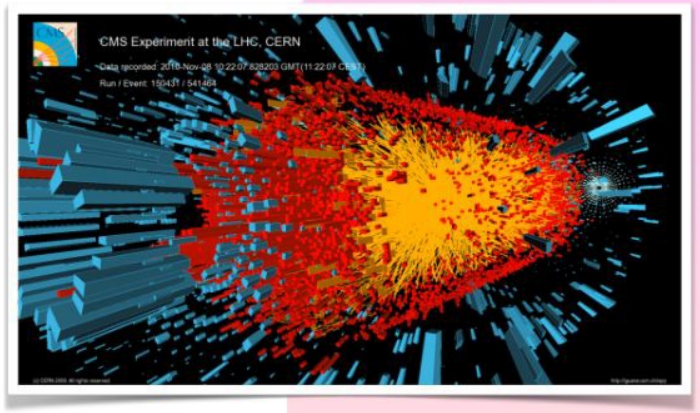

HEP FM Ecosystem

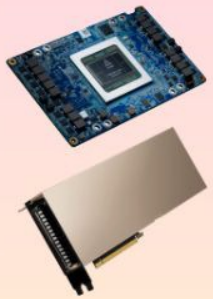

Training

HEP Foundation Model

Adaptation

Tasks

Jet Tagging

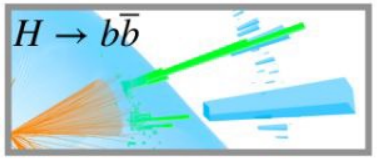

Tracking

Particle-Flow Reconstruction

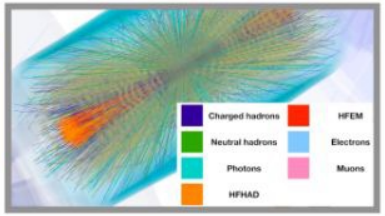

Pileup Mitigation

ν Energy Regression

ν Event Identification

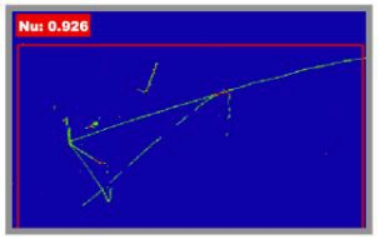

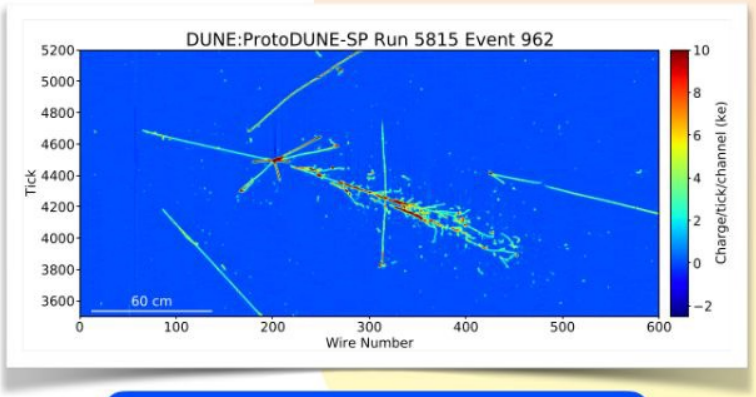

Intensity Frontier Data

[Image credit: Javier Duarte]

Towards a discussion

Many more challenges

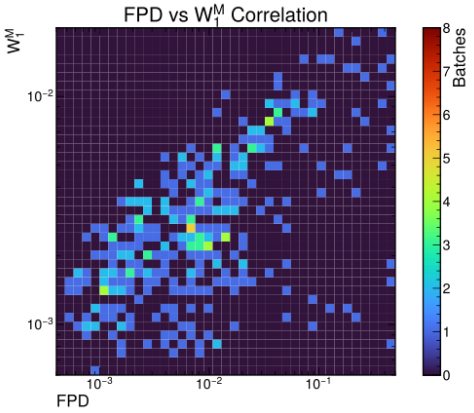

Evaluation of gen. models
[Compare metrics 2211.10295]

Decorrelation
[*Ethical AI* in Science]
[e.g. 2211.02486]

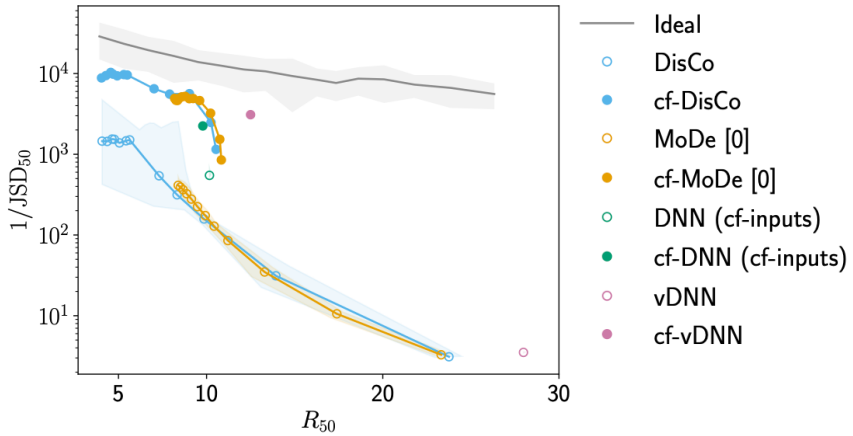

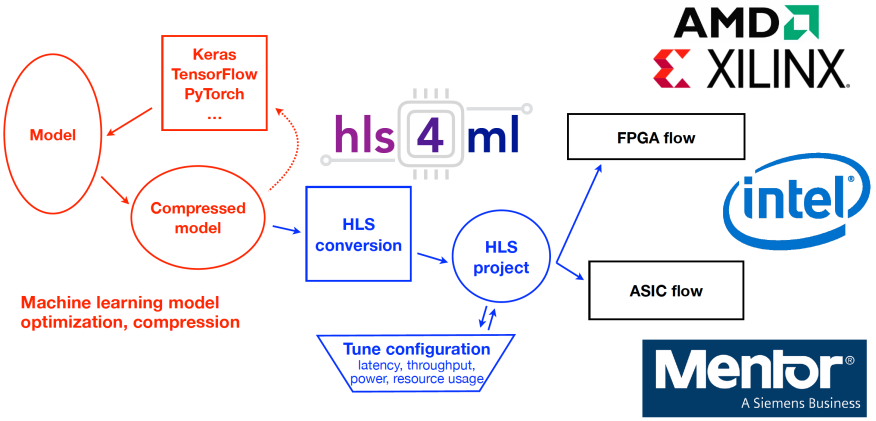

Offline → online
[On-the-edge, 1804.06913, hls4ml]

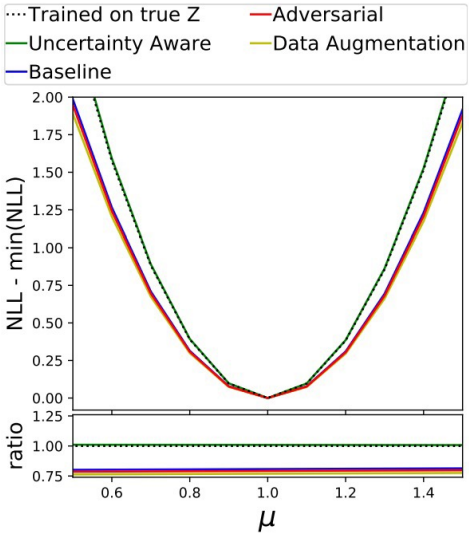

Making scientific decisions in the presence of uncertainties
[e.g. 2105.08742]

& Social challenges

Fast-moving ML ↔ **Slow** Experiment time scale

ML@HEP **competitive** ↔ ***Open Science*** @ Experiment

Need faster **concept-to-production** cycle

& Opportunities

- AI as a *muse* to science
  - ML to suggest new theories [active learning]
- Human-in-the-loop AI
  - Optimal detector design assisted by AI
- Differentiable programming → differentiable physics
- Data analysis in theory space [simulation-based inference]
- Diverse AI-assisted search portfolio [rigor/bias/automation]
- More use of GNNs & Transformers
- Impact of diffusion & foundation models – relevance of *language* aspect? [Feynman diagrams?]
- …

The HEP-AI ecosystem

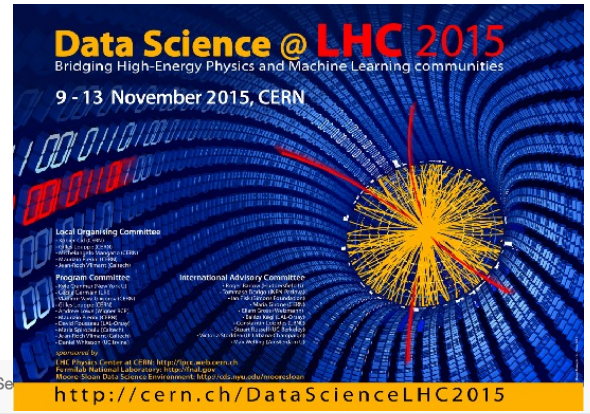

- Workshops & long-term collaborations (with industry)
  - Synergies & cross-pollination
  - Catalyst for R&D
  - Evaluate & compare
  - Community consensus

- Common benchmarks & metrics
  - Top-tagging reference data
  - CaloChallenge
  - Anomaly challenges
  - JetNet

Findable Accessible Interoperable Reusable

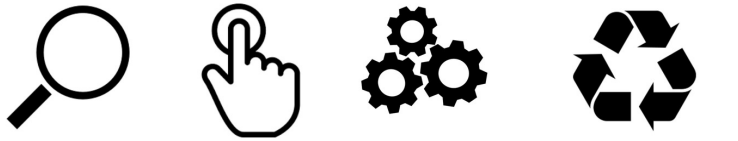

Journal of Brief Ideas — Home | New idea | Trending ideas | All ideas | About

Create standalone simulation tools to facilitate collaboration between HEP and machine learning community

By Kyle Cranmer, Tim Head, jean-roch vlimant, Vladimir Gligorov, Maurizio Pierini, Gilles Louppe, Andrey Ustyuzhanin, Balázs Kégl, Peter Elmer, Juan Pavez, Amir Farbin, Sergei Gleyzer, Steven Schramm, Lukas Heinrich, Michael Williams, Christian Lorenz Müller, Daniel Whiteson, Peter Sadowski, Pierre Baldi

dslhc | machinelearning | datascience | open data | simulation

Discussions at recent workshops have made it clear that one of the key barriers to collaboration between high energy physics and the machine learning community is access to training data. Recent successes in data sharing through the HiggsML and Flavours of Physics Kaggle challenges have borne much fruit, but required significant effort to coordinate.

While static simulated datasets are useful for challenges, in the course of investigating new machine learning techniques it is advantageous to be able to generate training data on demand (e.g. Refs. 1, 2, 3).
Therefore we recommend efforts be made to produce the ingredients required to facilitate such collaboration:

- Specific challenges for HEP experiments should be fully specified such that minimal domain-specific knowledge is required to attack them.
- Stand-alone simulators should be made open source. They should be developed to be easy to use without domain-specific expertise, while still being representative of real experimental challenges. Such a simulation will permit non-HEP researchers to generate realistic HEP datasets for training and testing. These simulators could range from truth-le
sensor arrays.
- Performance metric
solutions.

Sign in with ORCID

Authors

Kyle Cranmer, Tim Head, jean-roch vlimant, Vladimir Gligorov, Maurizio Pierini, Gilles Louppe, Andrey Ustyuzhanin, Balázs Kégl, Peter Elmer, Juan Pavez, Amir Farbin, Sergei Gleyzer, Steven Schramm, Lukas Heinrich, Michael Williams, Christian Lorenz Müller, Daniel Whiteson, Peter Sadowski, Pierre Baldi

Metadata

DOI 10.5281/zenodo.46864

Published: 26 Feb, 2016

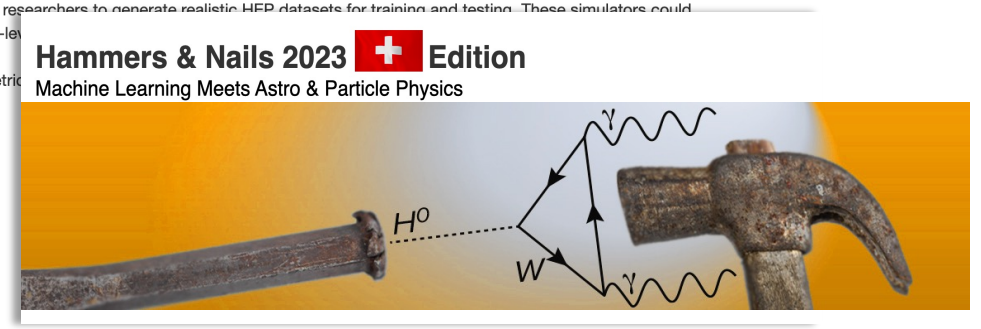

Summary

- Surrogates to efficiently model complex systems
- Transformative: automation & acceleration
- Inject physics into AI $\Leftrightarrow$ Interpretability
- Innovation $\rightarrow$ Exploitation

Outlook:

Attack problems which were considered unsolvable

ATLAS EXPERIMENT

syn·er·gy | ˈsinərjē |

UNIVERSITÉ DE GENÈVE
FACULTY OF SCIENCE

PIs

TG

PhD students

Tomke Schröer

Malte Algren

Lukas Ehrke

Matthew Leigh

Debajyoti Sengupta

Sam Klein

postdocs

Knut Zoch

Kinga Wozniak

Johnny Raine

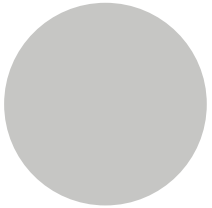

This could be you !

Slava Voloshynovskiy

Guillaume Quétant

Mariia Drozdova

Ivan Oleksiyuk

Olga Taran

François Fleuret

Bálint Máté

Atul Kumar Sinha

Daniele Paliotta

Full spectrum: physics - CS

FONDATION ERNEST BONINCHI

Alexander von Humboldt Stiftung/Foundation

Swiss National Science Foundation

Sinergia

Swiss Government Excellence Scholarships
STUDY IN SWITZERLAND